

Discussion 7

Hash Indexes &
Project 3 Intro
EECS 484

Logistics

- Homework 3 due Nov 3 at 11:55 PM EST
- Project 3 due Nov 10 at 11:55 PM EST
- Midterm exam in the process of grading
 - Please do not discuss the exam
- Today
 - Indexing & Hash index
 - Intro to MongoDB (project 3)
 - MapReduce (project 3)

Indexing

Indexes

- Our database has lots of nice data
 - How do we get the data as fast as possible
 - Indexing!
- Indexes are a type of data structure that allow us to search for, insert, and delete data in a table
 - Take an input k (not necessarily a key for the table, just a search key)
 - Set of columns
 - If search key contains a primary key then we call it the primary index
 - Otherwise it is the secondary index
 - Find data entries k^* for each k
 - Can have more than one data entry per search key k



Records

- All data are stored in a unit of space called a page
 - Multiple records can exist in a page
 - Often times needs multiple pages to fit all records in a table
- Pages don't have to be always full
 - Clustered indexes may keep pages $\frac{2}{3}$ full to make inserting and deletion faster

Data Entries k^*

- Take search key k and find data entries k^* (from the index file) for k
- Alternative 1: Data entry k^* is the actual record itself
 - Only 1 k^* can be resolved using Alternative 1 (no way to handle duplicates) ⇒ or you will have lots of duplicates
- Alternative 2: Data entry k^* is (k, rid) where rid is the record ID for the record with the search key k ↑ search key pointer to the location
- Alternative 3: Data entry k^* is $(k, \text{list of rids})$ where each rid is the record ID for the record with the search key k

one entry points to one record, one page
contains many records, to fetch one record
you need to access the whole page

Clustering

*sorted
actual data*

- **Clustered index**

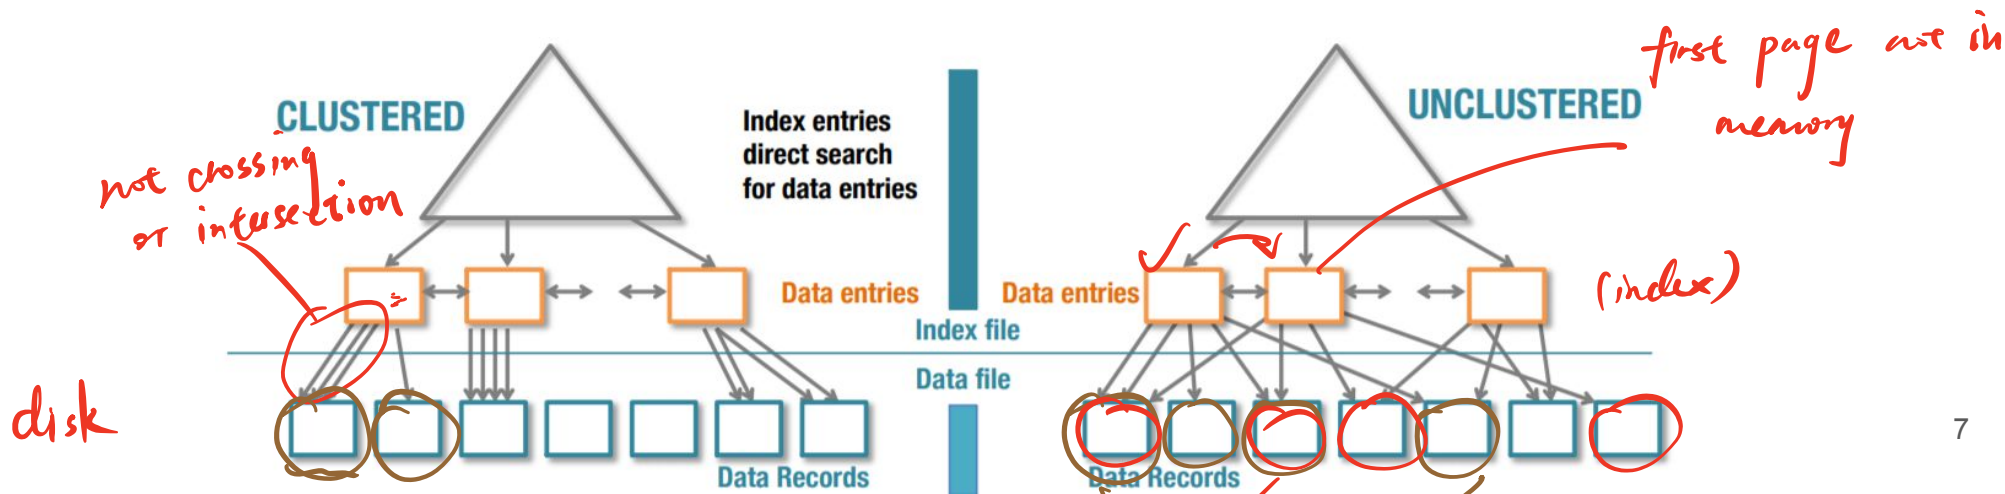
- The **data records** are in the **same order** (or very close to the same order) as the **data entries**
 - ~~★~~ Alternative 1 implies clustered index - each data entry is the data record
- Have to read each data record page once only

index

- **Unclustered index**

- No guaranteed order
- Have to **read each data record page once** for each data record we retrieve

*(in order) : finish reading first page
go to next page*



reading for first data entry: 2 pages

4 pages

read same record multiple times

Linear Hashing

Linear Hashing


take binary transformation
and look at last
2 bits

$$N = 2$$

$$H_i(x) = x \pmod{N \cdot 2^i}$$

$$\text{Level} = 0$$

- Family of hash functions
- Split one bucket at a time upon an overflow
- N = fixed base number of buckets *usually 2*
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow



① split when you create new overflow bucket.

② split when insert into an overflow bucket

H_0 : last bit of binary

H_1 : last 2 bits

Linear Hashing


- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow

$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

$$\text{Level} = 0$$

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$



Linear Hashing


- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow
- **Insert 23***

$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

$$\text{Level} = 0$$

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$



Linear Hashing


- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow
- **Insert 23***
 - $H_0(23) = 23 \pmod{2} = 1 \pmod{2} = 0b1$

$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

$$\text{Level} = 0$$

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$



Linear Hashing


- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow
- **Insert 23^***
 - $H_0(23) = 23 \pmod{2} = 1 \pmod{2} = 0b1$

$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

$$\text{Level} = 0$$

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$



Linear Hashing


- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow
- Insert 23^*
 - $H_0(23) = 23 \pmod{2} = 1 \pmod{2} = 0b1$
- Done :)

$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

$$\text{Level} = 0$$

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$



Linear Hashing


- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow
- Insert 12^*
 - $H_0(12) = 12 \pmod{2} = 0 \pmod{2} = 0b0$

$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

$$\text{Level} = 0$$

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$



Linear Hashing


- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow
- **Insert 12^***
 - $H_0(12) = 12 \pmod{2} = 0 \pmod{2} = 0b0$

$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

$$\text{Level} = 0$$

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$



Linear Hashing


- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow
- **Insert 12***
 - $H_0(12) = 12 \pmod{2} = 0 \pmod{2} = 0b0$

$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

$$\text{Level} = 0$$

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$



Linear Hashing

- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow
- Insert 12^*
 - $H_0(12) = 12 \pmod{2} = 0 \pmod{2} = 0b0$


$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

$$\text{Level} = 0$$

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$

$$H_1(x) = x \pmod{N * 2^1} = x \pmod{4}$$



Linear Hashing

- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow
- Insert 12^*
 - $H_0(12) = 12 \pmod{2} = 0 \pmod{2} = 0b0$
- Rehash $4^*, 10^*, 12^*$
 - $H_1(4) = 4 \pmod{4} = 0 \pmod{4} = 0b00$
 - $H_1(10) = 10 \pmod{4} = 2 \pmod{4} = 0b10$
 - $H_1(12) = 12 \pmod{4} = 0 \pmod{4} = 0b00$


$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

$$\text{Level} = 0$$

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$

$$H_1(x) = x \pmod{N * 2^1} = x \pmod{4}$$



Linear Hashing

- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow
- Insert 12^*
 - $H_0(12) = 12 \pmod{2} = 0 \pmod{2} = 0b0$
- Done :)


$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

$$\text{Level} = 0$$

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$

$$H_1(x) = x \pmod{N * 2^1} = x \pmod{4}$$



Linear Hashing

- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow
- **Insert 8***
 - $H_0(8) = 8 \pmod{2} = 0 \pmod{2} = 0b0$
 - $H_1(8) = 8 \pmod{4} = 0 \pmod{4} = 0b00$

$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

$$\text{Level} = 0$$

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$

$$H_1(x) = x \pmod{N * 2^1} = x \pmod{4}$$

$H_1 \ H_0$

00 0

4*	12*
----	-----

Full :(

Next = 1

01 1

1*	23*
----	-----

10

10*	
-----	--

Linear Hashing

- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow
- **Insert 8^***
 - $H_0(8) = 8 \pmod{2} = 0 \pmod{2} = 0b0$
 - $H_1(8) = 8 \pmod{4} = 0 \pmod{4} = 0b00$


$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

$$\text{Level} = 0$$

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$

$$H_1(x) = x \pmod{N * 2^1} = x \pmod{4}$$



Linear Hashing

- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow
- **Insert 8^***
 - $H_0(8) = 8 \pmod{2} = 0 \pmod{2} = 0b0$
 - $H_1(8) = 8 \pmod{4} = 0 \pmod{4} = 0b00$


$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

$$\text{Level} = 0$$

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$

$$H_1(x) = x \pmod{N * 2^1} = x \pmod{4}$$



Linear Hashing

- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow

- Insert 8^*
 - $H_0(8) = 8 \pmod{2} = 0 \pmod{2} = 0b0$
 - $H_1(8) = 8 \pmod{4} = 0 \pmod{4} = 0b00$
 - Remember we split the next node always even though it might not be overflow


$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

$$\text{Level} = 0$$

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$

$$H_1(x) = x \pmod{N * 2^1} = x \pmod{4}$$



Linear Hashing

- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow

- **Insert 8^***

- $H_0(8) = 8 \pmod{2} = 0 \pmod{2} = 0b0$
- $H_1(8) = 8 \pmod{4} = 0 \pmod{4} = 0b00$
- Remember we split the next node always even though it might not be overflow


$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

Level = 1

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$

$$H_1(x) = x \pmod{N * 2^1} = x \pmod{4}$$



Linear Hashing

- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow
- **Insert 8^***
 - $H_0(8) = 8 \pmod{2} = 0 \pmod{2} = 0b0$
 - $H_1(8) = 8 \pmod{4} = 0 \pmod{4} = 0b00$
 - $H_1(1) = 1 \pmod{4} = 1 \pmod{4} = 0b01$
 - $H_1(23) = 23 \pmod{4} = 3 \pmod{4} = 0b11$


$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

$$\text{Level} = 1$$

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$

$$H_1(x) = x \pmod{N * 2^1} = x \pmod{4}$$



Linear Hashing

- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow
- **Insert 8^***
 - $H_0(8) = 8 \pmod{2} = 0 \pmod{2} = 0b0$
 - $H_1(8) = 8 \pmod{4} = 0 \pmod{4} = 0b00$
 - $H_1(1) = 1 \pmod{4} = 1 \pmod{4} = 0b01$
 - $H_1(23) = 23 \pmod{4} = 3 \pmod{4} = 0b11$
 - Still have overflow page; that's ok


$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

$$\text{Level} = 1$$

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$

$$H_1(x) = x \pmod{N * 2^1} = x \pmod{4}$$



Linear Hashing

- Family of hash functions
- Split one bucket at a time upon an overflow
- $N = \text{fixed}$ base number of buckets
- Level = current level in hash family
- Next = pointer to next bucket to be split
- Split policy: split on overflow
- **Insert 8^***
 - $H_0(8) = 8 \pmod{2} = 0 \pmod{2} = 0b0$
 - $H_1(8) = 8 \pmod{4} = 0 \pmod{4} = 0b00$
 - $H_1(1) = 1 \pmod{4} = 1 \pmod{4} = 0b01$
 - $H_1(23) = 23 \pmod{4} = 3 \pmod{4} = 0b11$
 - Still have overflow page; that's ok
 - Done :)


$$N = 2$$

$$H_i(x) = x \pmod{N * 2^i}$$

$$\text{Level} = 1$$

$$H_0(x) = x \pmod{N * 2^0} = x \pmod{2}$$

$$H_1(x) = x \pmod{N * 2^1} = x \pmod{4}$$




Extendible Hashing

Extendible Hashing

- Use directory of pointers


- Split on overflow
- Once we're out of room in directory, double size
- Global depth = number of bits considered globally
- Local depth = number of bits considered locally

$$GD \geq LD$$




Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- Insert 5*




Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- Insert 5*
 - 5=0b101




Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- Insert 5*
 - 5=0b101




Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- Insert 5*
 - 5=0b101
 - Done :)




Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- Insert 10*




Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- Insert 10*
 - 10=0b1010




Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- Insert 10*
 - 10=0b1010



Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- **Insert 10***
 - 10=0b1010




Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally


- **Insert 10***

- 10=0b1010
- New directories point to "Split image"




Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- **Insert 10***
 - 10=0b1010
 - New directories point to "Split image"
 - Try to insert again




Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- Insert 10*
 - 10=0b1010
 - New directories point to "Split image"
 - Try to insert again




Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- **Insert 10***
 - 10=0b1010
 - New directories point to "Split image"
 - Try to insert again




Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- Insert 10*
 - 10=0b1010
 - New directories point to "Split image"
 - Try to insert again




Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- Insert 10*
 - 10=0b1010
 - New directories point to "Split image"
 - Try to insert again
 - Done :)




Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- Insert 7*




Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- **Insert 7***
 - **7=0b111**




Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- **Insert 7***
 - 7=0b111




Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- **Insert 7***
 - 7=0b111



Extendible Hashing


- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- Insert 7*
 - 7=0b111



Extendible Hashing

- Use directory of pointers
 - Split on overflow
 - Once we're out of room in directory, double size
 - Global depth = number of bits considered globally
 - Local depth = number of bits considered locally
- Insert 7*
 - 7=0b111
 - Done :)

☆ what if there are still overflow page after split? \Rightarrow keep splitting



Hashing Overview

- Directory size can double in extendible hashing
 - Linear hashing only adds one bucket at a time
 - Global depth \geq local depth always
 - # pointers to any specific bucket $\leq 2^{GD-LD}$
 - Overflow pages only in rare cases (unavoidable collisions)
- Linear Hashing only adds one bucket at a time
 - Better memory usage
 - Doesn't avoid overflow pages in many cases
 - Over time we minimize overflow pages

e.g. duplicates




Hash browns < hash tables

MongoDB and NoSQL

NoSQL

- Not Only SQL!
 - Non relational databases
 - Use very different data structures compared to traditional relational databases
- Reading: <https://www.mongodb.com/nosql-explained>
- Different data models used by different distributions
- Data in a common set doesn't need to adhere to a schema

HOW TO WRITE A CV




Leverage the NoSQL boom

MongoDB

- Instead of rows we have documents
 - Fields and values
 - I.E "Age": "7"
- Instead of relations/tables we have collections
- Use Javascript instead of SQL to interact
- Data is in JSON
- <https://docs.mongodb.com/manual/reference/sql-comparison/>


= table in SQL



JSON

- JavaScript Object notation
- {Key: Value}
 - `var data = {"Name": "Alice", "Major": "CS", "University": "UofM", "Hobby": "Beating MSU"};`
 - Can have nested key values as well:
 - `{"Location": {"City": "Ann Arbor", "State": "Michigan", "Country": "USA"}};`
- Retrieve data by `data["Name"]`
 - Returns "Alice"
- JSON Objects are not ordered

nested key-value pair



note: JSON has nothing to do with Jason Momoa, but again, they sound similar

MapReduce

MapReduce

processing large amount of data in parallel



- Programming paradigm to calculate aggregate analytics
- Designed to be highly parallelized, which works great when we have a large set of data and many processors
 - Historically performed on Hadoop (distributed storage) framework
- Often performs worse compared to an equivalent serial algorithm working on a single machine
 - But we will use it anyway with MongoDB in P3, to get practice with the paradigm

MapReduce

- Needed for queries 7 and 8
- Mapper
 - Run on each record
 - Emits (a key and a value)
- Reducer
 - Run on each mapped object
 - Combine values in aggregate that share key to get value of interest
 - **It's possible to have multiple calls to reduce!**
- Finalizer
 - Specific to MongoDB
 - Performs any last calculations before returning in final form

intermediate key-value

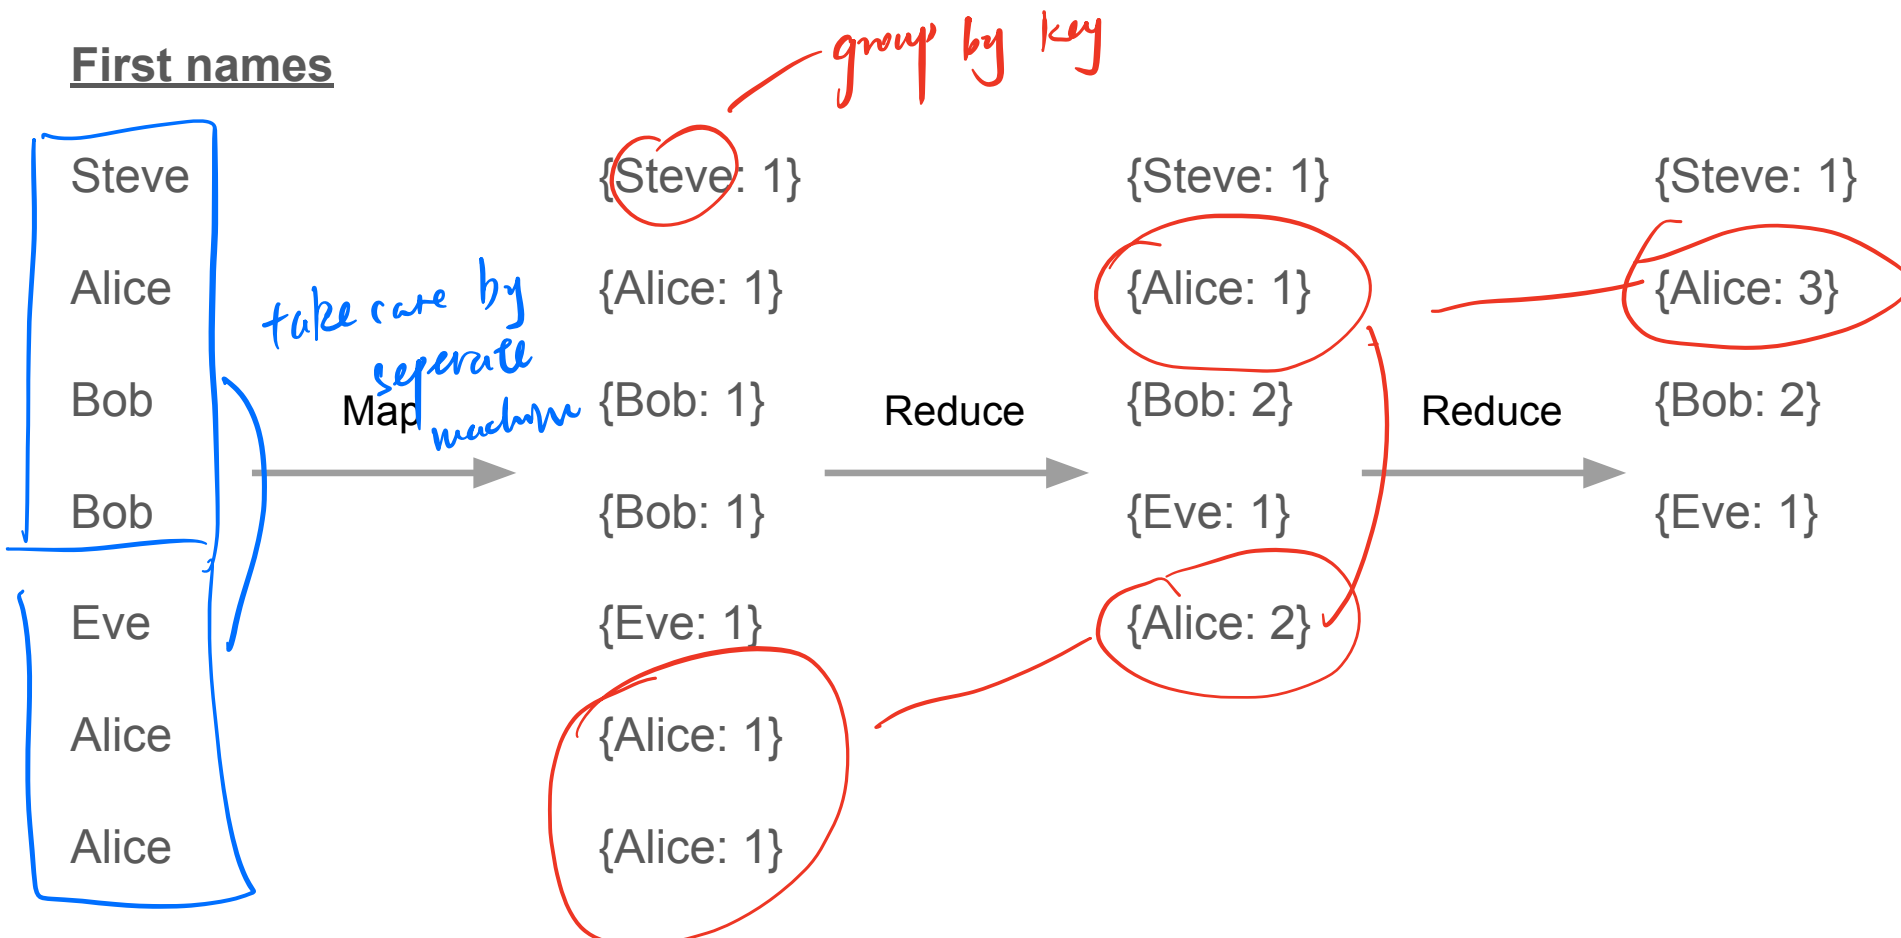
grouping :

key [x, y, z]



MapReduce

First names



Project 3 Intro

Project 3

- Part A: Java code to export database to JSON
 - Need to perform this on CAEN just like in Project 2
- Part B: MongoDB Queries
 - You can run entirely on your local machine (requires installing MongoDB)
 - We won't be able to help with specific installation issues though
 - Use the eecs484 server through CAEN
 - Set up your project on CAEN and edit the Makefile as specified in the spec
 - Run commands which will connect to a MongoDB server setup for you on the eecs484 server
 - Write the queries!
 - Lots of helpful references in the spec to various MongoDB documentation

General Tips

- Try to focus on getting the solution correct rather than doing it in the fanciest way
 - Yes you can use an aggregate, group, out pipeline or you can iterate a couple times
 - No efficiency tests
 - No private tests
 - Most important thing: make sure you understand how and why your code works
- Take the time to get familiar with Javascript
 - Documentation will be your best friend
- Query 5 is the hardest
 - Has a similar concept to how you dealt with the friends relation in Query 6 on Project 2
 - Completely different code but similar work around needed

Query 7

key is month

- Find number of users born in each month using map reduce
- Mapper
 - Given access to a user (this) what can we return that will give us useful information?
 - Emit Tuple: (key, value)
 - Hint: Think about a way to mark that this user was born in this month
- Reducer
 - Given access to a set of values that correspond to a key, how do we combine these values?
 - Return value
 - Should match the same time of value output by the mapper
 - Hint: The value output at the end of this method should be the number of users (that we know of) born in the month denoted by key
- Finalizer
 - Get the final answers for each key ready to return
 - No change needed for query 7 :)

Query 8

- Find average friend count per city

city as the key

- Mapper

- Given access to a user (this) what can we return that will give us useful information?
- Emit Tuple: (key, value)
- Hint: Value can be a tuple in it of itself!

- Reducer

- Given access to a set of values that correspond to a key, how do we combine these values?
- Return value - should match the same type of value output by the mapper
- Hint: The reducer can be called multiple times during execution
 - We can take intermediate sums, but not intermediate averages

✗ don't calculate avg in reducer

- Finalizer

- Get the final answers for each key ready to return
- Small change needed for query 8
- Hint: We can't compute an average in the reducer, but we can in the finalizer

← calculate here

- Find average of {5, 3, 9, 5, 13} given that they all map to the same key
 - Incorrect solution:
 - Reducer: Find average of 5 and 3 (4)
 - Reducer: Find average of 9 and 5 and 13 (9)
 - Reducer: Find the average of the averages (gives us 6.5 when the real answer is 7)
 - Correct solution
 - Reducer: sums 5 and 3, and outputs 8
 - Reducer: sums 9 and 5 and 13, and outputs 27
 - Reducer: sums 8 and 27, and outputs 35
 - Finalizer: computes the average $35/5$ to output 7 as the average

Get started on HW4 and Project 3!