

Discussion 9

Sorting
EECS 484

Logistics

- Homework 5 is due on Thursday, Nov 17 at 11:55 PM
- Project 3 is due this Thursday, Nov 10 at 11:55 PM
- Project 4 is out Nov 10 and due on Dec 8. You will have 4 weeks to complete P4!
- Final exam is on Dec 13, which immediately follows the due date of this project, so please be aware of this when scheduling your time!

External Sorting

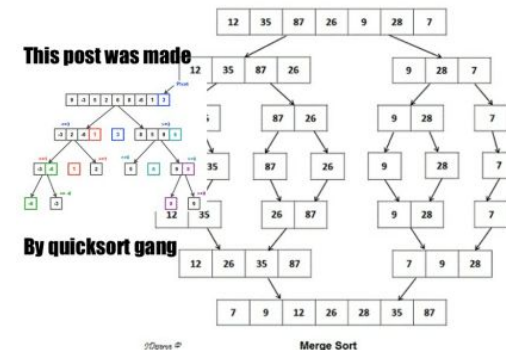
External Sorting

- Sorting is nice
 - We have lots of nice algorithms that will sort for us
 - Quicksort, Mergesort, Heapsort, etc.
 - We can do this very quickly with lots of data - $O(N \log N)$
- But what if we have too much data to fit in RAM?
 - We can still sort but it will be so so slow :(
 - Need some way to *externally* sort the data on the disk while dealing with limited fast memory

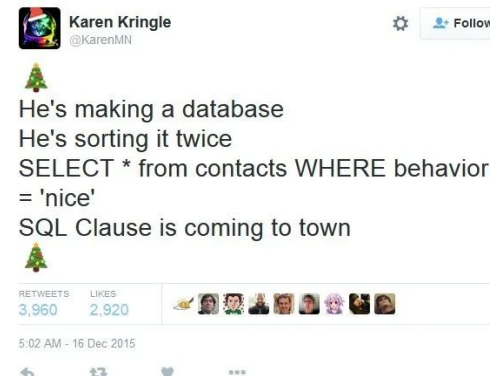
When someone asks you to prove the average case time complexity of merge sort:

$O(\text{no})$

Imagine requiring



$O(n)$ memory to sort



General External Merge Sort


- Step 1:
 - Have a large dataset of N pages that you would like to sort using B buffer pages
- Step 2:
 - Divide the dataset into $\lceil N/B \rceil$ runs (each of which is B pages long)
- Step 3:
 - Sort each run by itself normally using your favorite algorithm
 - We can fit the entire run of B pages into our RAM so no problem
- Step 4:
 - Sort the runs amongst each other
 - We can merge $B-1$ runs at a time
 - $B-1$ pages for each run plus 1 page to store the output
 - Each run is larger than 1 page though!
 - Load the first (sorted) page of each run and once it's empty, read the next page
 - Similarly, write the output buffer each time we run out of space and keep going

disk

$N \gg B$

in memory

$\lceil \frac{N}{B} \rceil$




Step 1

- Have a dataset

a bar = one element

Suppose $B = 4$ and
each page can hold
2 bars in full.

$4 \times 2 = 8$ bars




Step 2

4 pages (8 bars) - long


- Divide the data into $\lceil N/B \rceil$ runs
 - Each is B pages long, i.e. each run is technically supposed to have 8 bars
 - (for simplicity we only show 4 smallest bars in each run)

Suppose $B = 4$ and
each page can hold
2 bars in full.




Step 3

- Sort each run individually (for simplicity we only show 4 smallest bars in each run)



Step 4

- Sort the runs with each other
 - B-1 runs at a time




(for simplicity we only show 4
smallest bars in each run)

Step 4


Suppose $B = 4$ and
each page can hold
2 bars in full.

- Sort the runs with each other
 - $B-1$ runs at a time



$B-1$

Load 1 (sorted) page at
a time from each run ~~✖✖~~ $1 \times 2 = 2 \text{ bars}$



Single output page




(for simplicity we only show 4
smallest bars in each run)


Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.

- Sort the runs with each other
 - $B-1$ runs at a time



Take minimum element
from all loaded pages ✖
Remember Merge Sort




(for simplicity we only show 4
smallest bars in each run)


Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.

- Sort the runs with each other
 - $B-1$ runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort




(for simplicity we only show 4
smallest bars in each run)


Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.

- Sort the runs with each other
 - B-1 runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort




(for simplicity we only show 4
smallest bars in each run)


Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.

- Sort the runs with each other
 - B-1 runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort




(for simplicity we only show 4
smallest bars in each run)

Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.

- Sort the runs with each other
 - $B-1$ runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort




(for simplicity we only show 4
smallest bars in each run)




Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.


- Sort the runs with each other
 - $B-1$ runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort



Output page full




(for simplicity we only show 4
smallest bars in each run)


Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.

- Sort the runs with each other
 - $B-1$ runs at a time




Take minimum element
from all loaded pages
Remember Merge Sort



Write to disk
Empty page and continue




(for simplicity we only show 4
smallest bars in each run)




Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.


- Sort the runs with each other
 - $B-1$ runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort




(for simplicity we only show 4
smallest bars in each run)




Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.


- Sort the runs with each other
 - $B-1$ runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort



(for simplicity we only show 4
smallest bars in each run)




Step 4




Suppose $B = 4$ and
each page can hold
2 bars in full.

- Sort the runs with each other


- B-1 runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort




(for simplicity we only show 4
smallest bars in each run)





Step 4

Suppose $B = 4$ and
each page can hold
2 bars in full.


- Sort the runs with each other
 - $B-1$ runs at a time



Take minimum element
from all loaded pages
Remember Merge Sort




(for simplicity we only show 4
smallest bars in each run)



Step 4



Suppose $B = 4$ and
each page can hold
2 bars in full.

- Sort the runs with each other
 - $B-1$ runs at a time




(for simplicity we only show 4
smallest bars in each run)

Continue till runs sorted
Will need to make more
passes if more than $B-1$ runs



9 → 3 → 1



General External Merge Sort Math ✓

merge cost

- We have a dataset with N pages

- We'll use B buffer pages
- We'll have $\lceil N/B \rceil$ runs initially
- We need to make passes over the runs until entire dataset is sorted

$$\lceil \log_{B-1} \lceil \frac{N}{B} \rceil \rceil + 1 = \text{pass for sorting each run at the beginning}$$

- We merge $B-1$ runs together at a time
- That means we have $\lceil \lceil N/B \rceil / (B-1) \rceil$ merged runs afterwards
- Each time we make a pass we've merged all runs in sets of size $B-1$
- We must continue to do this till we have 1 output dataset in sorted order
- Takes $1 + \lceil \log_{B-1} \lceil N/B \rceil \rceil$ passes

- Total IO cost is #passes * $2N$

→ read/write pages to output

- Each pass we read each page and write each page in a new sorted order

N read N write

Example

$$N = 900$$

$$B = 18.$$

- We have a large dataset of 900 pages. We are going to use 18 buffer pages
 - How many passes will be required while performing a general external merge sort?

Example

- We have a large dataset of 900 pages. We are going to use 18 buffer pages
 - How many passes will be required while performing a general external merge sort?
 - $N=900, B=18$

Example

- We have a large dataset of 900 pages. We are going to use 18 buffer pages
 - How many passes will be required while performing a general external merge sort?
 - $N=900, B=18$
 - $\text{ceiling}(N/B)=50$

Example

- We have a large dataset of 900 pages. We are going to use 18 buffer pages
 - How many passes will be required while performing a general external merge sort?
 - $N=900, B=18$
 - $\text{ceiling}(N/B)=50$
 - $B-1=17$

Example

- We have a large dataset of 900 pages. We are going to use 18 buffer pages
 - How many passes will be required while performing a general external merge sort?
 - $N=900, B=18$
 - $\text{ceiling}(N/B)=50$
 - $B-1=17$
 - $\text{\#passes} = 1 + \text{ceiling}(\log_{B-1}(\text{ceiling}(N/B))) = 1 + \text{ceiling}(\log_{17}(50))$

Example

- We have a large dataset of 900 pages. We are going to use 18 buffer pages
 - How many passes will be required while performing a general external merge sort?
 - $N=900, B=18$
 - $\text{ceiling}(N/B)=50$
 - $B-1=17$
 - $\text{\#passes} = 1 + \text{ceiling}(\log_{B-1}(\text{ceiling}(N/B))) = 1 + \text{ceiling}(\log_{17}(50)) = 1 + \text{ceiling}(1.38\text{ish}) = 1 + 2 = 3$

Example

- We have a large dataset of 900 pages. We are going to use 18 buffer pages
 - How many passes will be required while performing a general external merge sort?
 - $N=900, B=18$
 - $\text{ceiling}(N/B)=50$
 - $B-1=17$
 - $\text{\#passes} = 1 + \text{ceiling}(\log_{B-1}(N/B)) = 1 + \text{ceiling}(\log_{17}(50)) = 1 + \text{ceiling}(1.38\text{ish}) = 1 + 2 = 3$
 - How many IO operations?

Example

- We have a large dataset of 900 pages. We are going to use 18 buffer pages
 - How many passes will be required while performing a general external merge sort?
 - $N=900, B=18$
 - $\text{ceiling}(N/B)=50$
 - $B-1=17$
 - $\text{\#passes} = 1 + \text{ceiling}(\log_{B-1}(N/B)) = 1 + \text{ceiling}(\log_{17}(50)) = 1 + \text{ceiling}(1.38\text{ish}) = 1 + 2 = 3$
 - How many IO operations?
 - $\text{\#IO} = 2N * \text{\#passes}$

Example

- We have a large dataset of 900 pages. We are going to use 18 buffer pages
 - How many passes will be required while performing a general external merge sort?
 - $N=900, B=18$
 - $\text{ceiling}(N/B)=50$
 - $B-1=17$
 - $\text{\#passes} = 1 + \text{ceiling}(\log_{B-1}(N/B)) = 1 + \text{ceiling}(\log_{17}(50)) = 1 + \text{ceiling}(1.38\text{ish}) = 1 + 2 = 3$
 - How many IO operations?
 - $\text{\#IO} = 2N * \text{\#passes} = 2 * 900 * 3 = 5400$

Replacement Sort

$$2N \cdot (1 + \lceil \log_{B-1} (\# \text{ runs}) \rceil)$$



Lower this number


make each run longer

- We have lots and lots of runs
 - More runs → More passes → more IO cost
- What if we had longer runs?
 - Fewer passes!
- Replacement sort helps solve this
 - Let's minimize the number of runs in step 0
 - Rest of the sort is the same
 - We still have B pages
 - Set aside 1 for input and 1 for output
 - B-2 buffer pages for current set
 - Continually read input and add to current set (make sure current set is full at all times)
 - Output smallest element in current set that is larger than largest value in output set
 - Write output buffer to disk when full
 - When no such element exists, end the run and start a new run
 - Continue till everything is in a run. Then continue with external merge sort as normal

Applies to initial pass
when creating runs

Step 1


- Have a dataset



Step 2

Suppose $B = 5$ and
each page can hold
2 bars in full.

- Allocate pages for current set, input, and output pages



Input Page




Current Set



Output Page

Step 3

- Read data into current set using input page



Input Page



Current Set



Suppose $B = 5$ and each page can hold 2 bars in full.



Output Page

Step 3

- Read data into current set using input page



Input Page



Current Set


Suppose $B = 5$ and each page can hold 2 bars in full.




Output Page

Step 3

- Read data into current set using input page



Input Page



Current Set




Suppose $B = 5$ and each page can hold 2 bars in full.



Output Page

Step 3

- Read data into current set using input page




Suppose $B = 5$ and each page can hold 2 bars in full.





Step 3

- Read data into current set using input page



Input Page



Current Set



Suppose $B = 5$ and each page can hold 2 bars in full.





Output Page

Step 3

- Read data into current set using input page



Input Page



Current Set

Suppose $B = 5$ and each page can hold 2 bars in full.





Output Page


Step 4

Suppose $B = 5$ and
each page can hold
2 bars in full.

- Take minimum element from current set greater than largest element in current run and put in output page



Input Page



Current Set



Output Page




Current Run

Step 4

Suppose $B = 5$ and
each page can hold
2 bars in full.


- Take minimum element from current set greater than largest element in current run and put in output page





Step 4

Suppose $B = 5$ and
each page can hold
2 bars in full.

- Take minimum element from current set greater than largest element in current run and put in output page



Input Page



Current Set

Refill Current Set!



Output Page




Current Run


Step 4

Suppose $B = 5$ and
each page can hold
2 bars in full.

- Take minimum element from current set greater than largest element in current run and put in output page



Input Page



Current Set



Output Page




Current Run



Step 4

Suppose $B = 5$ and
each page can hold
2 bars in full.

- Take minimum element from current set greater than largest element in current run and put in output page



Input Page




Current Set

Refill Current Set!



Output Page
Write Output Page!




Current Run (on disk)

Step 4

Suppose $B = 5$ and
each page can hold
2 bars in full.



- Take minimum element from current set greater than largest element in current run and put in output page




Step 4

Suppose $B = 5$ and
each page can hold
2 bars in full.

- Take minimum element from current set greater than largest element in current run and put in output page



Input Page



Current Set



Output Page





Current Run (on disk)


Step 4

Suppose $B = 5$ and
each page can hold
2 bars in full.

- Take minimum element from current set greater than largest element in current run and put in output page



Input Page




Current Set

Refill Current Set!



Output Page




Current Run (on disk)


Step 4

Suppose $B = 5$ and
each page can hold
2 bars in full.


- Take minimum element from current set greater than largest element in current run and put in output page




Input Page



Current Set



Output Page



Current Run (on disk)




Step 4

Suppose $B = 5$ and
each page can hold
2 bars in full.

- Take minimum element from current set greater than largest element in current run and put in output page




Input Page
Refill Input Page!



Current Set
Refill Current Set!



Output Page
Write Output Page!





Current Run (on disk)


Step 4

Suppose $B = 5$ and each page can hold 2 bars in full.


- Take minimum element from current set greater than largest element in current run and put in output page




Input Page



Current Set



Output Page




Current Run (on disk)



Step 4

Suppose $B = 5$ and
each page can hold
2 bars in full.

- Take minimum element from current set greater than largest element in current run and put in output page




Input Page




Current Set

Refill Current Set!



Output Page




Current Run (on disk)


Step 4

Suppose $B = 5$ and each page can hold 2 bars in full.


- Take minimum element from current set greater than largest element in current run and put in output page




Input Page



Current Set



Output Page




Current Run (on disk)




Step 4

Suppose $B = 5$ and each page can hold 2 bars in full.

- Take minimum element from current set greater than largest element in current run and put in output page




Input Page



Current Set



Output Page




Current Run (on disk)



Step 4

Suppose $B = 5$ and each page can hold 2 bars in full.


- Take minimum element from current set greater than largest element in current run and put in output page




Input Page



Current Set



Output Page



Current Run (on disk)


Step 4

Suppose $B = 5$ and each page can hold 2 bars in full.

- Take minimum element from current set greater than largest element in current run and put in output page



Input Page




Current Set



Output Page

guaranteed be sorted





Current Run (on disk)



Step 4

Suppose $B = 5$ and each page can hold 2 bars in full.


- Take minimum element from current set greater than largest element in current run and put in output page




Input Page



Current Set



Output Page




Current Run (on disk)

Step 4

Suppose $B = 5$ and each page can hold 2 bars in full.



- Take minimum element from current set greater than largest element in current run and put in output page





Step 5

- Write last of output page to disk and save run

Suppose $B = 5$ and each page can hold 2 bars in full.




Input Page



Current Set



Output Page





Current Run (on disk)

*you have longer run > B
page long*



Step 5

- Create new run and repeat

Suppose $B = 5$ and each page can hold 2 bars in full.




Input Page



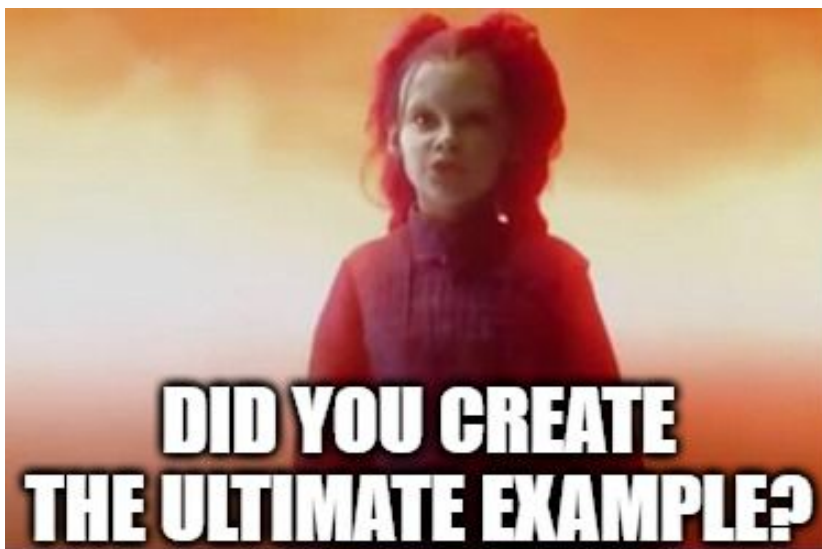
Current Set



Output Page



Current Run (on disk)



Replacement Sort Math

$\rightarrow B$

- Average length of a run is $2 \cdot (B-2)$ pages

$2(B-2)$ - page-long runs

- Proof left as an exercise to the reader

- Fewer but longer runs

- $\text{ceiling}(N / (2 \cdot (B-2)))$ runs

$$1 + \left\lceil \log_{B-1} \left\lceil \frac{N}{2B-4} \right\rceil \right\rceil$$

- Takes $1 + \text{ceiling}(\log_{B-1} \text{ceiling}(N / (2 \cdot (B-2))))$ passes



- Should be a smaller number!
- Not always in all cases though

- Total IO cost is #passes * $2N$

- Smaller because fewer passes!

Example

- Suppose we have 4 buffer pages and the following dataset:
 - 12, 15, 5, 30, 33, 51, 8, 1, 2, 7
- Each page can store 2 elements
- How many runs will we have?



Example

- Suppose we have 4 buffer pages and the following dataset:
 - 12, 15, 5, 30, 33, 51, 8, 1, 2, 7
- Each page can store 2 elements
- How many runs will we have?
- 4 buffer pages -> 1 input, 1 output page -> 2 pages for current set

Example

- Suppose we have 4 buffer pages and the following dataset:
 - 12, 15, 5, 30, 33, 51, 8, 1, 2, 7
- Each page can store 2 elements
- How many runs will we have?
- 4 buffer pages -> 1 input, 1 output page -> 2 pages for current set
- Current set = {12, 15, 5, 30}

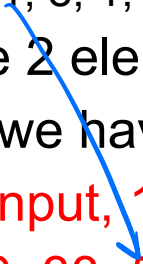
Example

- Suppose we have 4 buffer pages and the following dataset:
 - 12, 15, 5, 30, 33, 51, 8, 1, 2, 7
- Each page can store 2 elements
- How many runs will we have?
- 4 buffer pages -> 1 input, 1 output page -> 2 pages for current set
- Current set = {12, 15, 5, 30}
- Run 1 = {}


Example

- Suppose we have 4 buffer pages and the following dataset:
 - 12, 15, 5, 30, 33, 51, 8, 1, 2, 7
- Each page can store 2 elements
- How many runs will we have?
- 4 buffer pages -> 1 input, 1 output page -> 2 pages for current set
- Current set = {12, 15, 30, 33}
- Run 1 = {5}


Example

- Suppose we have 4 buffer pages and the following dataset:
 - 12, 15, 5, 30, 33, 51, 8, 1, 2, 7
 - Each page can store 2 elements
 - How many runs will we have?
 - 4 buffer pages -> 1 input, 1 output page -> 2 pages for current set
 - Current set = {15, 30, 33, 51}
 - Run 1 = {5, 12}
- 


Example

- Suppose we have 4 buffer pages and the following dataset:
 - 12, 15, 5, 30, 33, 51, 8, 1, 2, 7
 - Each page can store 2 elements
 - How many runs will we have?
 - 4 buffer pages -> 1 input, 1 output page -> 2 pages for current set
 - Current set = {30, 33, 51, 8}
 - Run 1 = {5, 12, 15}
- 

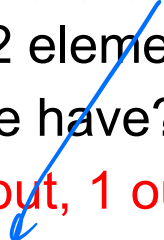
Example

- Suppose we have 4 buffer pages and the following dataset:
 - 12, 15, 5, 30, 33, 51, 8, 1, 2, 7
 - Each page can store 2 elements
 - How many runs will we have?
 - 4 buffer pages -> 1 input, 1 output page -> 2 pages for current set
 - Current set = {33, 51, 8, 1}
 - Run 1 = {5, 12, 15, 30}
- 

Example

- Suppose we have 4 buffer pages and the following dataset:
 - 12, 15, 5, 30, 33, 51, 8, 1, 2, 7
 - Each page can store 2 elements
 - How many runs will we have?
 - 4 buffer pages -> 1 input, 1 output page -> 2 pages for current set
 - Current set = {51, 8, 1, 2}
 - Run 1 = {5, 12, 15, 30, 33}
- 

Example

- Suppose we have 4 buffer pages and the following dataset:
 - 12, 15, 5, 30, 33, 51, 8, 1, 2, 7
 - Each page can store 2 elements
 - How many runs will we have?
 - 4 buffer pages -> 1 input, 1 output page -> 2 pages for current set
 - Current set = {8, 1, 2, 7}
 - Run 1 = {5, 12, 15, 30, 33, 51}
- 

Example

- Suppose we have 4 buffer pages and the following dataset:
 - 12, 15, 5, 30, 33, 51, 8, 1, 2, 7
- Each page can store 2 elements
- How many runs will we have?
- 4 buffer pages -> 1 input, 1 output page -> 2 pages for current set
- Current set = {8, 1, 2, 7}
- Run 1 = {5, 12, 15, 30, 33, 51}
- No more elements in current set greater than largest element in run :(
 - Flush output and start a new run

Example

- Suppose we have 4 buffer pages and the following dataset:
 - 12, 15, 5, 30, 33, 51, 8, 1, 2, 7
- Each page can store 2 elements
- How many runs will we have?
- 4 buffer pages -> 1 input, 1 output page -> 2 pages for current set
- Current set = {8, 1, 2, 7}
- Run 2 = {}
- Run 1 = {5, 12, 15, 30, 33, 51}

Example

- Suppose we have 4 buffer pages and the following dataset:
 - 12, 15, 5, 30, 33, 51, 8, 1, 2, 7
- Each page can store 2 elements
- How many runs will we have?
- 4 buffer pages -> 1 input, 1 output page -> 2 pages for current set
- Current set = {8, 2, 7} - No more data!
- Run 2 = {1}
- Run 1 = {5, 12, 15, 30, 33, 51}

Example

- Suppose we have 4 buffer pages and the following dataset:
 - 12, 15, 5, 30, 33, 51, 8, 1, 2, 7
- Each page can store 2 elements
- How many runs will we have?
- 4 buffer pages -> 1 input, 1 output page -> 2 pages for current set
- Current set = {8, 7} - No more data!
- Run 2 = {1, 2}
- Run 1 = {5, 12, 15, 30, 33, 51}

Example

- Suppose we have 4 buffer pages and the following dataset:
 - 12, 15, 5, 30, 33, 51, 8, 1, 2, 7
- Each page can store 2 elements
- How many runs will we have?
- 4 buffer pages -> 1 input, 1 output page -> 2 pages for current set
- Current set = {8} - No more data!
- Run 2 = {1, 2, 7}
- Run 1 = {5, 12, 15, 30, 33, 51}

Example

- Suppose we have 4 buffer pages and the following dataset:
 - 12, 15, 5, 30, 33, 51, 8, 1, 2, 7
- Each page can store 2 elements
- How many runs will we have?
- 4 buffer pages \rightarrow 1 input, 1 output page \rightarrow 2 pages for current set
- Current set = {} - No more data!
- Run 2 = {1, 2, 7, 8}
- Run 1 = {5, 12, 15, 30, 33, 51}
- Done :)
- 2 Runs

) lead runs with different length
on avg : $2(B-2)$ - length runs.

Blocked I/O

$$1 + \left\lceil \log_{\left\lfloor \frac{B}{b} \right\rfloor - 1} \left\lceil \frac{N}{B} \right\rceil \right\rceil$$

sequential I/O cheaper
than random I/O

- Pages in memory live in segments called Blocks

- Often cheaper to load an entire Block into memory at a time

(contain multiple pages)

- How to handle

- Instead of caring about B buffer pages lets take into account our blocks

- Block consists of b pages

- We instead have $\text{floor}(B/b)$ buffer "blocks"

- Floor since we can't have half a block and we're given an upper bound on the number of pages (and blocks) that can fit into memory at the same time

- #passes = $1 + \text{ceiling}(\log_{\text{floor}(B/b)-1}(\text{ceiling}(N/B)))$

- More passes but less per page costs

- Everything else is the same as usual

$$\left\lfloor \frac{B}{b} \right\rfloor$$

floor: can't read
 $\frac{1}{3}$ of block


$$\left\lfloor \frac{B}{b} \right\rfloor - 1$$

output buffer block


$\Rightarrow \log \text{ base } \downarrow \text{ pass } \uparrow \Rightarrow \text{I/O } \uparrow$

do $\left\lfloor \frac{B}{b} \right\rfloor$ -way merge

non-blocked I/O



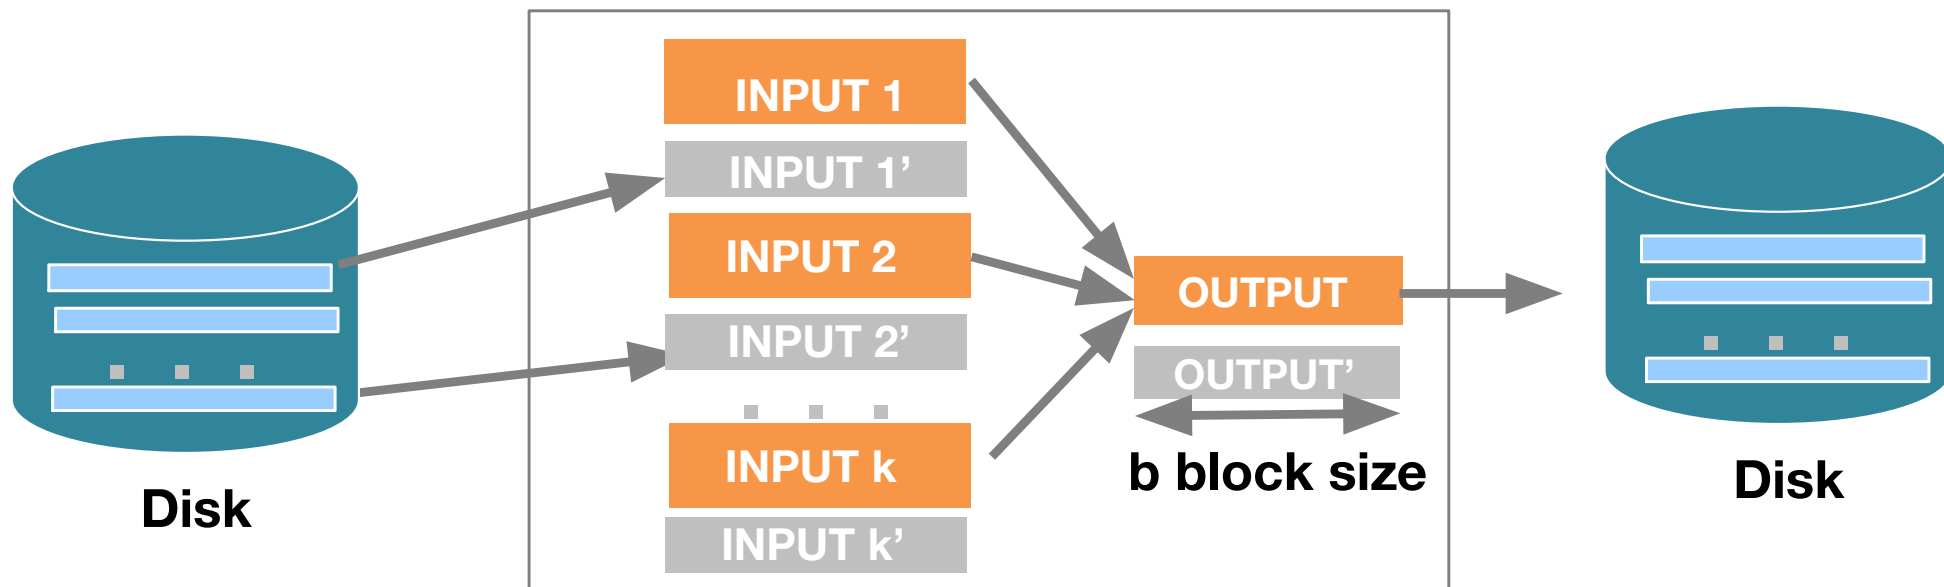
blocked I/O (suppose $b = 2$)



Double Buffering

keep CPU and disk working all the time
disk read in orange / grey
CPU do grey / orange

- I/O processing needs time
- Overlap CPU and I/O processing



$2*(K+1)*b$ main memory buffers, k-way merge

Extra Examples

- We have a 565 GB dataset with page size of 8 KB and 2 GB RAM
 - How many IOs required to sort?
 - Hint: figure out how many buffer pages we can fit in RAM along with how many pages our dataset spans
- Using replacement sort at step 0, how many IOs do we need now?
- How many IOs will we need if we use blocking (but not replacement sort)?
 - Block size of 64KB
- How many IOs will we need if we use blocking and replacement sort at step 0?
 - Block size of 64KB
- Answers on last slide

Extra Example Answers

- We have a 565 GB dataset with page size of 8 KB and 2 GB RAM. How many IOs required to sort?
 - Number of pages in dataset (N) = $565\text{GB} / (8\text{KB} / \text{page}) = 74,055,680$ data pages
 - Number of buffer pages (B) = $2\text{GB} / (8\text{KB} / \text{page}) = 262,144$ buffer pages
 - # passes = $1 + \text{ceiling}(\log_{B-1}(\text{ceiling}(N/B))) = 2$
 - # IOs = # passes * $2N = 296,222,720$ IOs
- Using replacement sort at step 0, how many IOs do we need now?
 - # passes = $1 + \text{ceiling}(\log_{B-1}(\text{ceiling}(N/(2*(B-2))))) = 1 + \text{ceiling}(\log_{B-1}(142)) = 2$
 - # IOs = # passes * $2N = 296,222,720$ IOs
 - The same??? (Take a look at the log to see why) *fewer passes not guaranteed*
 - Think about which number would I have to decrease (N or B) to have replacement sort result in fewer passes than the previous question

Extra Example Answers

- How many IOs will we need if we use blocking (but not replacement sort)?
 - Block size of 64KB
 - # pages in a block (b) = $(64 \text{ KB/block}) / (8 \text{ KB/page}) = 8 \text{ pages/block}$ $b=8$
 - # passes = $1 + \text{ceiling}(\log_{\text{floor}(B/b)-1}(\text{ceiling}(N/B))) = 2$
 - Again ???
 - # IOs = # passes * $2N = 296,222,720$ IOs *total I/O is same but each page is cheaper*
- How many IOs will we need if we use blocking and replacement sort at step 0?
 - Block size of 64KB
 - # passes = $1 + \text{ceiling}(\log_{\text{floor}(B/b)-1}(\text{ceiling}(N/(2*(B-2))))) = 1 + \text{ceiling}(\log_{\text{floor}(B/b)-1}(142)) = 2$
 - Again ???
 - # IOs = # passes * $2N = 296,222,720$ IOs
- We can sort a CRAZY amount of data with just 2GB of RAM in just 2 passes
 - Yay modern computing :)
 - Try reducing the RAM to 128MB and see how many passes you have to make with each attempt *B is large here \Rightarrow pass # doesn't change much*
 - *B smaller now \Rightarrow # pass may change*

block I/O + replacement sort $1 + \left\lceil \log_{\lfloor \frac{B}{b} \rfloor - 1} \left\lceil \frac{N}{2(B-2)} \right\rceil \right\rceil$

Get started on HW5!