

# Discussion 11

Transactions  
EECS 484

# Logistics


- HW 6 Due Dec. 1st at 11:55 pm
- Project 4 Due Dec. 8th at 11:55 pm
- Final Exam on Dec. 13th, Tuesday 7-9PM
  - Only multiple choice questions
  - Covers second half of course (Indexing and beyond)
  - Practice exams posted on Canvas
  - Final Review Session - Dec. 7th, lecture time
  - Exam time and location for SSD students sent via email
    - If you applied for SSD accommodation and didn't receive any email about it, please email us at [eeecs484staff@umich.edu](mailto:eeecs484staff@umich.edu)
  - More logistics posted later on Piazza
- Today
  - Transactions

# Transaction Management

# Transactions

⇒ all / none






- A transaction is an atomic unit of work
  - Either everything in a transaction happens or it doesn't
  - Multiple actions in a transaction
- Transactions can interleave actions
  - Can do action 2 of transaction 1 at the same time as action 4 of transaction 2
  - Need to ensure that we avoid any inconsistencies
    - Same result as doing transactions serially (consecutively)



Good User's transaction:

G1: Give the first 4  
students an A+


G2: Give the last 4  
students an A+



Evil User's transaction:

E1: Give the first 4  
students an F










E2: Give the last 4  
students an F



Good User's transaction:

G1: Give the first 4  
students an A+

G2: Give the last 4  
students an A+




Evil User's transaction:

E1: Give the first 4  
students an F

E2: Give the last 4  
students an F










Interleaving: G1



Good User's transaction:

G1: Give the first 4 students an A+

G2: Give the last 4 students an A+




Evil User's transaction:

E1: Give the first 4 students an F

E2: Give the last 4 students an F










Interleaving: G1, E1



Good User's transaction:

G1: Give the first 4 students an A+

G2: Give the last 4 students an A+




Evil User's transaction:

E1: Give the first 4 students an F

E2: Give the last 4 students an F

Interleaving: G1, E1, E2














Good User's transaction:

G1: Give the first 4 students an A+

G2: Give the last 4 students an A+



Evil User's transaction:

E1: Give the first 4 students an F

E2: Give the last 4 students an F

Interleaving: G1, E1, E2, G2


# ACID

- Properties we need to enforce to ensure the database is valid
- Atomicity
  - All actions in the transaction happen or none of them happen
- Consistency
  - If we start with a consistent database and perform a consistent transaction we have a consistent database at the end
- Isolation
  - Each transaction appears to occur serially
- Durability
  - If a transaction occurs, its effects persist

# Scheduling

- Schedules are a list of ordered actions across transactions
  - Can interleave actions from various transactions
  - Serial schedule = no interleaving of transactions
  - Serializable schedule = result matches what some serial schedule would have produced (all reads and final states) *result same + intermediate same*
- Different serial schedules for the same transactions can have different results
  - All are assumed to be okay

T1	T2
Read(A)	
	Read(B)
	Write(B)
Write(A)	
	Commit
Commit	



# Conflicts

*Don't care about commit/abort/running*

- Conflicts happen when we try to access a certain resource multiple times in non-compatible ways
  - All happen with writes - updating data
  - **Write-Read (WR)** - could indicate a dirty data
  - **Read-Write (RW)** - could indicate an unrepeatable read
  - **Write-Write (WW)** - could indicate overwriting of uncommitted data
- Conflicts **do not** necessarily mean anomalies
  - **WR, RW, WW conflicts only indicate the order (precedence) that these operations need to be executed**
  - Only if we have an inconsistent database after the schedule do we have an anomaly

*conflict  $\neq$  anomaly*

# Conflict Serializability

(don't consider abort for precedence graph) But conflict still occurs even aborted  
Conflict serializable  $\Leftrightarrow$  no cycle

- A schedule is conflict serializable if and only if its precedence graph is acyclic
- Precedence graph is a series of connected nodes for committed transactions
  - A node for each committed transaction
  - Arc from  $T_x$  to  $T_y$  if some action in  $T_x$  precedes and conflicts with some action in  $T_y$
- All conflict serializable schedules are serializable

- Serializable: all reads and the final state is what some complete serial schedule of committed transactions would have produced.
- Not the other way around!
- Testing conflict serializability is much easier than testing serializability

(graph)

Conflict serializability  $\rightarrow$  serializability

(try every possibility)

serializability  $\Rightarrow$  Conflict serializability



# Conflicts and Conflict Serializability Example


T1	T2
Write(B)	
Read(B)	
	Write(B)
	Write(A)
	Commit
Write(B)	
Commit	

Diagram illustrating a schedule with conflicts between transactions T1 and T2. Conflicting operations are circled in red. Arrows indicate the sequence of operations and the type of conflict (WW, RW, or WW).

- T1: Write(B) (circled), Read(B) (circled), Write(B) (circled), Commit
- T2: Write(B) (circled), Write(A), Commit

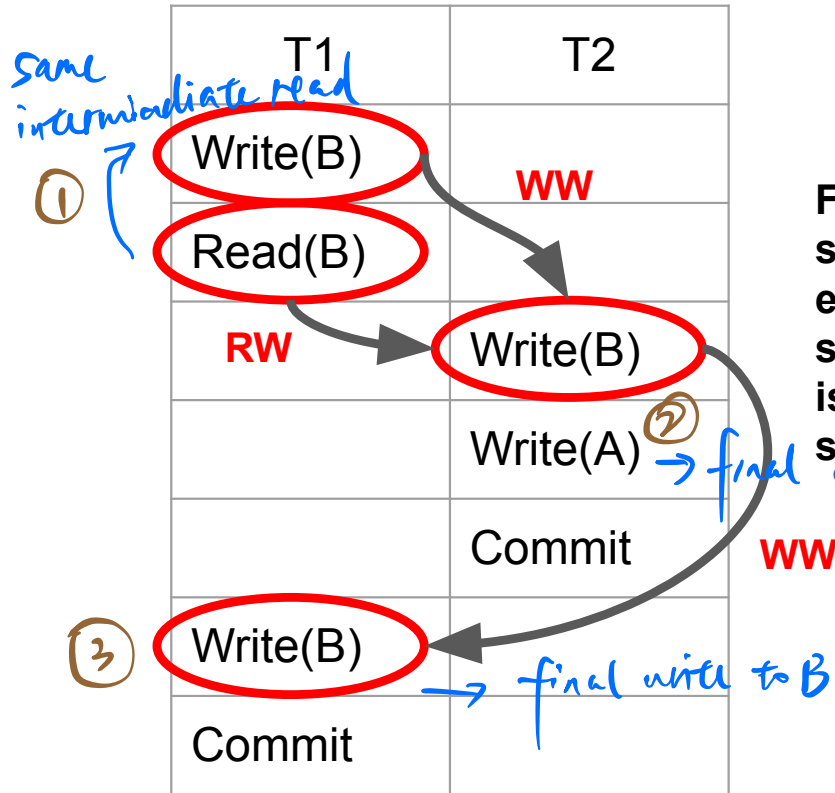
Conflicts identified:

- WW (Write-Write) conflict between T1's Write(B) and T2's Write(B).
- RW (Read-Write) conflict between T1's Read(B) and T2's Write(B).
- WW (Write-Write) conflict between T1's Write(B) and T2's Write(B).



A cycle in the precedence graph indicates this schedule is **not conflict serializable**

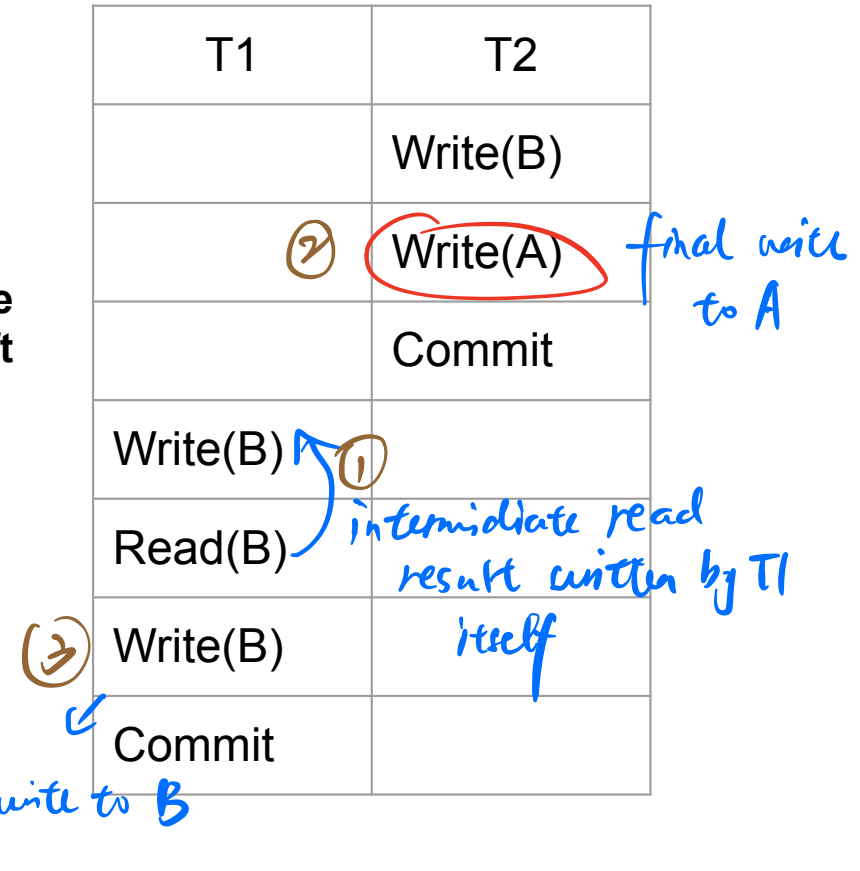
# Not Conflict Serializable BUT Serializable



Not conflict serializable

*Same*

From the previous slide, we established that the schedule on the left is not conflict serializable



but serializable

# Not Conflict Serializable BUT Serializable

T1	T2
Write(B)	
Read(B)	
	Write(B)
	Write(A)
	Commit
Write(B)	
Commit	

Not conflict serializable

**All Reads() still read the same value (i.e. the B value written by T1)**

**The same transactions still make the same final Write() to A and B**


T1	T2
	Write(B)
	Write(A)
	Commit
Write(B)	
Read(B)	
Write(B)	
Commit	

but serializable



# Recoverable Schedules and ACA

ACA  $\rightarrow$  recoverable



- Recoverable schedules


- Any transaction, T1, that reads a change from a different transaction, T2, must **commit after T2**
- Not necessarily serializable or vice versa

T1 read from T2, commit after


- Avoid Cascading Aborts (ACA)

- Transactions **only read changes by committed transactions**
  - We don't have to abort multiple transactions if one is aborted
- ACA implies recoverable but not vice versa


not recoverable



already committed  
no chance for  
rollback



recoverable

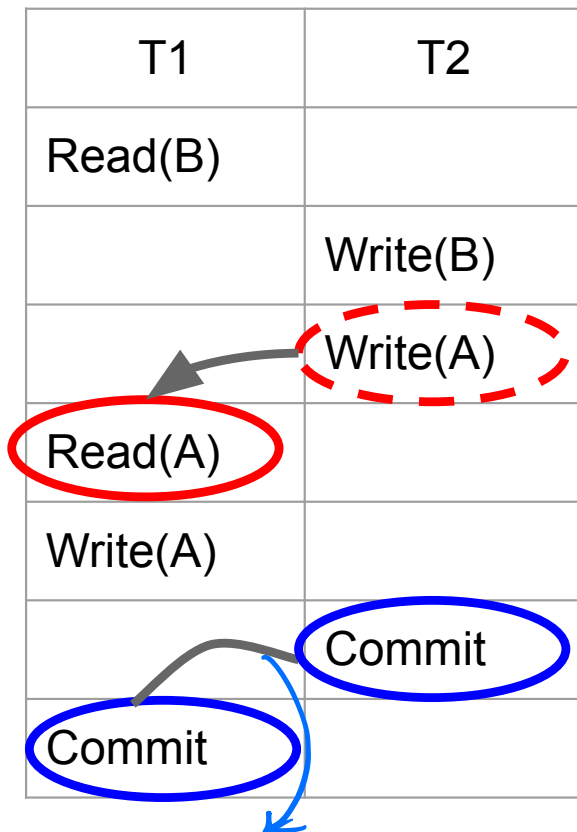


recoverable  
but not ACA

(cascading abort)

still have chance to abort after T2 abort

# Recoverable and ACA Schedule Examples



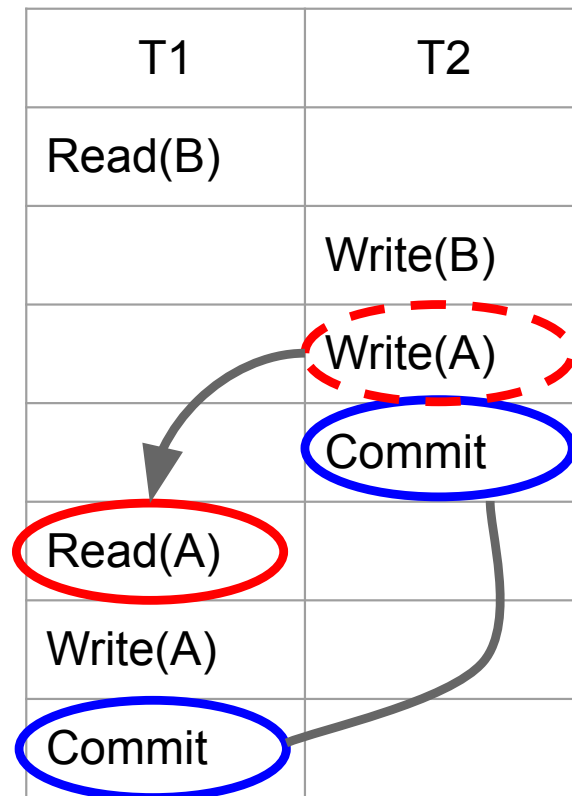
Recoverable but not ACA

**Left:** T1 read uncommitted change made by T2.

- If T2 aborts, T1 has to abort as well. Hence not ACA.

- T1 only commits after T2 commits. Hence it's recoverable.

**Right:** T1 only reads committed changes. Hence it avoids cascading aborts (ACA)



ACA (and thus also recoverable)

# Two Phase Locking (2PL)

- Locking allows system to ensure one transaction occurs before another
  - Use locks, where a transaction must get a lock before proceeding
    - If not available, wait ✱
  - Two types of locks
    - Shared (read) locks: multiple transactions can hold same lock at the same time
      - Can only read resource
    - Exclusive (write) locks: one transaction can hold lock at a time
      - Can read or write resource
      - Prevents other transactions from reading resource while we are writing


# Two Phase Locking

- 2PL



- If a transaction releases any lock, it can no longer acquire any new locks
- Guarantees conflict-serializability (and thus serializability)
- May sometimes cause deadlocks

- Strict 2PL (S2PL)

- 2PL + Hold all locks acquired until the end of the transaction
- Guarantees conflict serializability (and thus serializability)
- Guarantees ACA (and thus recoverability)



others will action after you  
COMMIT !!



# Example Problems

- List the conflicts between T1 and T2
- Is the schedule conflict serializable?
- Is the schedule serializable?
- Is the schedule recoverable?
- Is the schedule ACA?
- Is the schedule 2PL?
- Is the schedule S2PL?

T1	T2
Read(A)	
Write(A)	
	Read(A)
	Write(A)
Read(B)	
Write(B)	
Commit	
	Read(B)
	Write(B)
	Commit

# Example Problems

- **List the conflicts between T1 and T2**

RW conflict from T1 to T2 on A

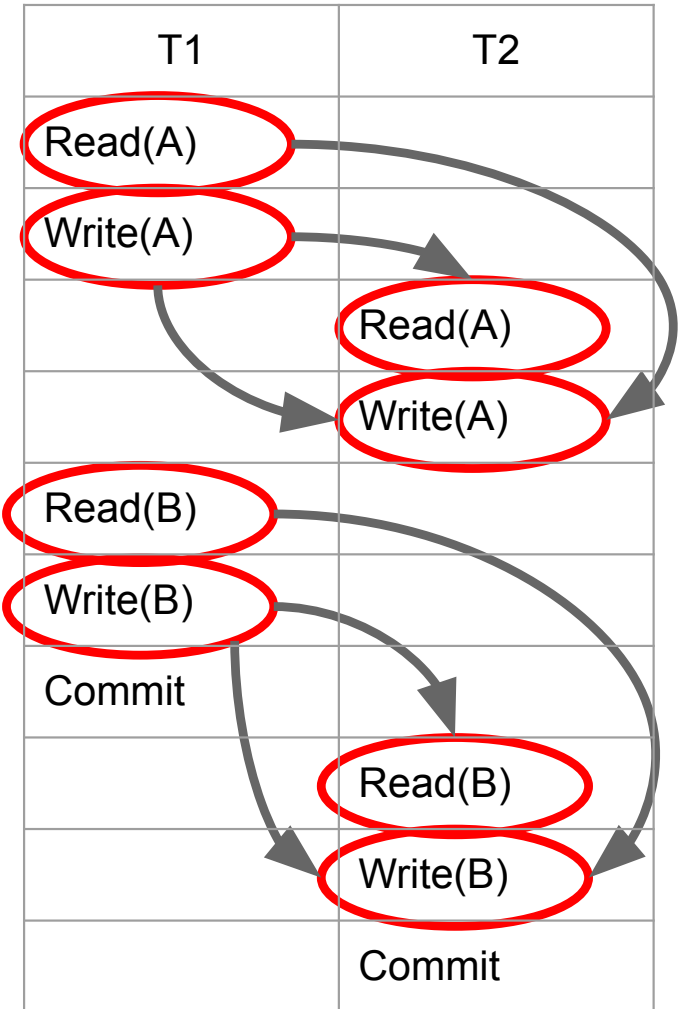
WR conflict from T1 to T2 on A

WW conflict from T1 to T2 on A

RW conflict from T1 to T2 on B

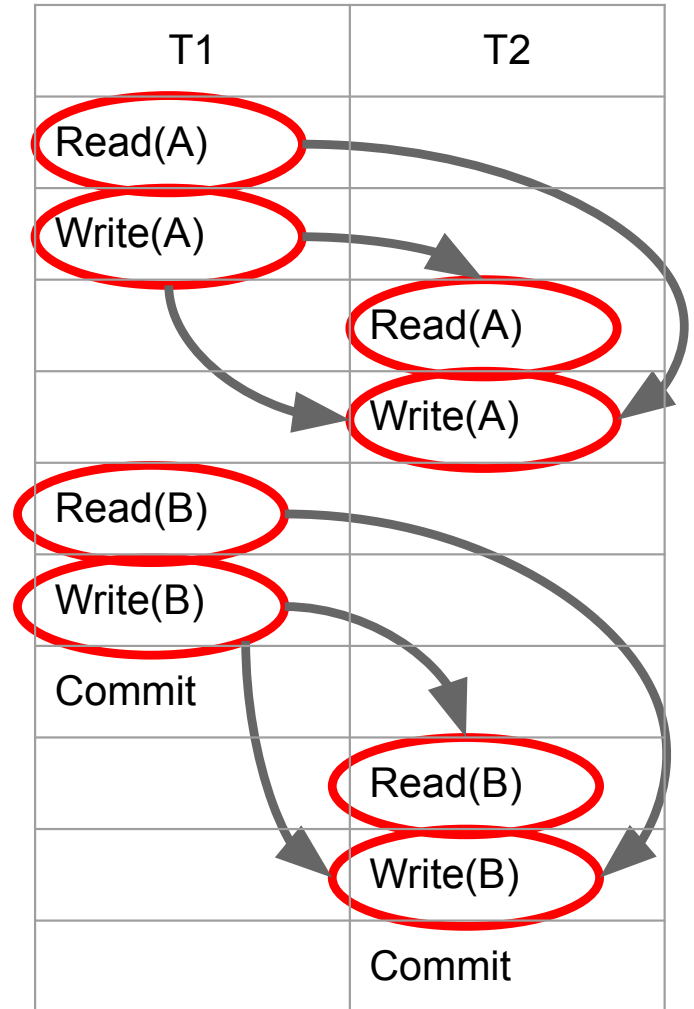
WR conflict from T1 to T2 on B

WW conflict from T1 to T2 on B



# Example Problems

- Is the schedule conflict serializable?
- Yes *all T1 → T2 no cycle*
- Draw the precedence graph
  - Or examine the dependencies
  - No dependencies from T2 to T1
    - For sure acyclic then



# Example Problems

- **Is the schedule serializable?**
- Yes. We get for free since conflict serializable :)

T1	T2
Read(A)	
Write(A)	
	Read(A)
	Write(A)
Read(B)	
Write(B)	
Commit	
	Read(B)
	Write(B)
	Commit



# Example Problems

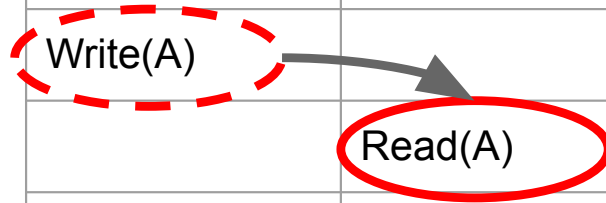
- **Is the schedule recoverable?**
- T2 reads output of T1 (A and B writes)
- T1 commits before T2
- Recoverable

T1	T2
Read(A)	
Write(A)	
	Read(A)
	Write(A)
Read(B)	
Write(B)	
Commit	Read(B)
	Write(B)
	Commit

# Example Problems

- **Is the schedule ACA?**
- T2 reads output of T1 that is not committed
- Not ACA: if T1 aborts, T2 would have to abort as well

T1	T2
Read(A)	
Write(A)	
	Read(A)
	Write(A)
Read(B)	
Write(B)	
Commit	
	Read(B)
	Write(B)
	Commit



# Example Problems

- Is the schedule 2PL?

- Assume we acquire locks at the necessary step
- T1 needs shared lock on A in step 1 - okay
- T1 acquires exclusive lock on A in step 2 - okay
- T2 needs shared lock on A - :(
  - T1 holds exclusive lock on A
  - If T1 lets go, then T1 cannot acquire another lock
  - T1 needs shared and exclusive lock on B later

- Not 2PL

- If we acquire all necessary locks at the beginning and when available, then this can be 2PL

No need for exam

at  $t=0$  T1 grabs exclusive locks on both A & B

T1	T2
Read(A)	
Write(A)	
Read(B)	Read(A)
Write(B)	Write(A)
Commit	
	Read(B)
	Write(B)
	Commit

$t=3$   
 R(A) lock  
 ↓ upgrade  
 W(A) lock  
 T1 releases locks  
 so T2 can grab R(A) lock  
 grab R(B) lock  
 ⇒ Not allowed by 2PL

# Example Problems

- **Is the schedule S2PL?**
  - No since it's not 2PL

T1	T2
Read(A)	
Write(A)	
	Read(A)
	Write(A)
Read(B)	
Write(B)	
Commit	
	Read(B)
	Write(B)
	Commit

# Extra Problems


- List the conflicts between T1 and T2
- Is the schedule conflict serializable?
- Is the schedule serializable?
- Is the schedule recoverable?
- Is the schedule ACA?
- Is the schedule 2PL?
- Is the schedule S2PL?
- Solution on final slide!

T1	T2
Read(A)	
	Read(B)
	Write(B)
	Read(C)
	Write(C)
	Commit
Read(B)	
Write(B)	
Write(A)	
Commit	

# Extra Problems

- List the conflicts between T1 and T2
  - WW, RW, WR From T2->T1 on B
- Is the schedule conflict serializable?
  - Yes, no dependencies from T1->T2
- Is the schedule serializable?
  - Yes, since conflict serializable
- Is the schedule recoverable?
  - Yes, since T1 reads outputs of T2 and T2 commits first
- Is the schedule ACA?
  - Yes, since T1 reads B only after T2 commits

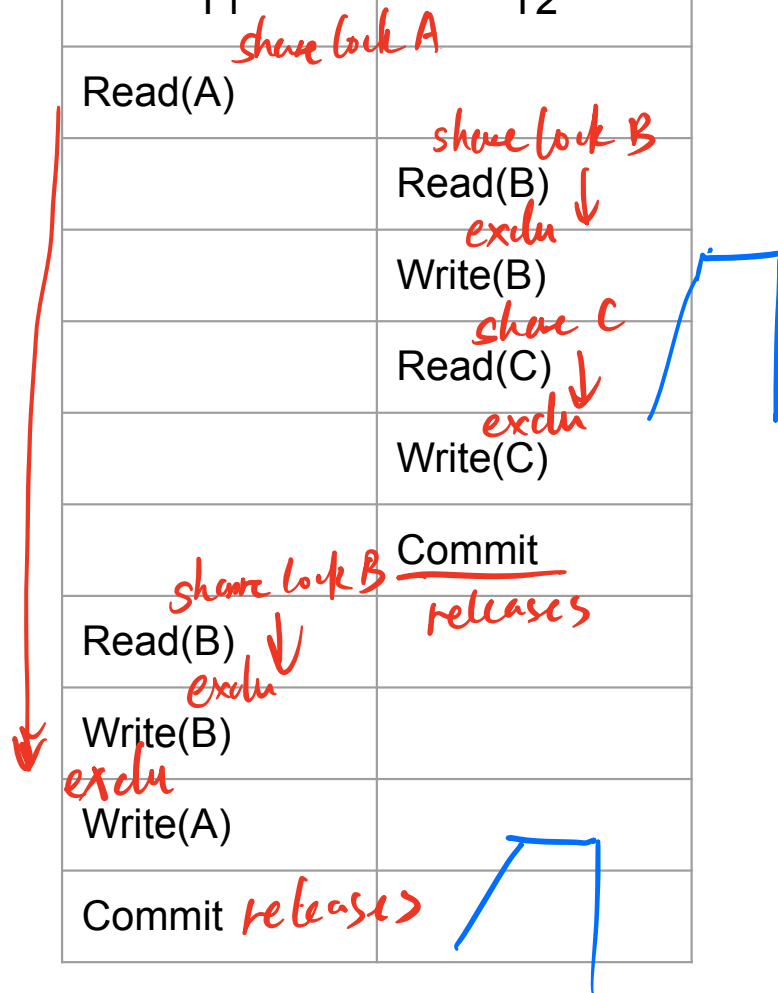
T1	T2
Read(A)	
	Read(B)
	Write(B)
	Read(C)
	Write(C)
	Commit
Read(B)	
Write(B)	
Write(A)	
Commit	



# Extra Problems Cont.

- Is the schedule 2PL? - yes
  - T1 acquires shared lock on A in step 1
  - T2 acquires shared lock on B in step 2
  - T2 acquires exclusive lock on B in step 3
  - T2 acquires shared lock on C in step 4
  - T2 acquires exclusive lock on C in step 5
  - T2 commits and lets go of all locks in step 6
  - T1 acquires shared lock on B in step 7
  - T1 acquires exclusive lock on B in step 8
  - T1 acquires exclusive lock on A in step 9
  - T1 commits and lets go of all locks in step 10
- Is the schedule S2PL?
  - Yes, T1 and T2 only release exclusive locks on commit

T1	T2
Read(A)	
	Read(B)
	Write(B)
	Read(C)
	Write(C)
	<u>Commit</u>
Read(B)	
Write(B)	
Write(A)	
Commit	



**Get started on  
P4!**