

The Relational Model: Database Definition and Integrity Constraints

Chapter 3

Jennifer Widom Lectures

- Introduction and Relational Databases
- Relational Algebra
- SQL
- Constraints and Triggers

<https://www.edx.org/course/databases-5-sql>

Relational and Other Data Models

- **DBMS using the relational DM**

- IBM DB2
- MS SQL Server
- Informix
- Oracle
- Sybase
- Microsoft Access
- Tandem
- Teradata
- SQLite
- MySQL
- PostgreSQL...

- Other data models

- ✧ Hierarchical
 - IBM IMS
- ✧ Network
 - IDMS, IDS
- ✧ Object-oriented
 - ObjectStore
- ✧ Object-relational
 - Oracle
- ✧ ...



Relational (Data) Model

- The most widely-used model today
- **Data model** = a collection of concepts for describing data
 - A collection of **relations**
 - **Relation** = set of records – think of it as a table with rows and columns
 - attributes*
 - tuples*

Students

sid	name	login	age
13	Lisa	lsimp	40
41	Bart	bart	20

Courses

cid	cname	Cr.
E-484	EECS484	4
E-584	EECS584	3

Enrolled

sid	cid	Grade
41	E-484	A-
13	E-584	A+



Relational (Data) Model

- **Schema** = a description of data in terms of a data model
 - Every relation has a schema
 - Specifies the **name** of the **relation**, the **name** and **type** of the **columns** (or fields or attributes)
 - Each row also called a **tuple** or a record

Students(sid:string, name:string, login:string, age:integer)
Courses(cid:string, cname:string, credits:integer)
Enrolled(sid:string, cid:string, grade:string)

Students			
sid	name	login	age
13	Lisa	lsimp	40
41	Bart	bart	20

Courses		
cid	cname	Cr.
E-484	EECS484	4
E-584	EECS584	3

Enrolled		
sid	cid	Grade
41	E-484	A-
13	E-584	A+

not a row in DB \Rightarrow it's schema.




Relational (Data) Model

- **Schema** = a description of data in terms of a data model
- **Instance** = a table, with rows (aka tuples, records), and columns (aka fields, attributes) that match the schema
 - # of rows: cardinality
 - # of columns: degree or arity

Students

sid	name	login	age
13	Lisa	lsimp	40
41	Bart	bart	20

New Scenario: Olympic Games



- Some history:
 - Inspired by the ancient Olympic Games, which were held in Olympia, **Greece** (8th century BC).



Example:

Instance of Athlete Relation

AID	Name	Country	Sport
1	Mary Lou Retton	USA	Gymnastics
2	Jackie Joyner-Kersey	USA	Track
3	Guo Jingjing	China	Diving

What is the schema?

(aid: *integer*, name: *string*,
country: *string*, sport: *string*)

Example: Instance of Athlete Relation

AID	Name	Country	Sport
1	Mary Lou Retton	USA	Gymnastics
2	Jackie Joyner-Kersee	USA	Track
3	Guo Jingjing	China	Diving

Cardinality & Degree?

3 4

- (A) Cardinality: 3, Degree: 3
- (B) Cardinality: 3, Degree: 4
- (C) Cardinality: 4, Degree: 3

Relational Query Languages

- Supports simple, powerful **querying** of data
- Queries written **declaratively** *⇒ declare what would like to see*
 - In contrast to **procedural** methods
(C++)
- DBMS is responsible for **efficient evaluation**
 - System can optimize for efficient query execution, and still ensure that the answer does not change
- SQL is the standard database query language



Structured Query Language (SQL)

- Create a Table Create
- Add new records Insert
- Retrieve records Select
- Update records Update
- Delete records Delete

- Create a View Create
- Update a View Update

external schema.



Structured Query Language (SQL)



- Create a Table `Create`
 - Integrity Constraints
 - Enforcing Constraints
- Add new records `Insert`
- Retrieve records `Select`
- Update records `Update`
- Delete records `Delete`
- Create a View `Create`
- Update a View `Update`

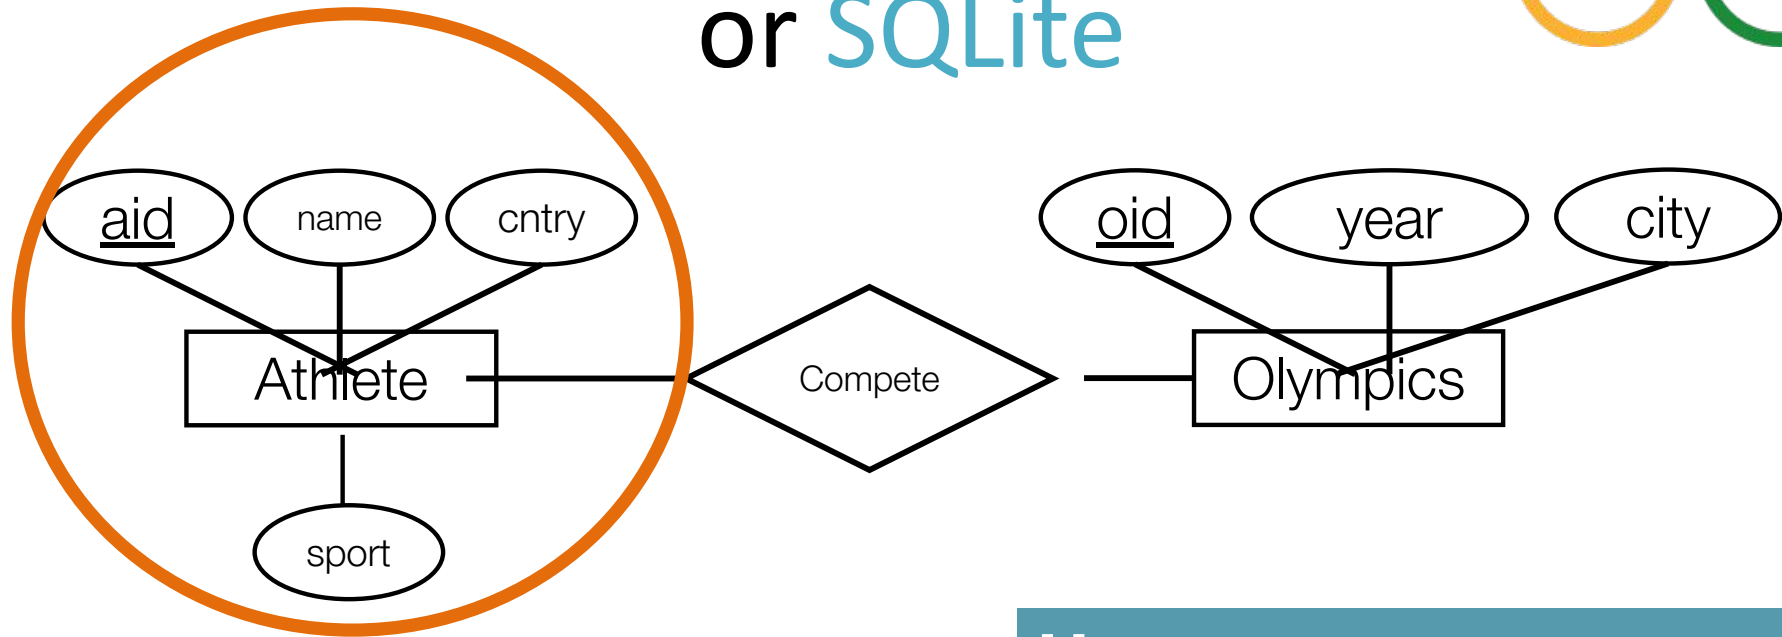



Create a Table (Relation)

. field = attribute = column .

```
CREATE TABLE table_name (  
    field1 TYPE,  
    field2 TYPE,  
    ... ..  
) ;
```

Try it out on paper
or SQLite




```
CREATE TABLE table_name
(
    field1 TYPE,
    field2 TYPE,
    ... ..
);
```

HINTS

- Examples of types:
- char(20), integer, real, (big). text, blob *binary large object*
- To create a DB named olympics:
>> sqlite3 olympics.db

Download SQLite:


<https://www.sqlite.org/download.html>




Creating Relations in SQL

- Create the Athlete relation
 - Domain constraint (type) enforced when tuples added or modified

```
CREATE TABLE Athlete
(aid INTEGER,
name CHAR(30),
country CHAR(20),
sport CHAR(20));
```







Creating Relations in SQL

- Create the Athlete relation
 - Domain constraint (type) enforced when tuples added or modified
- Create the Olympics relation

```
CREATE TABLE Athlete
(aid INTEGER,
name CHAR(30),
country CHAR(20),
sport CHAR(20));
```

```
CREATE TABLE Olympics
(oid INTEGER,
year INTEGER,
city CHAR(20));
```





Creating Relations in SQL

- Create the Athlete relation
 - Domain constraint (type) enforced when tuples added or modified
- Create the Olympics relation
- Create the Compete relation

```
CREATE TABLE Athlete  
(aid INTEGER,  
name CHAR(30),  
country CHAR(20),  
sport CHAR(20));
```

```
CREATE TABLE Olympics  
(oid INTEGER,  
year INTEGER,  
city CHAR(20));
```


```
CREATE TABLE Compete  
(aid INTEGER,  
oid INTEGER);
```

Structured Query Language (SQL)



- Create a Table `Create`
 - Integrity Constraints
 - Enforcing Constraints
- Add new records `Insert`
- Retrieve records `Select`
- Update records `Update`
- Delete records `Delete`
- Create a View `Create`
- Update a View `Update`





Creating Relations: Constraints

- How to specify certain attributes as **keys**?
 - e.g., athlete ID (aid) or olympics ID (oid)
 - DBMS must prevent duplicate keys, e.g. two athletes with the same ID in the database
- How to say that the Athlete ID and Olympic ID values in Compete relation must have **valid references**?

Integrity Constraints (ICs)

- IC: condition that must be true for **every** instance of the database; e.g., domain constraints
 - ICs are specified when schema is defined
 - ICs are checked when relations are modified
- A **legal** instance of a relation satisfies ***all*** specified ICs
 - DBMS must not admit illegal instances




Integrity Constraint:

Primary and Candidate Keys

- A **key** for a relation R
 - = **minimal** set of attributes A_1, \dots, A_n such that:
 - no two tuples in **(any instance of)** R can have the same values for A_1, \dots, A_n
- A set of attributes that satisfies the above condition, **without the minimal requirement**, is called a **superkey**.
 - \Rightarrow Every key is a superkey.
- A relation can have more than one key:
 - One is designated as **primary key**.
 - Others are called **candidate keys**.

Examples: {aid}, {ssn},

- {aid} is a key in the Athlete relation
- {ssn} is a key for Citizen relation




PRIMARY KEY Constraint


A couple of ways to specify a Primary Key constraint:

```
CREATE TABLE Athlete  
(aid INTEGER PRIMARY KEY,  
  name CHAR(30),  
  country CHAR(20),  
  sport CHAR(20));
```

```
CREATE TABLE Athlete  
(aid INTEGER,  
  name CHAR(30),  
  country CHAR(20),  
  sport CHAR(20),  
  PRIMARY KEY(aid));
```

⇒ for more than 1 attribute key






NOT NULL Constraint

Disallow null values for a field

```
CREATE TABLE Athlete
(aid INTEGER PRIMARY KEY,
name CHAR(30) NOT NULL,
country CHAR(20),
sport CHAR(20));
```



- NULL value = the value is unknown or inapplicable
- Example:
 - country and sport can be NULL (= not known or unspecified)
 - But, name must be specified.




Primary Keys Properties

- Can **never** be null (DBMS enforces this)
 - NO parts of a composite primary key can be NULL.
- Need not be an integer ID
 - though they often are for efficient search

An implementation detail

- IDs used as primary keys do not necessarily auto-increment in databases.
 - Additional features of SQL must be used to make them auto-increment. You will see that in the projects.



Candidate Keys

- Candidate keys specified using **UNIQUE**
- One of the candidate keys is chosen as the *primary key*.

```
CREATE TABLE Athlete  
  (aid INTEGER,  
   name CHAR(30) NOT NULL,  
   country CHAR(20),  
   sport CHAR(20),  
   UNIQUE (name, country),  
   PRIMARY KEY (aid));
```

*(candidate key
as a constraint !!!
more constraints ⇒ less likely
to have errors*

WARNING: If used carelessly, ICs can prevent storing instances that arise in practice!

too rigorous.

Foreign Keys in SQL

- Only people listed in Athletes relation should be allowed to compete

```
CREATE TABLE Compete
  (aid INTEGER,   oid INTEGER,
   PRIMARY KEY   (aid, oid),
   FOREIGN KEY   (aid) REFERENCES Athlete);
```

restrict aid domain.

- ... and only in games stored in the Olympics relation

```
CREATE TABLE Compete
  (aid INTEGER,   oid INTEGER,
   PRIMARY KEY   (aid, oid),
   FOREIGN KEY   (aid) REFERENCES Athlete,
   FOREIGN KEY   (oid) REFERENCES Olympics);
```

user implemented ptr.

Foreign Keys: Definition and Rules

- **Foreign key** = set of fields in one relation that is used to refer to a tuple in another relation.
- **Must refer to primary key of the second relation**
 - Like a 'logical pointer'

- **Example:**

```
CREATE TABLE Compete
(aid INTEGER,   oid INTEGER,
 PRIMARY KEY   (aid, oid),
 FOREIGN KEY (aid) REFERENCES Athlete,
 FOREIGN KEY (oid) REFERENCES Olympics);
```

```
CREATE TABLE Athlete
(aid INTEGER PRIMARY KEY,
 name CHAR(30),
 country CHAR(20),
 sport CHAR(20));
```



If all foreign key constraints are enforced, **referential integrity** (no dangling references) is achieved.

like no dangling references

Structured Query Language (SQL)



- Create a Table
 - Integrity Constraints
 - Enforcing Constraints
- Add new records
- Retrieve records
- Update records
- Delete records

- Create a View
- Update a View

Create

Insert

Select

Update

Delete


Create

Update



Enforcing ICs


- Whenever we modify the database
 - the DBMS must check for violations of ICs
- Enforcing Domain, Primary Key, Unique ICs is straightforward ^① ^② ^③ .
 - Reject offending UPDATE / INSERT command



Enforcing Referential Integrity

- If a Compete tuple is **inserted** with no corresponding Athlete aid:
 - Insert operation is REJECTED!

```
CREATE TABLE Compete
  (aid INTEGER,   oid INTEGER,
   PRIMARY KEY   (aid, oid),
   FOREIGN KEY   (aid) REFERENCES Athlete,
   FOREIGN KEY   (oid) REFERENCES Olympics);
```



Enforcing Referential Integrity

- If a Compete tuple is **inserted** with no corresponding Athlete aid:
 - Insert operation is REJECTED!
- What if an Athlete tuple is **deleted**? Possible actions:
 - ✓ – **Disallow deletion** if a Compete tuple refers to athlete
 - **Delete all** Compete tuples that refer to deleted athlete
 - **Set to default or null** value for all references to the deleted athlete
- Similar choices on update of primary key of Athlete

default: IC check happens at the end: set of transactions all completed.

Referential Integrity in SQL

- SQL supports all four options on deletes and updates.
- Default is **NO ACTION**: action is rolled back at commit;

check right away

Similar to **RESTRICT**: action is disallowed at delete/update time.

- **CASCADE**: also delete all tuples that refer to deleted tuple
- **SET NULL / SET DEFAULT**: sets foreign key value of referencing tuple

transaction: Valid
take money out of account A

invalid

add money to account B

Valid

```
CREATE TABLE Compete
  (aid INTEGER,   oid INTEGER,
   PRIMARY KEY   (aid, oid),
   FOREIGN KEY   (aid)
REFERENCES Athlete
ON DELETE CASCADE
ON UPDATE SET NULL)
```


What happens to the Compete relation if we add a new athlete (with a new ID) to the Athlete table? Nothing

Primary Keys and ICs

- The primary key (or any of its parts, if it is composite) cannot get NULL value.
 - True in most DBMS. In this class, we will stick with this rule
 - despite the fact that there are exceptions
- Although you can choose how to handle UPDATE/DELETE actions when there are references (via foreign keys), some options may be in conflict with the Primary Key constraints
 - The system will not allow changes that violate the Primary Key constraints.

Where do ICs Come From?

- Based on real-world enterprise being modeled
- An IC is a statement about **all** possible instances!
- We can **check** a database instance to see if an IC is **violated**, but we can **NEVER** infer that an IC is true by looking at an instance.
- Key and foreign key ICs are the most common
- Also table constraints and assertions *cross-table constraint*
general constraints



Destroying & Altering Relations

- To destroy the relation Olympics.
 - Schema information and tuples are deleted

```
DROP TABLE Olympics
```

- To alter the Athlete schema by adding a new column

```
ALTER TABLE Athlete  
ADD COLUMN age: INTEGER
```

- What do we put in the new field?
 - A **null** value: 'unknown' or 'inapplicable'

Relational Model: Summary

- A tabular representation of data
- Simple and intuitive
- The most widely used database model by far
- Integrity constraints can be specified by the DBA, based on application semantics. DBMS checks for violations.
 - Two important ICs: primary and foreign keys
 - We always have domain constraints
 - e.g. INTEGER fields must always contain integer values
- Views can be used for external schemas, and provide logical data independence

Students


sid	name	login	age
13	Lisa	lsimp	40
41	Bart	bart	20

Courses

cid	cname	Cr.
E-484	EECS484	4
E-584	EECS584	3

Enrolled

sid	cid	Grade
41	E-484	A-
13	E-584	A+



Terminology Parade

- **Database**: A set of relations or tables in the database:
 - **Relation**: Defined by:
 - **Schema**: Describes the columns and constraints
 - Relation name
 - Name and **domain** (i.e., type) for each column
 - E.g., Student (sid: integer, name: string, gpa: real)
 - **Instance**: A table, with rows (aka tuples, records), and columns (aka fields, attributes) that match the schema
 - # Rows = cardinality
 - # Columns = degree / arity
 - **Set semantics**: (classical relational model) *Every row is unique*
 - **Multiset semantics**: (modern systems, SQL) *Duplicate rows allowed*
- Set of tuples (no duplicates).*
- Exactly same entry twice in a table*

Integrity Constraints

- Describes conditions that must be satisfied by every legal instance
- Types of integrity constraints
 - Domain constraints
 - Primary key constraints
 - Foreign key constraints
 - General constraints