

Logical Database Design: Mapping ER to Relational

Chapter 3, Section 3.5

ER Model vs. Relational Model




- ER Model used for conceptual design
- Relational Model implemented by modern DBMS
- Important Step: Translate ER diagram to Relational schema

Recall ER Constructs



- Basic Constructs
 - Entity Sets
 - Relationship Sets
 - Attributes (of entities and relationships)
- Additional Constructs
 - ISA Hierarchies
 - Weak Entities
 - Aggregation
- Integrity Constraints
 - Key constraints
 - Participation constraints
 - Overlap / Covering constraints for ISA hierarchies

Review: ER relationship types


!!! my constraints offend
here





1-to-1



1-to Many
1-to-N




Many-to-
1
N-to-1




Many-to-Many
M-to-N

ER to Tables: Basics

- **Strong entities:**
 - key = primary key
- **Relationships:** keys come from the participating entities. E.g. for binary relationships:
 - 1-to-1: either key (other = candidate key)
 - 1-to-N: the key of the 'many' (N) part
 - M-to-N: both keys




Entity Sets to Tables



```
CREATE TABLE Citizens
    (ssn    CHAR(11),
     name   CHAR(20),
     bday   DATE,
     PRIMARY KEY (ssn) )
```

Question??



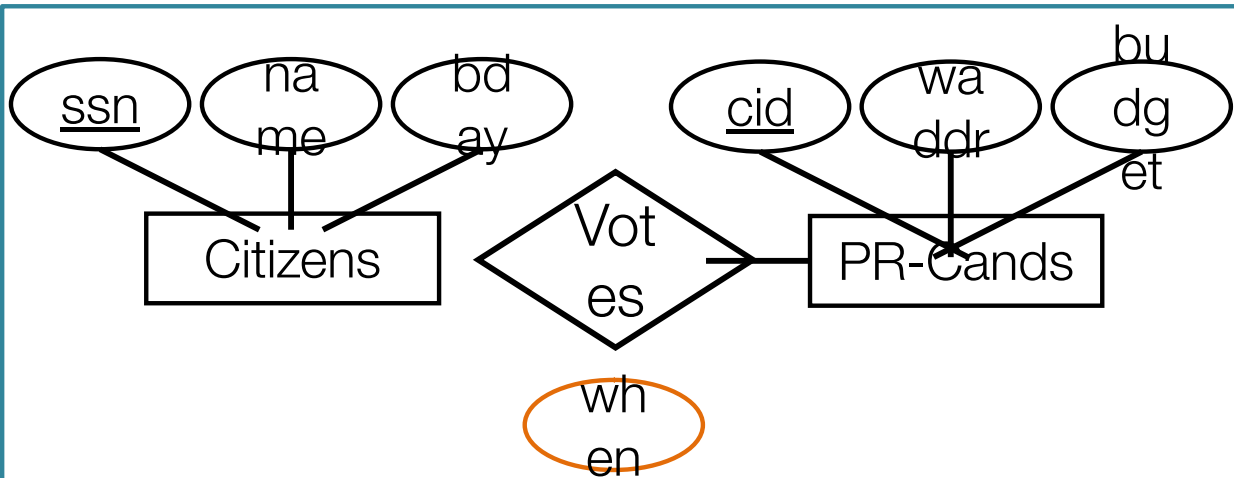
```
CREATE TABLE Citizens
    (ssn CHAR(11),
     name CHAR(20),
     bday DATE,
     PRIMARY KEY (ssn))
```

Can ssn have a null value?



- A. Yes
- B. No
- C. Depends

Relationship Sets to Tables



```
CREATE TABLE Votes (
    ssn      CHAR(11),
    cid      INTEGER,
    when     DATE,
    PRIMARY KEY (ssn, cid),
    FOREIGN KEY (ssn) REFERENCES Citizens,
    FOREIGN KEY (cid) REFERENCES PR-Cands)
```

- Foreign key: keys from participating entity sets
- Descriptive attributes


Generalizes to n-ary relationships

(we will see example later)
same col name => no need specify

Relationship Sets to Tables



*Same entity participate in
two different roles*




```
CREATE TABLE Represents(  
    elected_ssn CHAR(11),  
    cons_ssn CHAR(11),  
    PRIMARY KEY (elected_ssn, cons_ssn),  
    FOREIGN KEY (elected_ssn) REFERENCES Citizens(ssn),  
    FOREIGN KEY (cons_ssn) REFERENCES Citizens(ssn))
```

(s'ssn, ssn) ⇒ make no sense!!!

Note that you need to specify the column that you are referring to in the Citizens table, as Citizens does not have "elected_ssn" nor "cons_ssn".


Key Constraints



- Approach 1: **Three** Tables (Citizens, Votes, PR-Cands)

```
CREATE TABLE Votes
(
    ssn      CHAR(11),
    cid      INTEGER,
    when     DATE,
    PRIMARY KEY (ssn),
    FOREIGN KEY (ssn) REFERENCES Citizens,
    FOREIGN KEY (cid) REFERENCES PR-Cands)
```

Key Constraints




!!!
Each citizen can only vote once,
so OK to fold
'Votes'
relationship into
'Citizens' entity

- Approach 2: Two Tables (Citizen_Votes, PR-Cands)

```
CREATE TABLE Citizen_Votes (  
    ssn      CHAR(11),  
    name     CHAR(20),  
    bday     DATE,  
    when     DATE,  
    cid      INTEGER,  
    PRIMARY KEY (ssn),  
    FOREIGN KEY (cid) REFERENCES PR-Cands)
```

Question??



Each citizen can only vote once, so OK to fold 'Votes' relationship into 'Citizens' entity

- Approach 2: Two Tables (Citizen_Votes, PR-Cands)

```
CREATE TABLE Citizen_Votes (  
    ssn CHAR(11),  
    name CHAR(20),  
    bday DATE,  
    when DATE,  
    cid INTEGER,  
    PRIMARY KEY (ssn),  
    FOREIGN KEY (cid) REFERENCES PR-Cands)
```

*PR-cand table can't be null
cid in*

primary never null

foreign key may be null

Q: Can cid be null?

A – YES; B- NO


Q: Which approach is better?

A – 3 Tables

B – 2 Tables



Key Constraints



preferred with
high partici rate

partici
Constraint

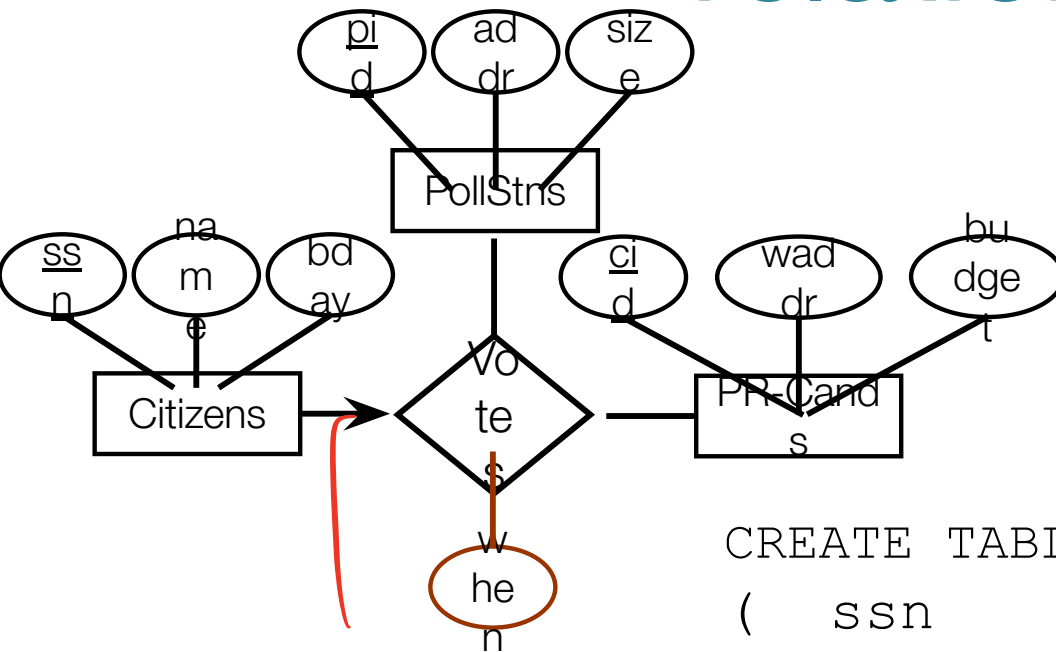
- What about one table?
- No! This is bad design.



e.g., For each citizen that votes for Candidate-X,
we have to store Candidate-X information

(cid, waddr, budget) => **REDUNDANCY!**

Key Constraints: N-ary relationships



key

```
CREATE TABLE Citizens_Votes
```

```
(  ssn      CHAR(11),
```

```
   name     CHAR(20),
```

```
   bday     DATE,
```

```
   when     DATE,
```

```
   pid      INTEGER,
```


```
   cid      INTEGER,
```

```
   PRIMARY KEY (ssn),
```

```
   FOREIGN KEY (cid) REFERENCES PR-Cands,
```

```
   FOREIGN KEY (pid) REFERENCES PollStns)
```

Participation Constraints



Using Approach 2

```
CREATE TABLE Citizen_Votes (
```

```
    ssn      CHAR(11),
```

```
    name     CHAR(20),
```

```
    bday     DATE,
```

```
    when     DATE,
```

```
    cid      INTEGER NOT NULL,
```

```
    PRIMARY KEY (ssn),
```

```
    FOREIGN KEY (cid) REFERENCES PR_Cands
```

```
    ON DELETE NO ACTION);
```

must vote.

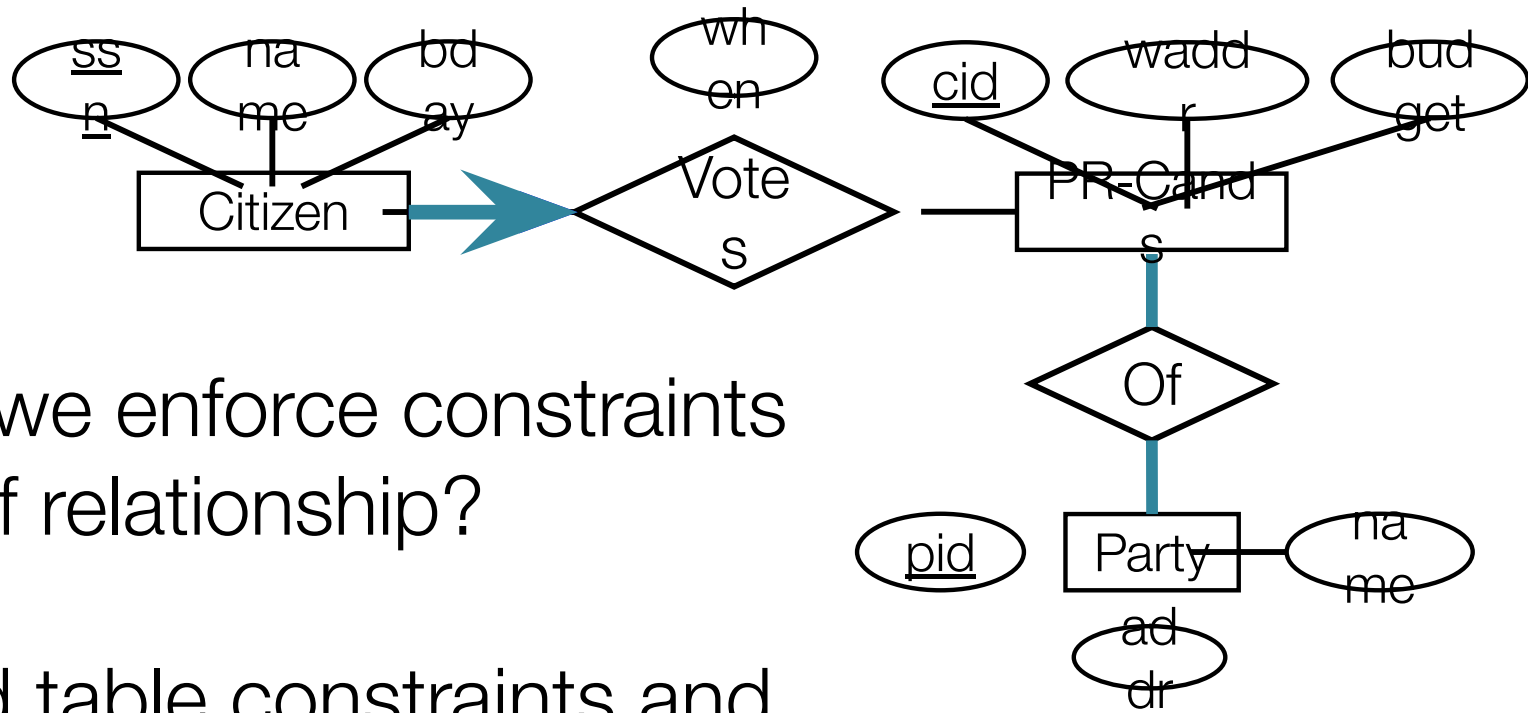
*⇒ just disallow
(Don't allow (and
withdraw))*

Can we enforce the participation constraint using Approach 1 (three tables)?



No!!!


Participation Constraints



Can we enforce constraints on Of relationship?

Need table constraints and assertions (later).

Mapping Participation Constraints (1-to-1)




```
CREATE TABLE RAB (  
    r1 Integer,  
    a1 Integer,  
    a2 Integer,  
    a3 Integer,  
    b1 Integer,  
    b2 Integer,  
    b3 Integer ...)
```

Key constraints?




Mapping Participation Constraints (1-to-1)



```
CREATE TABLE RAB (  
    r1 Integer,  
    a1 Integer,  
    a2 Integer,  
    a3 Integer,  
    b1 Integer NOT NULL,  
    b2 Integer,  
    b3 Integer,  
    UNIQUE (b1), PRIMARY KEY (a1))
```

candidate key

Weak Entities



- Approach 2: Combine weak entity and owning relationship into one relation
 - Delete all weak entities when an owner entity is deleted.

```
CREATE TABLE Expert_Staff (  
  ename      CHAR(20),  
  jobDesc    CHAR(40),  
  salary     REAL,  
  cid        INTEGER,  
  PRIMARY KEY (ename, cid),  
  FOREIGN KEY (cid) REFERENCES PR-Cands  
  ON DELETE CASCADE )
```

→ two table design


⇒ pair for weak entities

ISA Hierarchies: General Approach

separate table for parents and childs,

- Three relations:
 - Citizens (ssn, name, bday)
 - PR-Cands (ssn, waddr, budget)
 - Presidents (ssn, from, to)

- Queries:
 - Involving all citizens => Easy
 - Involving just PR-Cands => need to join PR-cands with Citizens to get some attributes




from base class

ISA Hierarchies: Alternative



- Two relations:
 - PR-Cands (ssn, name, bday, waddr, budget)
 - Presidents (ssn, name, bday, from, to)

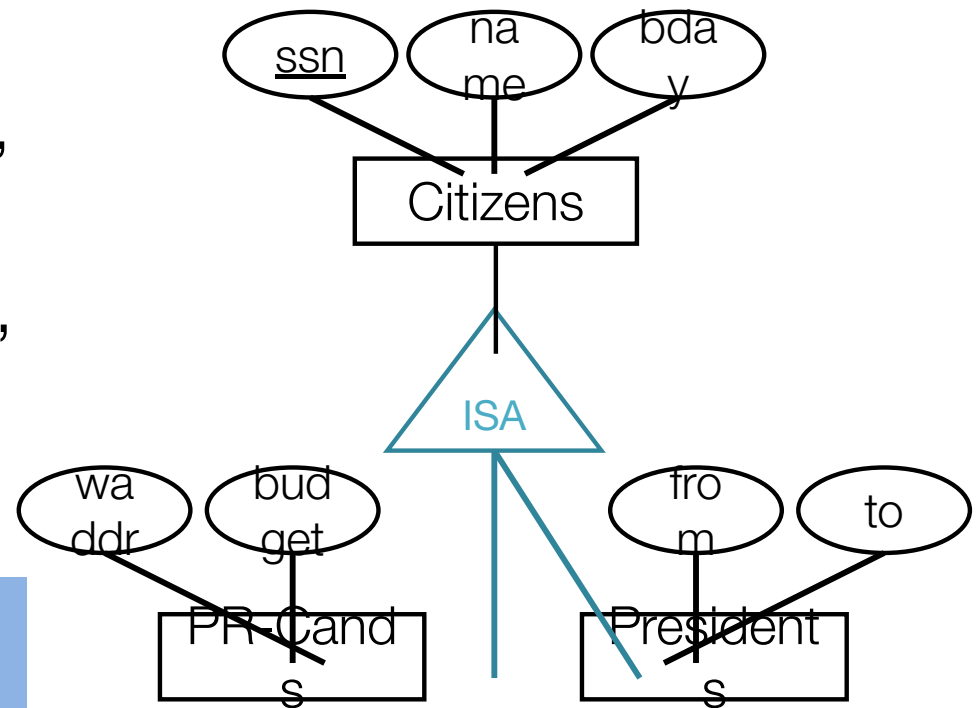
+ citizen (not pre/pre-can)





Question??

- Two relations:
 - PR-Cands (ssn, name, bday, waddr, budget)
 - Presidents (ssn, name, bday, from, to)




This design works best if we have:

- A. Total Covering, Overlapping
- B. Total Covering, Disjoint
- C. Partial Covering, Overlapping
- D. Partial Covering, Disjoint

ISA Hierarchies: Alternative



- Two relations:
 - PR-Cands (ssn, name, bday, waddr, budget)
 - Presidents (ssn, name, bday, from, to)

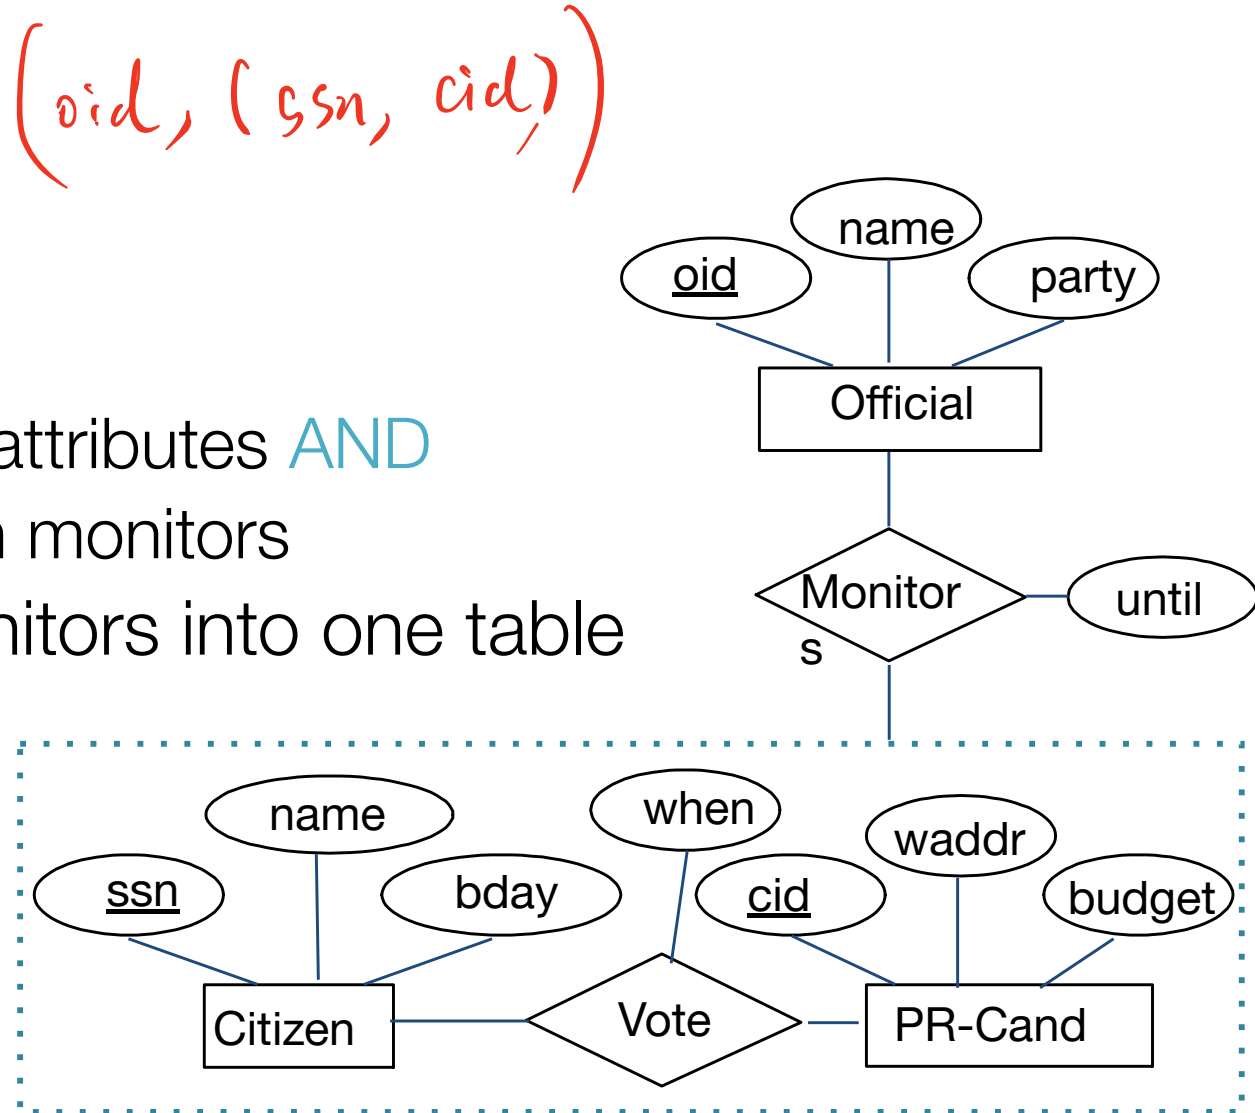


- Problems:
 - What if citizen is both?
 - Redundancy (pre + pre-cand).
 - What if citizen is neither?
 - Use General approach



Aggregation

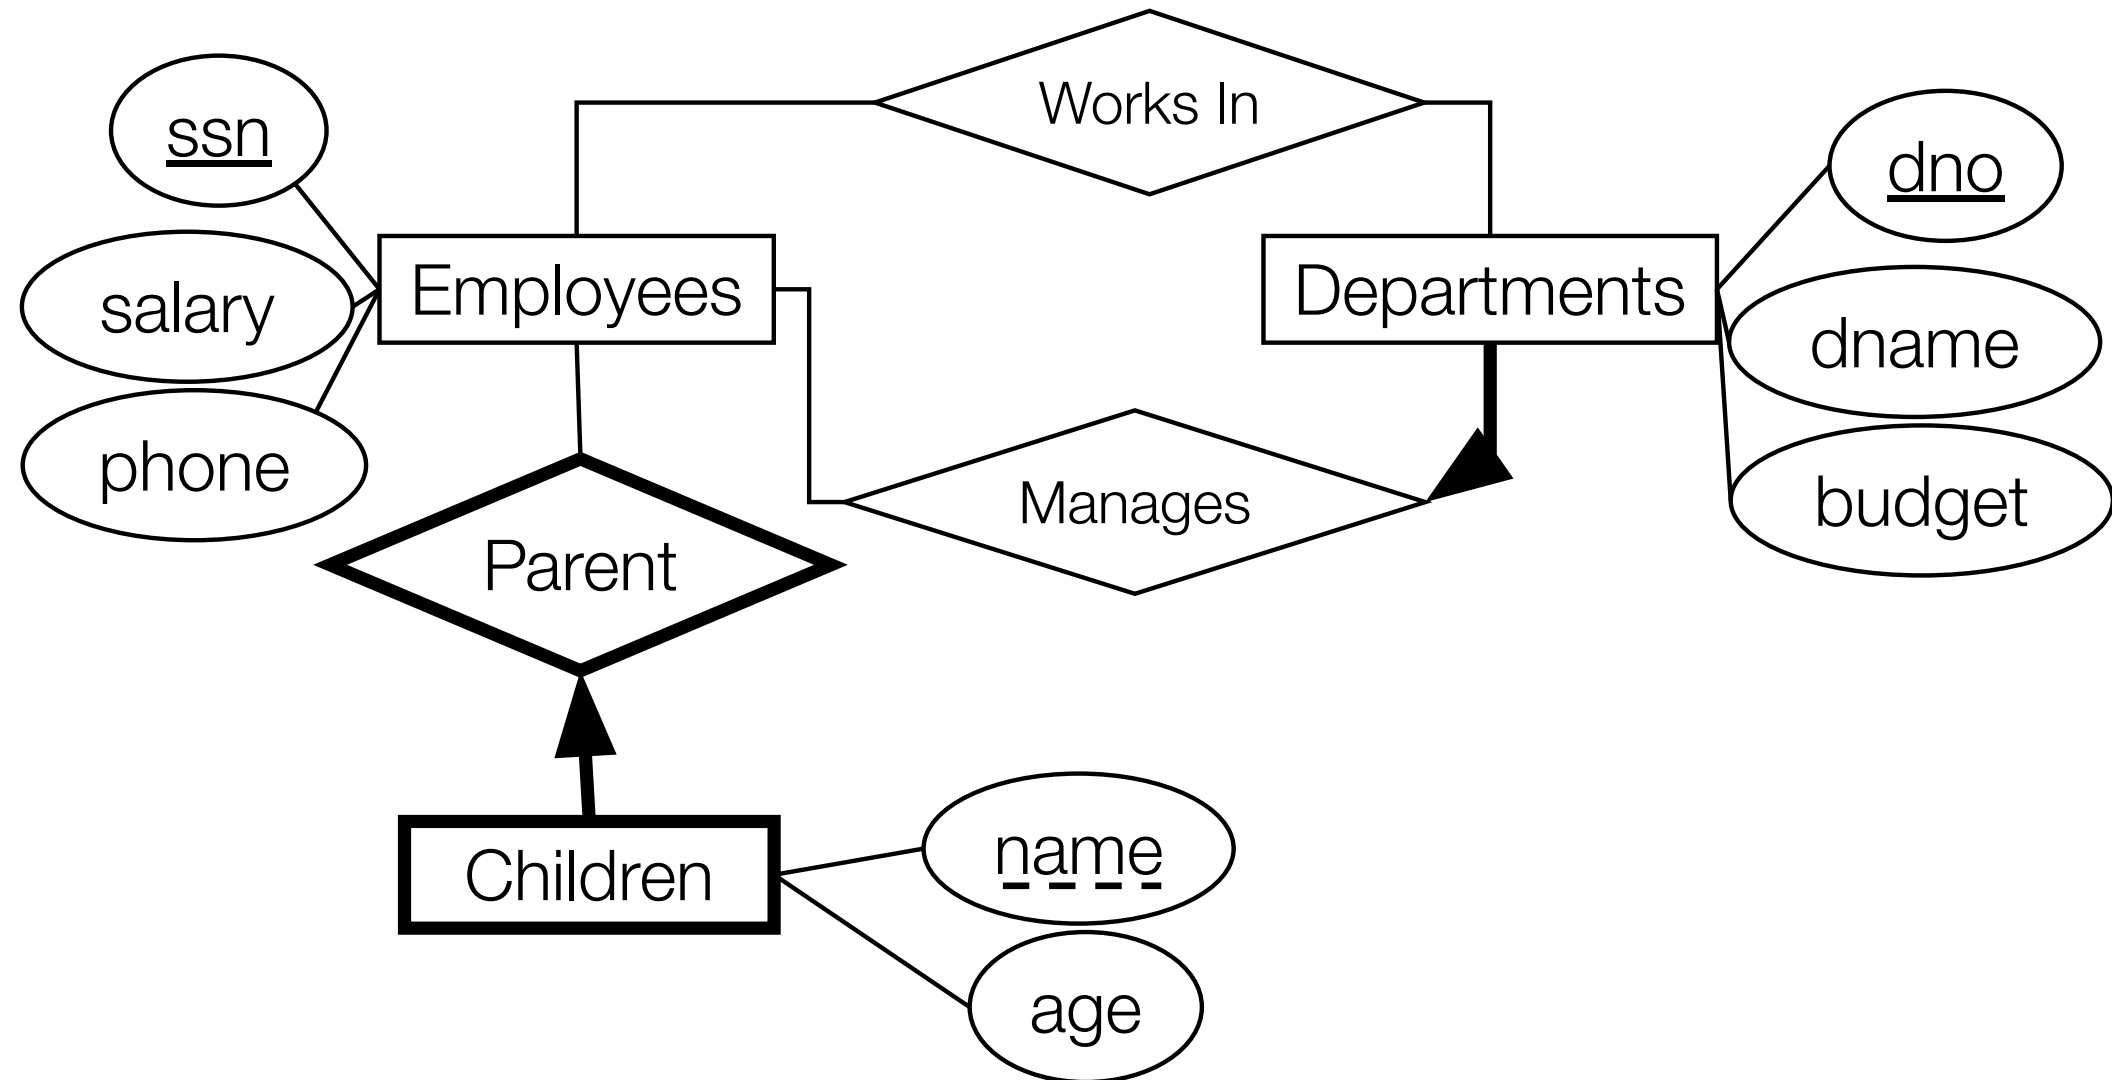
- Keep keys of **all** participating entity sets; decide primary
- Keys for Monitors
 - oid
 - (ssn, cid) *aggre.*
- Q: What if Vote:
 - has no descriptive attributes **AND**
 - total participation in monitors
- A: Fold Vote and Monitors into one table



Exercise – Part 1

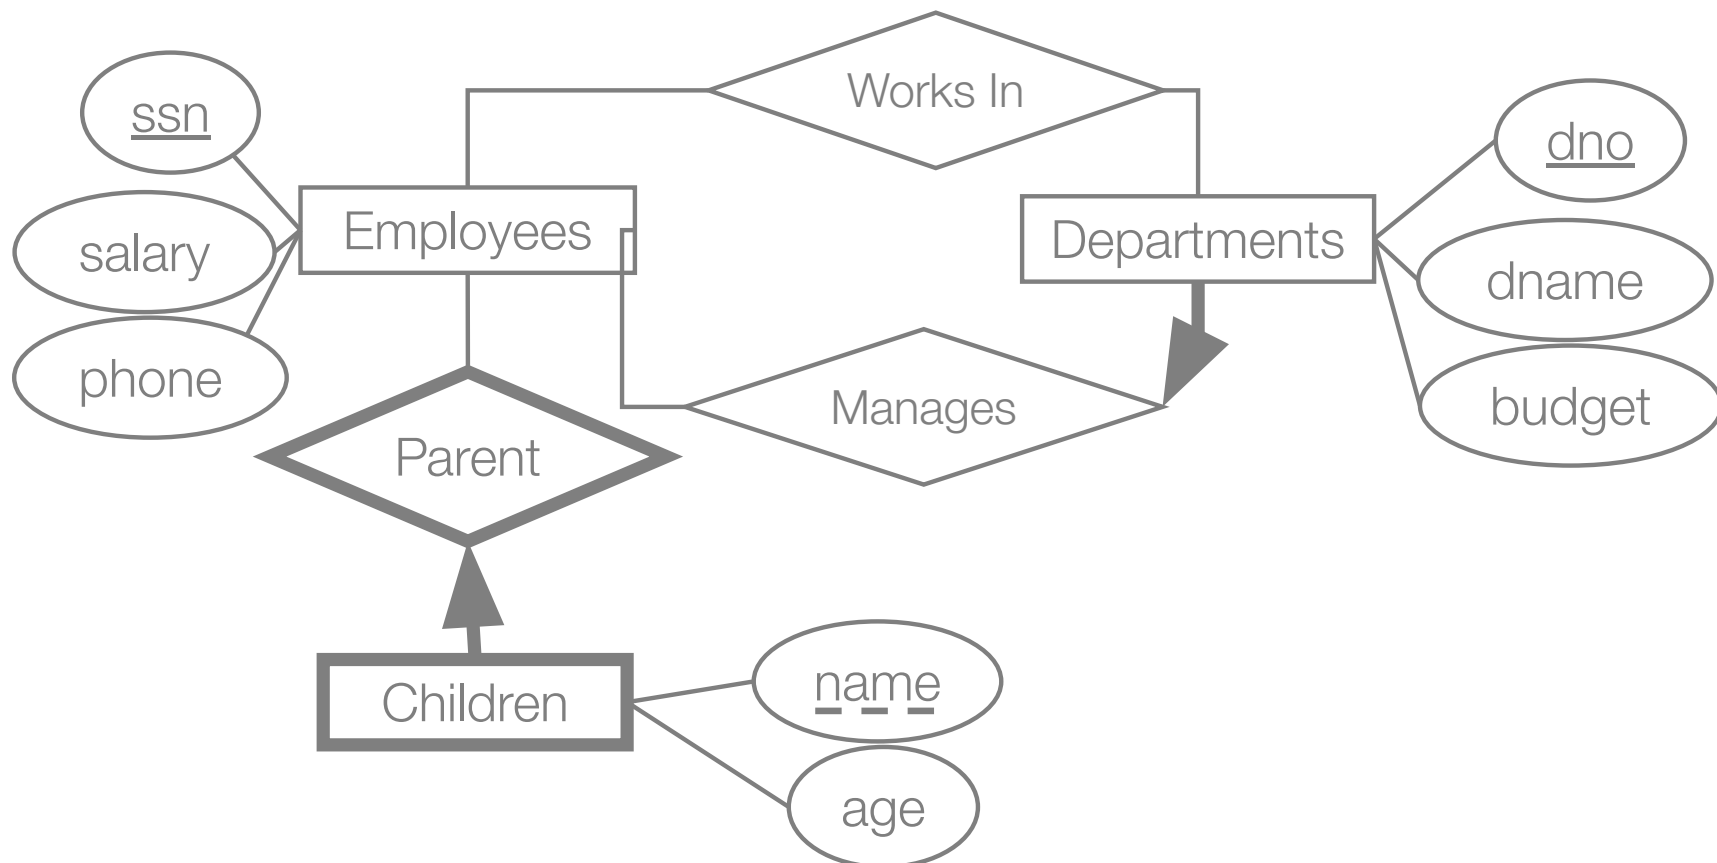
- A company database needs to store information about
 - **employees** (identified by ssn, with salary and phone attributes),
 - **departments** (identified by dno, with dname and budget attributes), and
 - **children of employees** (with name and age attributes).
- Employees work in (zero or more) departments
- Each department is managed by exactly one employee
- A child must be identified uniquely by name when the parent (who is an employee; assume only one parent works for the company) is known.
- We are not interested in information about a child once the parent leaves the company.
- Draw an ER diagram that captures this information

ER Diagram (one solution)



Exercise – Part 2

- Write SQL statements to create the corresponding relations, and to capture as many of the constraints as possible.



SQL DDL – One solution



```
CREATE TABLE Employees (  
  ssn INTEGER,  
  salary REAL,  
  phone CHAR(10),  
  PRIMARY KEY(ssn))
```

```
CREATE TABLE Works (  
  ssn INTEGER,  
  dno INTEGER,  
  PRIMARY KEY (ssn, dno),  
  FOREIGN KEY (ssn)  
    REFERENCES employees,  
  FOREIGN KEY (dno)  
    REFERENCES departments)
```

```
CREATE TABLE Departments (  
  dno INTEGER,  
  dname CHAR(20),  
  budget real,  
  manager INTEGER NOT NULL,  
  PRIMARY KEY (dno),  
  FOREIGN KEY (manager)  
    REFERENCES employees(ssn))
```

```
CREATE TABLE Children (  
  name CHAR(20),  
  age REAL,  
  parent INTEGER NOT NULL,  
  PRIMARY KEY(name, parent),  
  FOREIGN KEY(parent)  
    REFERENCES employees  
    ON DELETE CASCADE)
```

Integrity Constraints

- Describe conditions that must be satisfied by every legal instance
- Types of integrity constraints
 - Domain constraints
 - Primary key constraints
 - Foreign key constraints
 - General constraints

Table Constraints

- More general than key constraints
- Can use a query to express constraint
 - Constraints checked each time table updated
 - CHECK constraint always true for empty relation

```
CREATE TABLE Athlete
(aid INTEGER PRIMARY KEY,
name CHAR(30),
age INTEGER,
country CHAR(20),
sport CHAR(20)),
CHECK ( age >= 18
      AND age <= 80 )

CREATE TABLE Compete
(aid INTEGER, oid INTEGER,
PRIMARY KEY (aid, oid),
FOREIGN KEY (aid)
REFERENCES Athlete
CONSTRAINT noRussia
CHECK (`Russia' <>
      ( SELECT O.city
        FROM Olympics O
        WHERE O.oid=oid)))
```

*cross-table
constraint*

Try it out in sqlplus or sqlite

```
CREATE TABLE Sailors
( sid      INTEGER,
  sname    CHAR(10),
  rating   INTEGER,
  age      REAL,
  PRIMARY KEY (sid),
  CHECK (rating >= 1 AND rating <= 10));

INSERT INTO Sailors VALUES (1, 's1', 11, 25);
```

Do you get a constraint violation error?

Try it out in sqlplus or sqlite SQLite

(Example in book 5.7)

```
CREATE TABLE Sailors
( sid      INTEGER,
  sname    CHAR(10),
  rating   INTEGER,
  age      REAL,
  PRIMARY KEY (sid),
  CHECK (rating >= 1 AND rating <= 10));

INSERT INTO Sailors VALUES (1, 's1', 11, 25);
```

Do you get a constraint violation error?

> Error: CHECK constraint failed: Sailors



More general CHECK

- Interlake boats cannot be reserved.
- Note: these are not supported in Oracle or SQLite

```
CREATE TABLE Reserves
(  sname  CHAR(10),
   bid    INTEGER,
   day    DATE,
   PRIMARY KEY (bid, day),
   CONSTRAINT noInterlakeRes
   CHECK ('Interlake' NOT IN
         ( SELECT B.bname
           FROM Boats B
           WHERE B.bid=bid)))
```

Constraints Over Multiple Relations

- For general constraint over multiple tables, use an **assertion**

*Number of boats
plus number of
sailors is < 100*

```
CREATE ASSERTION smallClub  
CHECK  
((SELECT COUNT (S.sid) FROM Sailors S) +  
 (SELECT COUNT (B.bid) FROM Boats B) < 100)
```

Practical Considerations

- CHECK with subqueries and ASSERTIONS
 - Part of SQL standard (since 1992?)
 - But, they are not supported in many major databases
 - Because it's only been 30 years lolol
 - **Main concern:** Performance issues; CHECK or ASSERTION constraints over multiple tables are very slow
- Instead: **Triggers**
 - Most major database systems support them
 - Triggers are procedural

Active Databases & Triggers (5.8 in the book)

Trigger: Procedure that starts automatically if specified changes occur to the DBMS

- Three parts:
 - Event (activates the trigger)
 - Condition (test that is run when the trigger is activated)
 - Action (what happens if the trigger runs)
 - Before and After Triggers
- Trigger Execution
 - Row-level Triggers: Once per modified row
 - Statement-level Triggers: Once per SQL statement

Triggers

- When condition is checked, success/failure can be used to trigger arbitrary actions.
- Used for many things:
 - Often used to fill out fields in a form
 - Check complex actions (such as credit limit in a shopping application)
 - Check preferred customer status
 - Generate logs for auditing and security checks.
- Must specify when to check condition
 - Usually just before or just after an update

Try it out



- Modify [athlete_create.sql](#) so that it has the UPDATE and DELETE constraints in COMPETE relation as in the previous slide.
 - Modified file available in [athlete_modified.sql](#).
- Try the following and check COMPETE:

```
sqlite> DELETE FROM Athlete WHERE name='Michael  
Phelps';  
sqlite> UPDATE Athlete SET aid=5 WHERE aid=4;
```

Implementation Note. In SQLite, make sure you issued "PRAGMA foreign keys = ON;" command to enforce foreign key constraints. By default, SQLite ignores them for backward compatibility.

Implementation Notes



ORACLE®

- Oracle's sqlplus:
 - You cannot use NO ACTION constraints. They are the default and thus not needed.
 - String literals like 'USA' must use **single quotes**, not double quotes



- SQLite:
 - You need `PRAGMA foreign_keys = ON;` to enforce foreign key constraints. This is for backward compatibility.

Oracle Trigger Example



- First trigger executed before the activating statement, second executes after the activating statement.
- In combination with:
 - FOR EACH ROW - execute once per modified record
 - (default) - execute once per activating statement
- Activating statements:
 - INSERT
 - DELETE
 - UPDATE

```
CREATE TRIGGER init_count
BEFORE INSERT ON Student /* Event */
DECLARE
    count INTEGER; → declare
BEGIN /* Action */
    count := 0;
END;
```

statement-level

```
CREATE TRIGGER incr_count
AFTER INSERT ON Student /* Event */
FOR EACH ROW
WHEN (new.age >= 18) /* Condition */
BEGIN /* Action */
    count := count + 1;
END;
```


row-level trigger

Recall CASCADE constraints

```
CREATE TABLE Athlete  
(aid INTEGER PRIMARY KEY,  
 name CHAR(30),  
 country CHAR(20),  
 sport CHAR(20));
```

```
CREATE TABLE Olympics  
(oid INTEGER PRIMARY KEY,  
 year INTEGER,  
 city CHAR(20));
```

```
CREATE TABLE Compete  
  (aid INTEGER,  
   oid INTEGER,  
   PRIMARY KEY (aid, oid),  
   FOREIGN KEY (aid) REFERENCES Athlete  
     ON DELETE CASCADE  
   FOREIGN KEY (oid) REFERENCES Olympics  
  );
```



CASCADE Using Triggers

```
CREATE TABLE Compete
  (aid INTEGER,
   oid INTEGER,
   PRIMARY KEY  (aid, oid),
   FOREIGN KEY (aid) REFERENCES Athlete
   FOREIGN KEY (oid) REFERENCES Olympics);
```

```
CREATE OR REPLACE TRIGGER cascade_on_delete
AFTER DELETE ON Athlete
FOR EACH ROW
BEGIN
  DELETE FROM Compete
  WHERE Compete.aid = :OLD.aid;
END;
```



Trying out triggers



- Try out the file [athlete_trigger_cascade.sql](#) in sqlplus
- Also try it out in sqlite to see if it supports triggers the same way
- Try removing a row from athlete. Does it cascade to Compete via the trigger?
- Try dropping the trigger:
 - `DROP TRIGGER triggername;`
- What happens now on deleting a row from Athlete?

Triggers: Pitfalls and Pain

- Triggers can be recursive!
 - Chain of triggers can be hard to predict, which makes triggers difficult to understand and debug
- Errors with “mutating” table
 - A table that is currently being modified by an UPDATE, DELETE, or INSERT statement, or a table that might be updated by the effects of a DELETE CASCADE constraint
 - The session that issued the triggering statement cannot query or modify a mutating table
 - Used to prevent a trigger from seeing inconsistent data