


Hash-Based Indexes

Chapter 11

Index Design Space

Organization Structure for k^*

- Tree-based
 - (+) Range, equality search
- Hash-based
 - (+) Equality search
 - Static hashing
 - Extensible hashing
 - Linear hashing



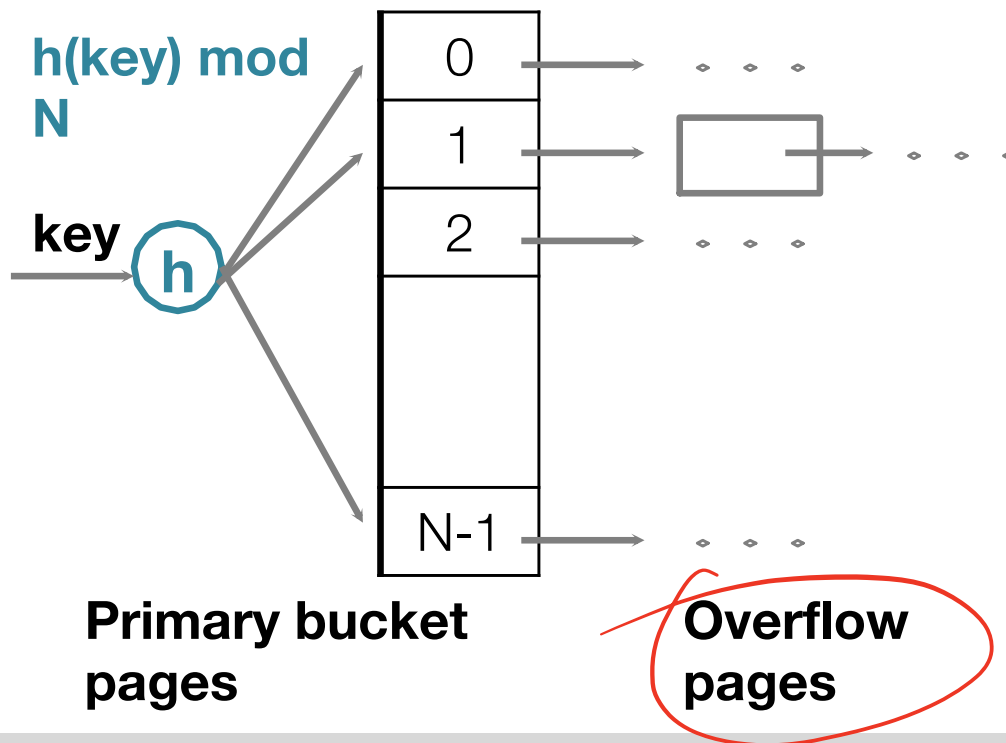
Data Entry (k^*) Contents

1. Actual Data record
index = file
2. $\langle k, \text{rid} \rangle$
actual records in a diff file
3. $\langle k, \text{list of rids} \rangle$

Static Hashing

- # primary bucket pages fixed, allocated sequentially, never de-allocated; overflow pages if needed.
- $h(\text{key}) \bmod N$ = bucket to which data entry with **key** belongs. (N = # of buckets)

bucket size is a page



Question??

The maximum length of overflow chain with static hashing, for a data set of d elements is

- A. 1
- B. $\log d$
- C. d ✓
- D. Infinity (there is no maximum)

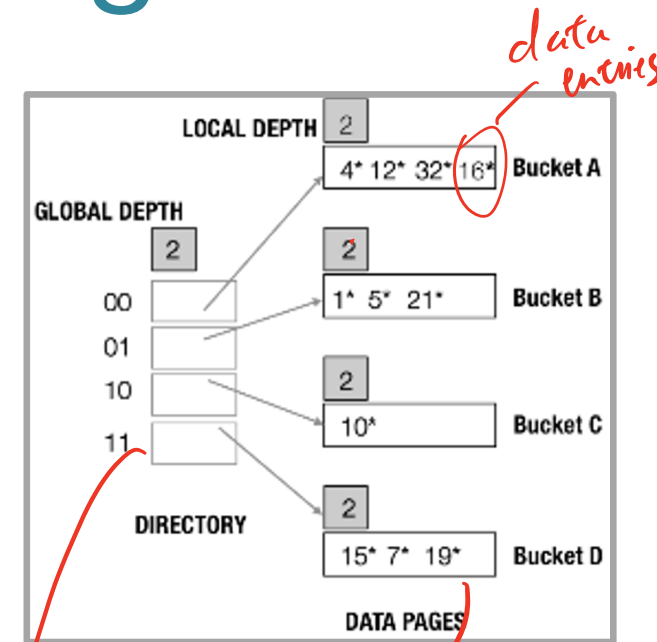
Static Hashing (Contd.)

- Static hashing with fixed # N of buckets can be problematic
 - Long overflow chains can develop (and hurt performance!)
- Might consider periodically doubling N and “rehashing” file
 - Slow when that happens: Entire file has to be read and written
 - Index unavailable while rehashing
- **Dynamic hashing** fixes above problems. Two dynamic hashing techniques: ** double N without rehashing*
 - Extensible hashing
 - Linear Hashing

Extensible Hashing

each bucket is one page in size

- Main Idea: Use a directory (array) of $N = 2^d$ pointers to buckets, where d is called global depth *$d=2$ here*
- Each bucket has a local depth, initially equal to d
- *Invariant:* local depth \leq global depth.
- Search for k :
 - Compute hash function $h(k)$ — an integer
 - Take last global depth # bits of $h(k)$ and use them to index into the directory to find the bucket and retrieve the data entry.



one bucket is a page
2^d buckets

Example: Search

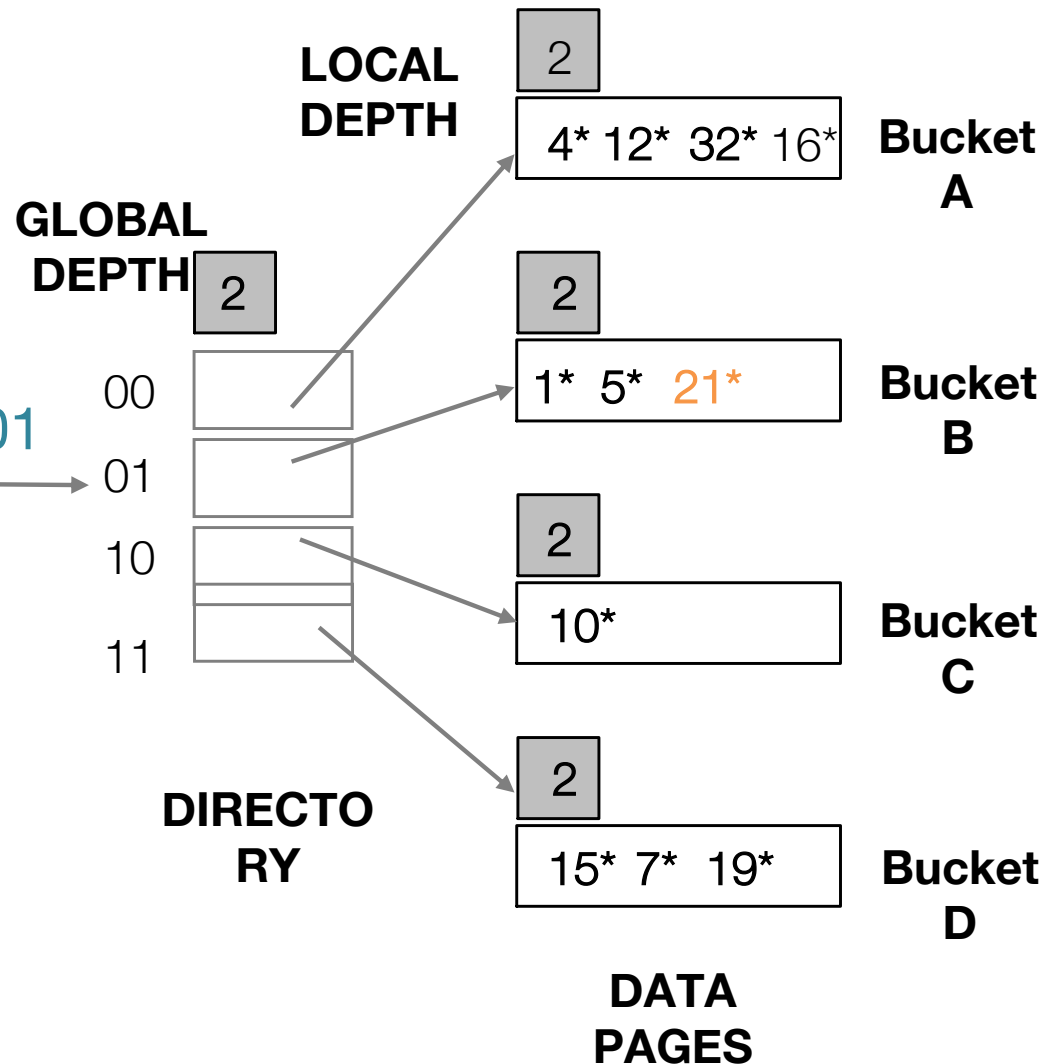
- Find 21^*

- Suppose

$$h(21^*) = 101\boxed{01}$$

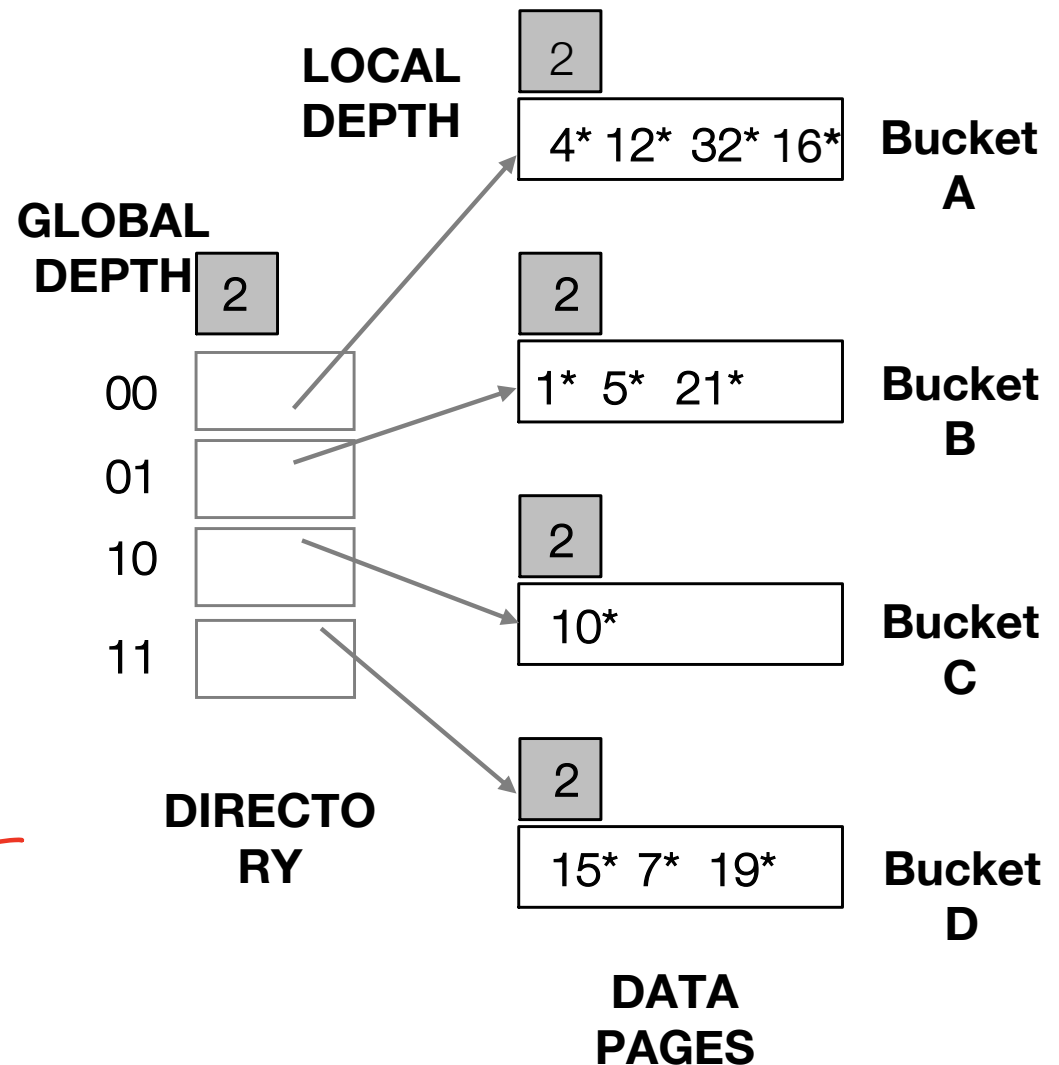
*look at 2 digits
since global
depth is 2*

$$h(21^*) = 10101$$



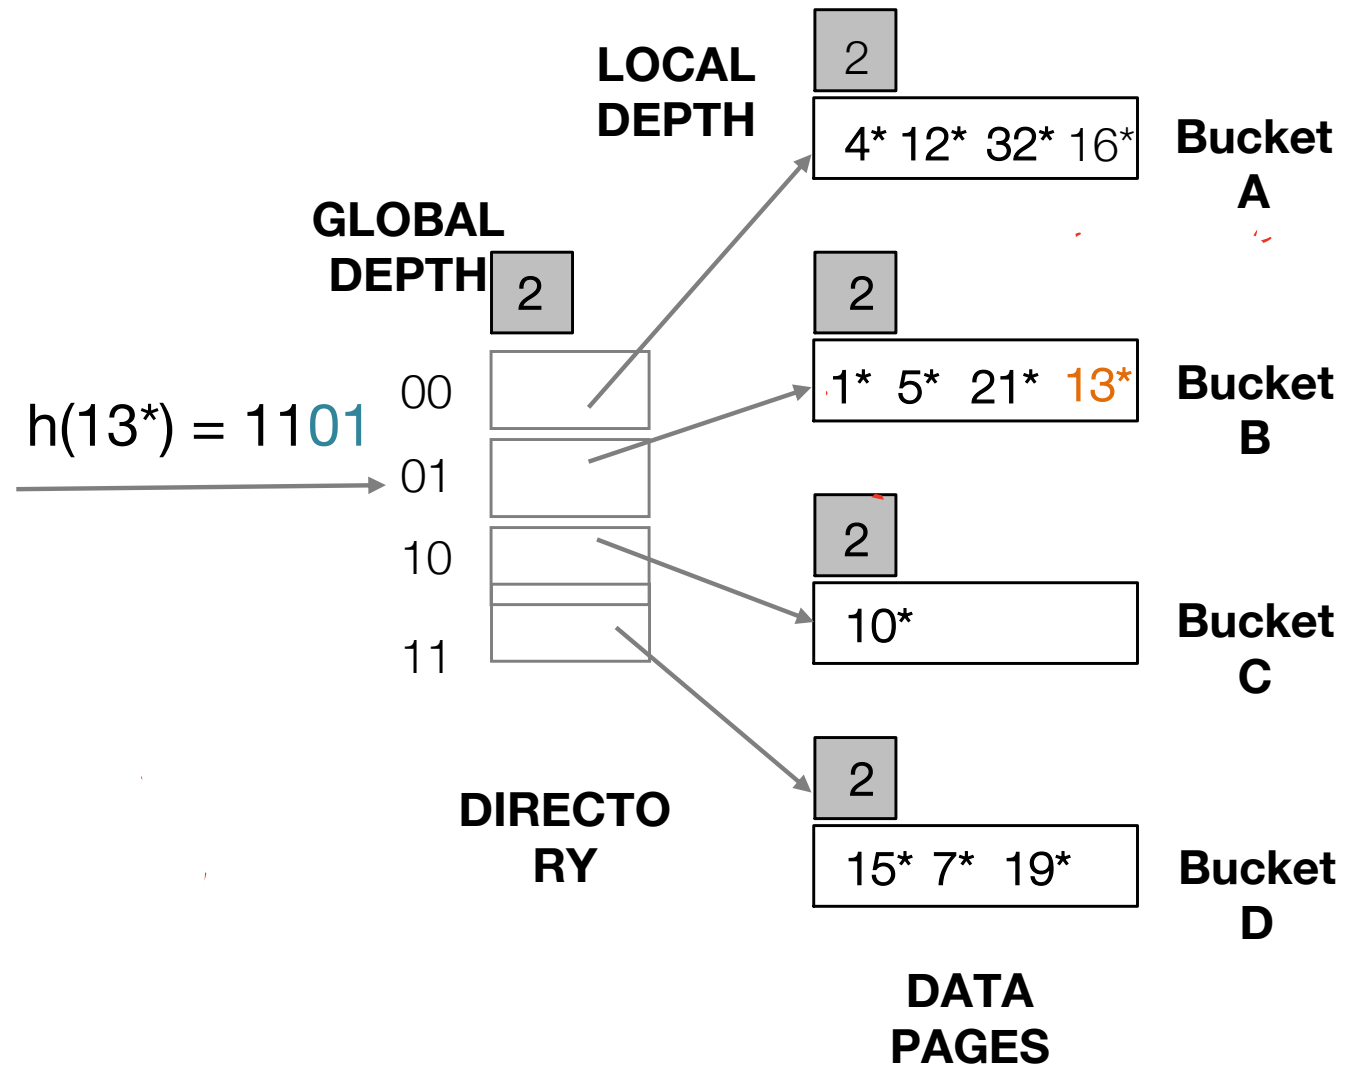
Extensible Hashing: Insert

- Insert k :
 - Take last global depth # bits of $h(k)$
 - If bucket has space, insert, done
 - If bucket is full:
 - Redistribute the values in the overflowed bucket into the two buckets using an extra bit from the index, incrementing their local depth.
 - If buckets' local depth exceeds global depth, restore invariant: double the directory size and increment global depth. Fix pointers.



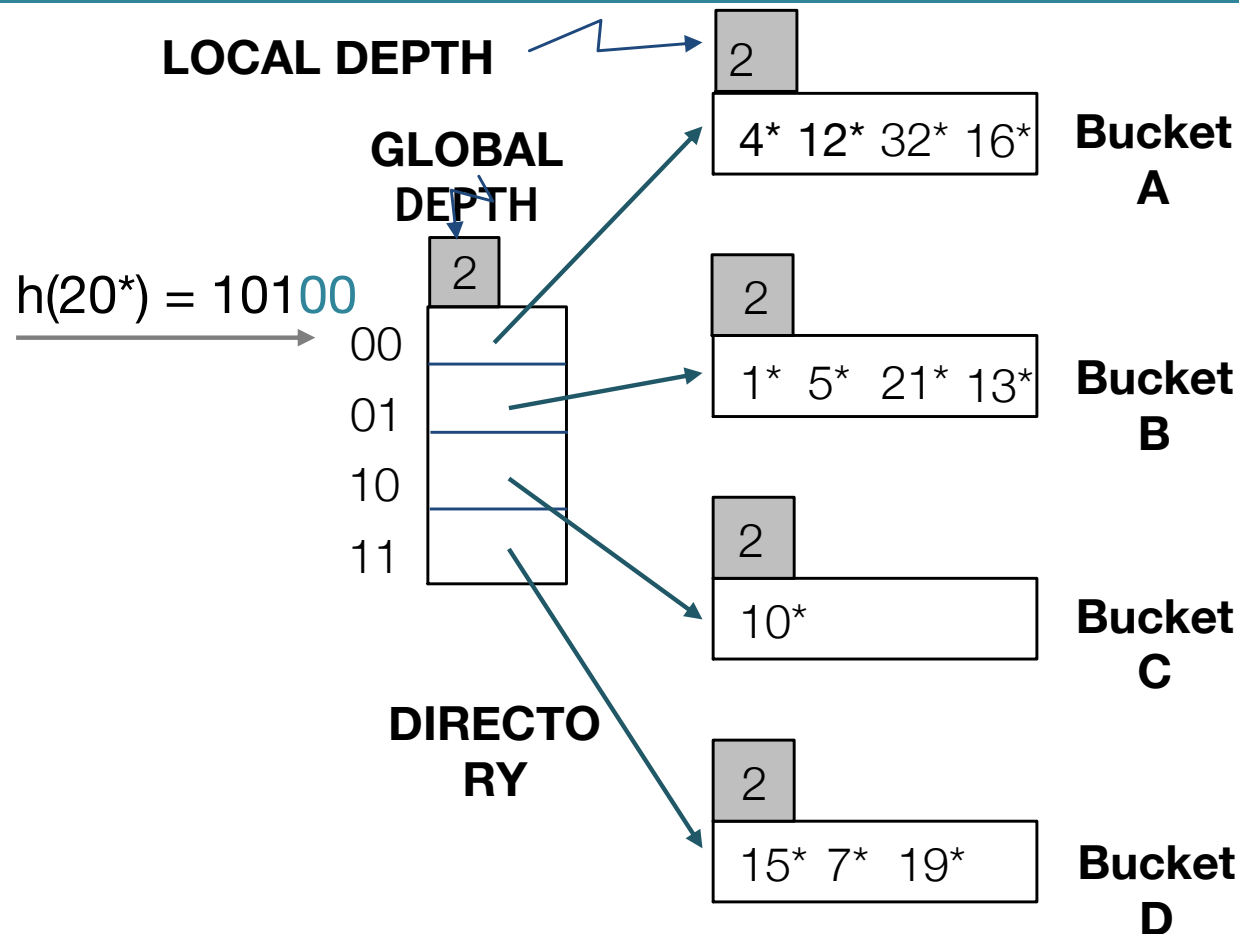
Example: Insertion into a free bucket

- Insert 13^*
- Suppose $h(13^*) = 1101$



Example: Insertion into a full bucket

Insert $20 = (10100)_2$

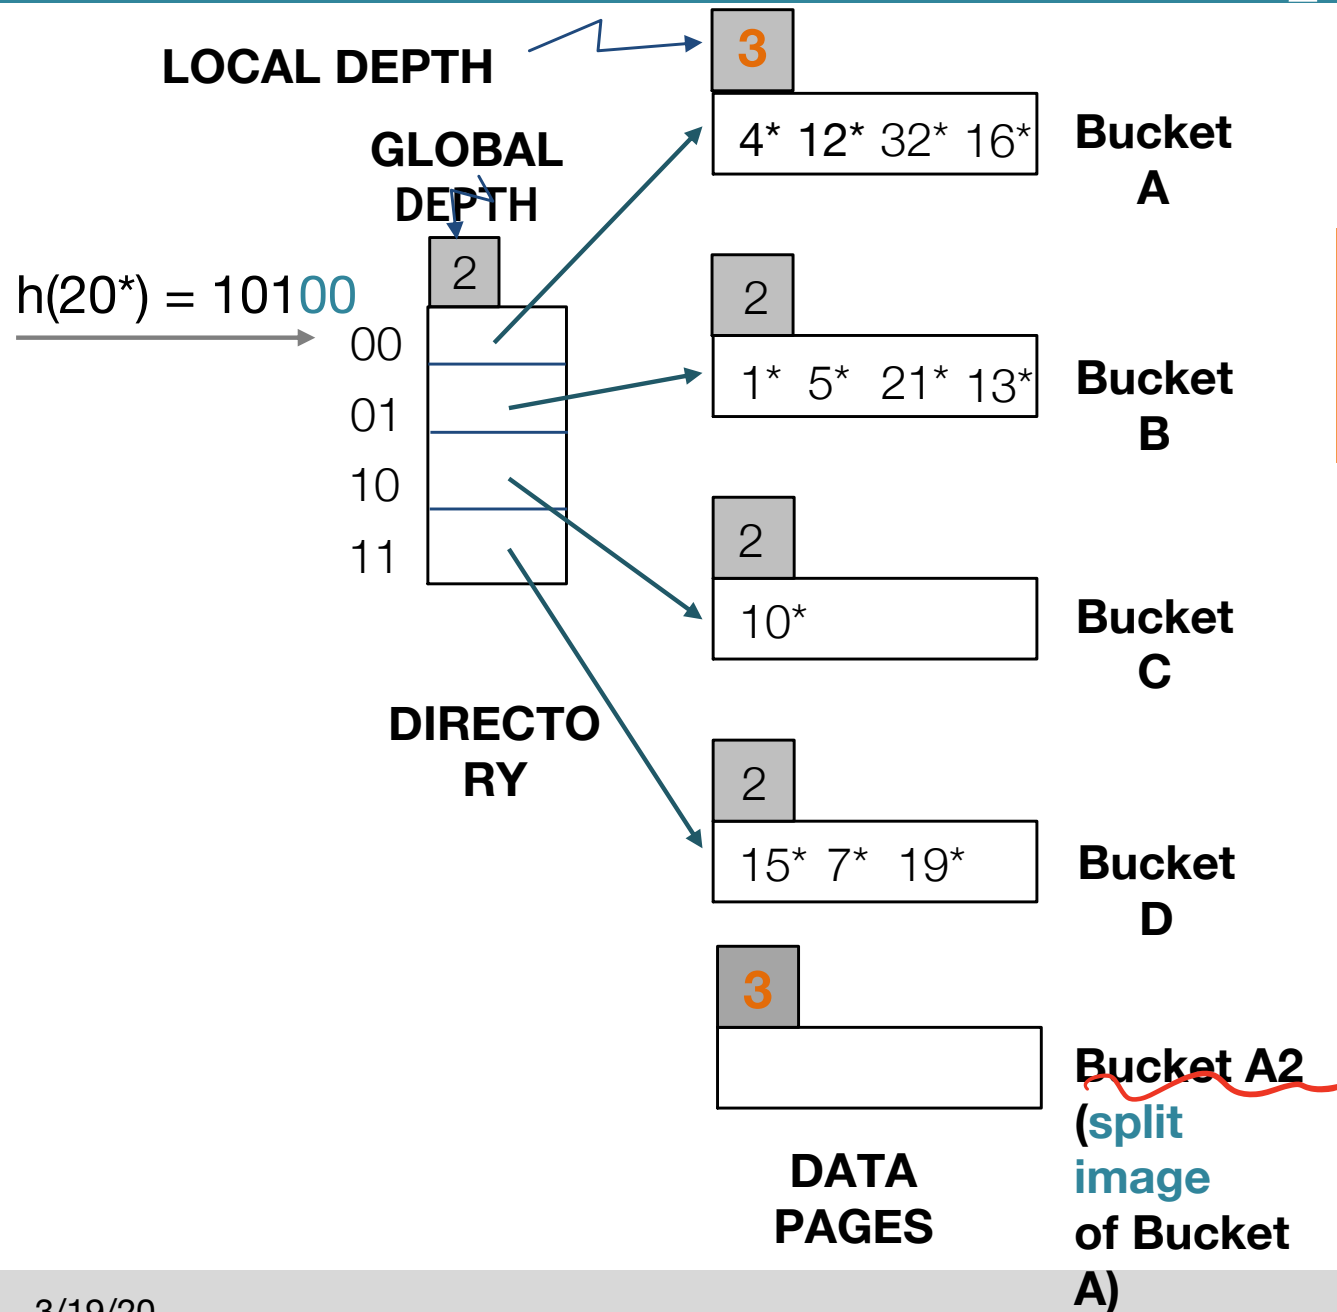


Bucket A is full! No space.

Bucket A is full! No space. Next step: Split Bucket A, incrementing its local depth

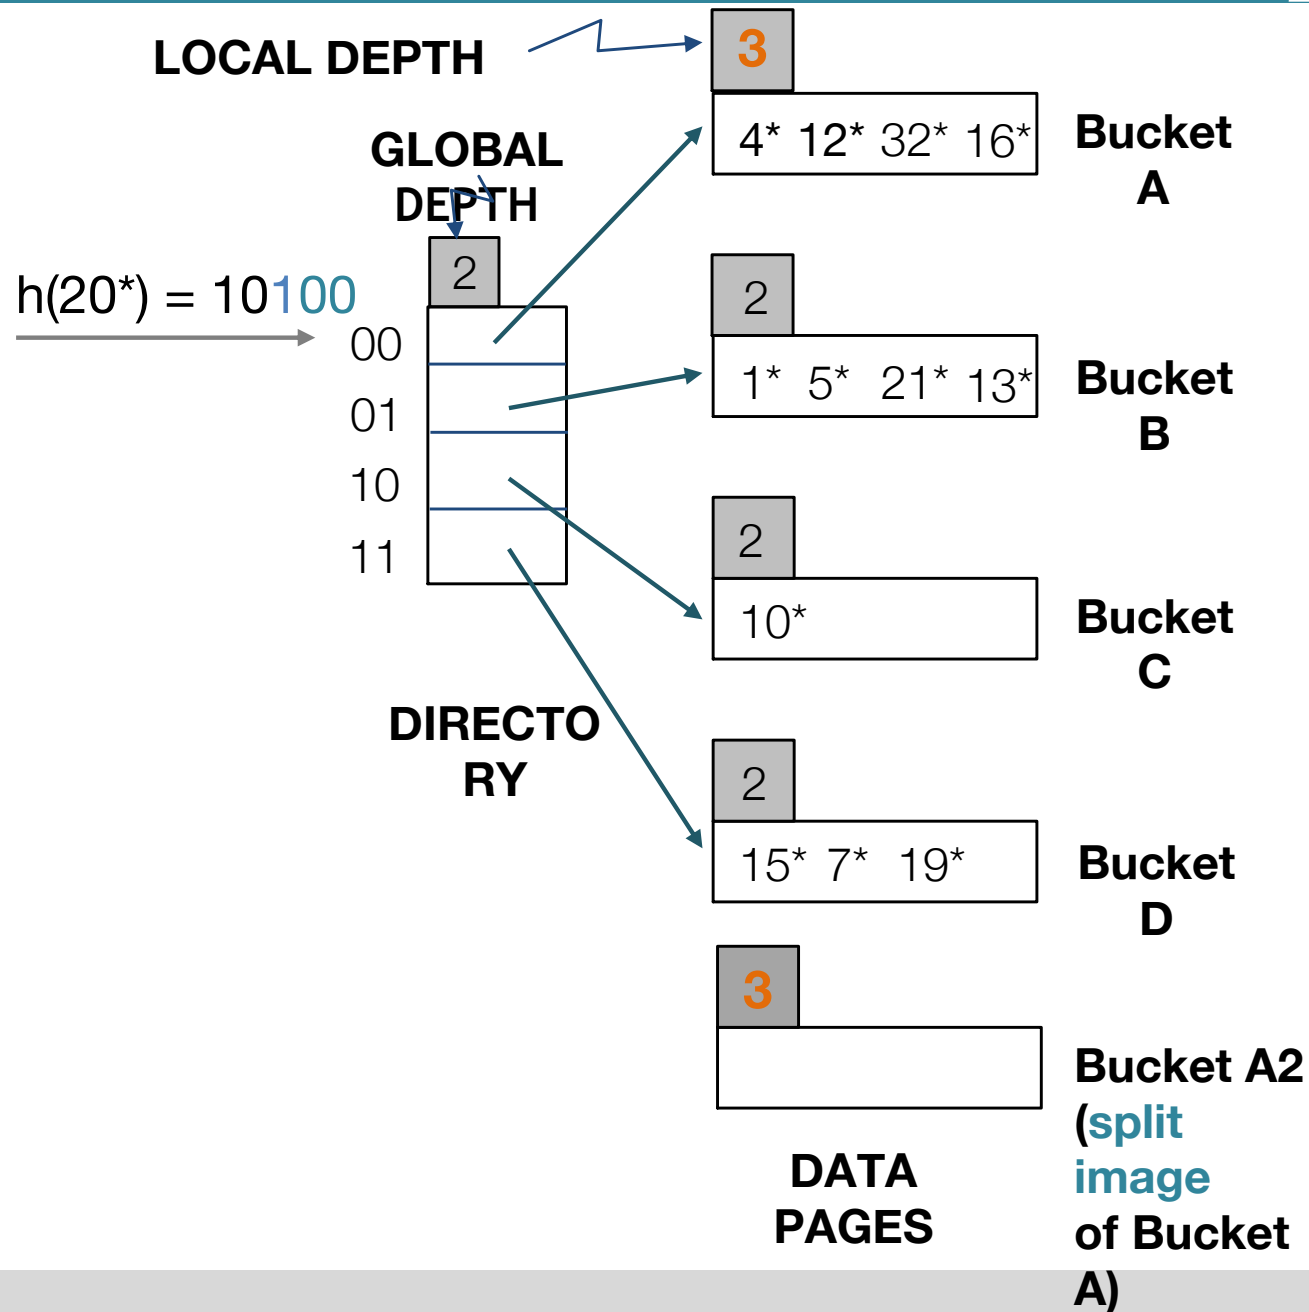
Example: Insertion into a full bucket

Insert $20 = (10100)_2$



Example: Insertion into a full bucket

Insert $20 = (10100)_2$



$$32 = (100000)_2$$

$$16 = (10000)_2$$

Redistribute entries of full bucket and inserted value using the new local depth

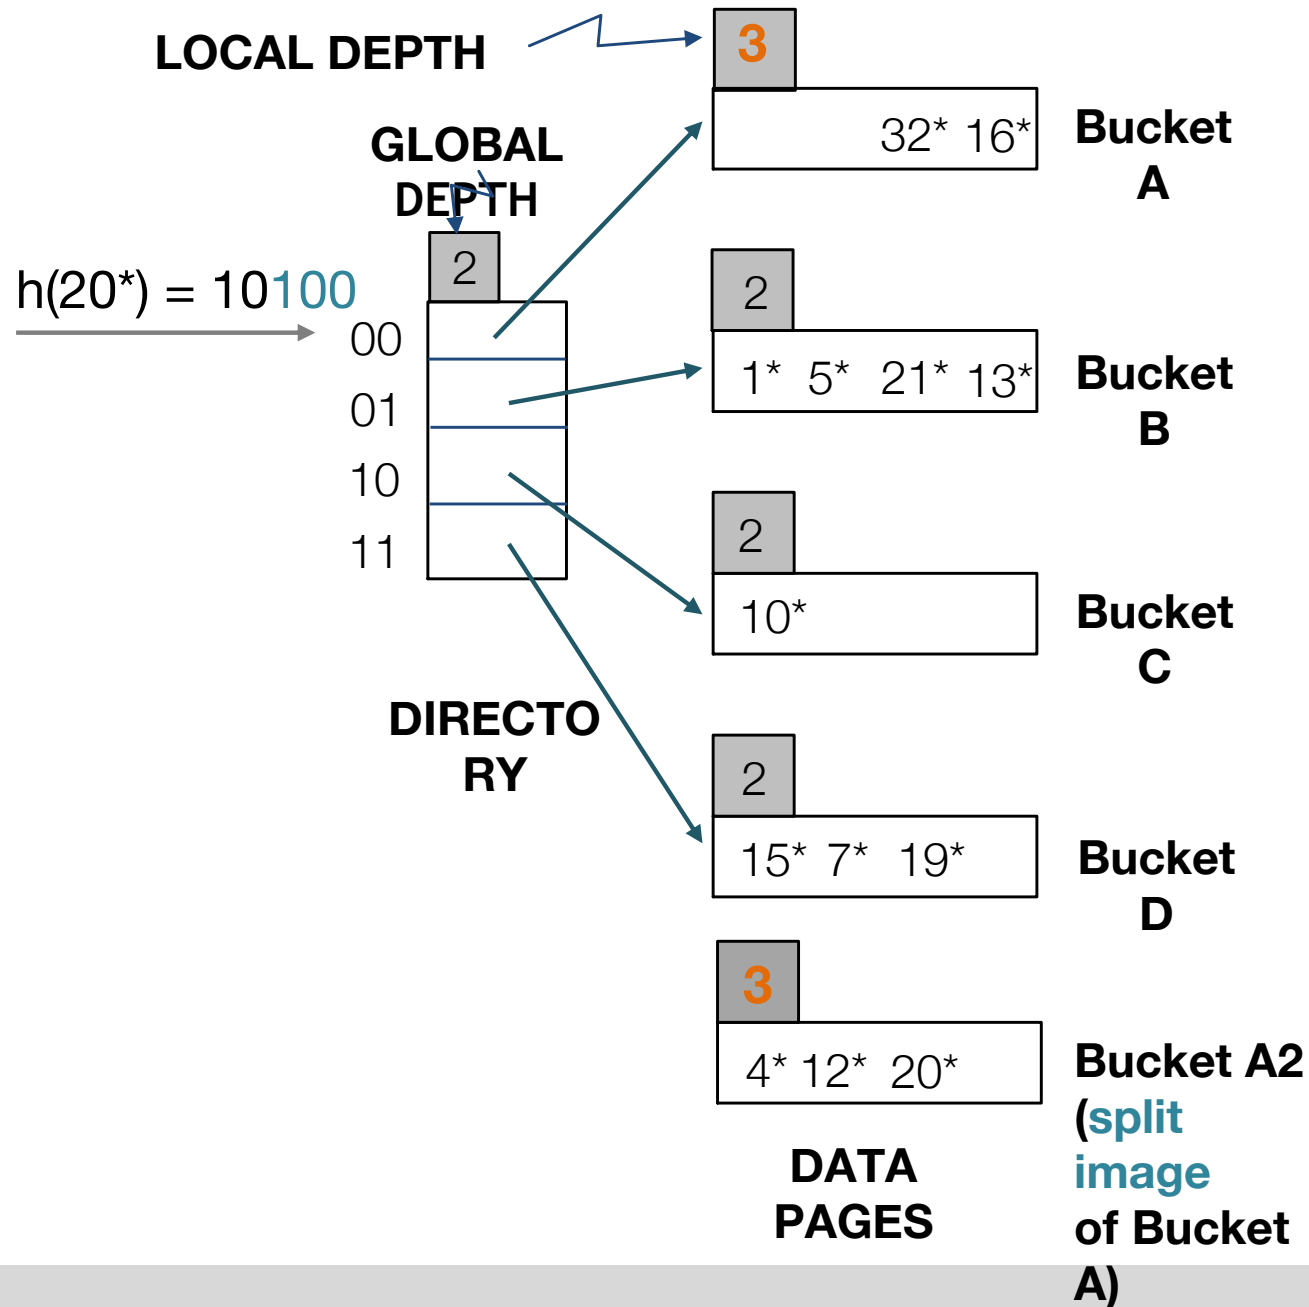
$$4 = (100)_2$$

$$12 = (1100)_2$$

$$20 = (10100)_2$$

Example: Insertion into a full bucket

Insert $20 = (10100)_2$



$$32 = (100000)_2$$

$$16 = (10000)_2$$

Values redistributed.
But, invariant is violated! Fix that by doubling the directory

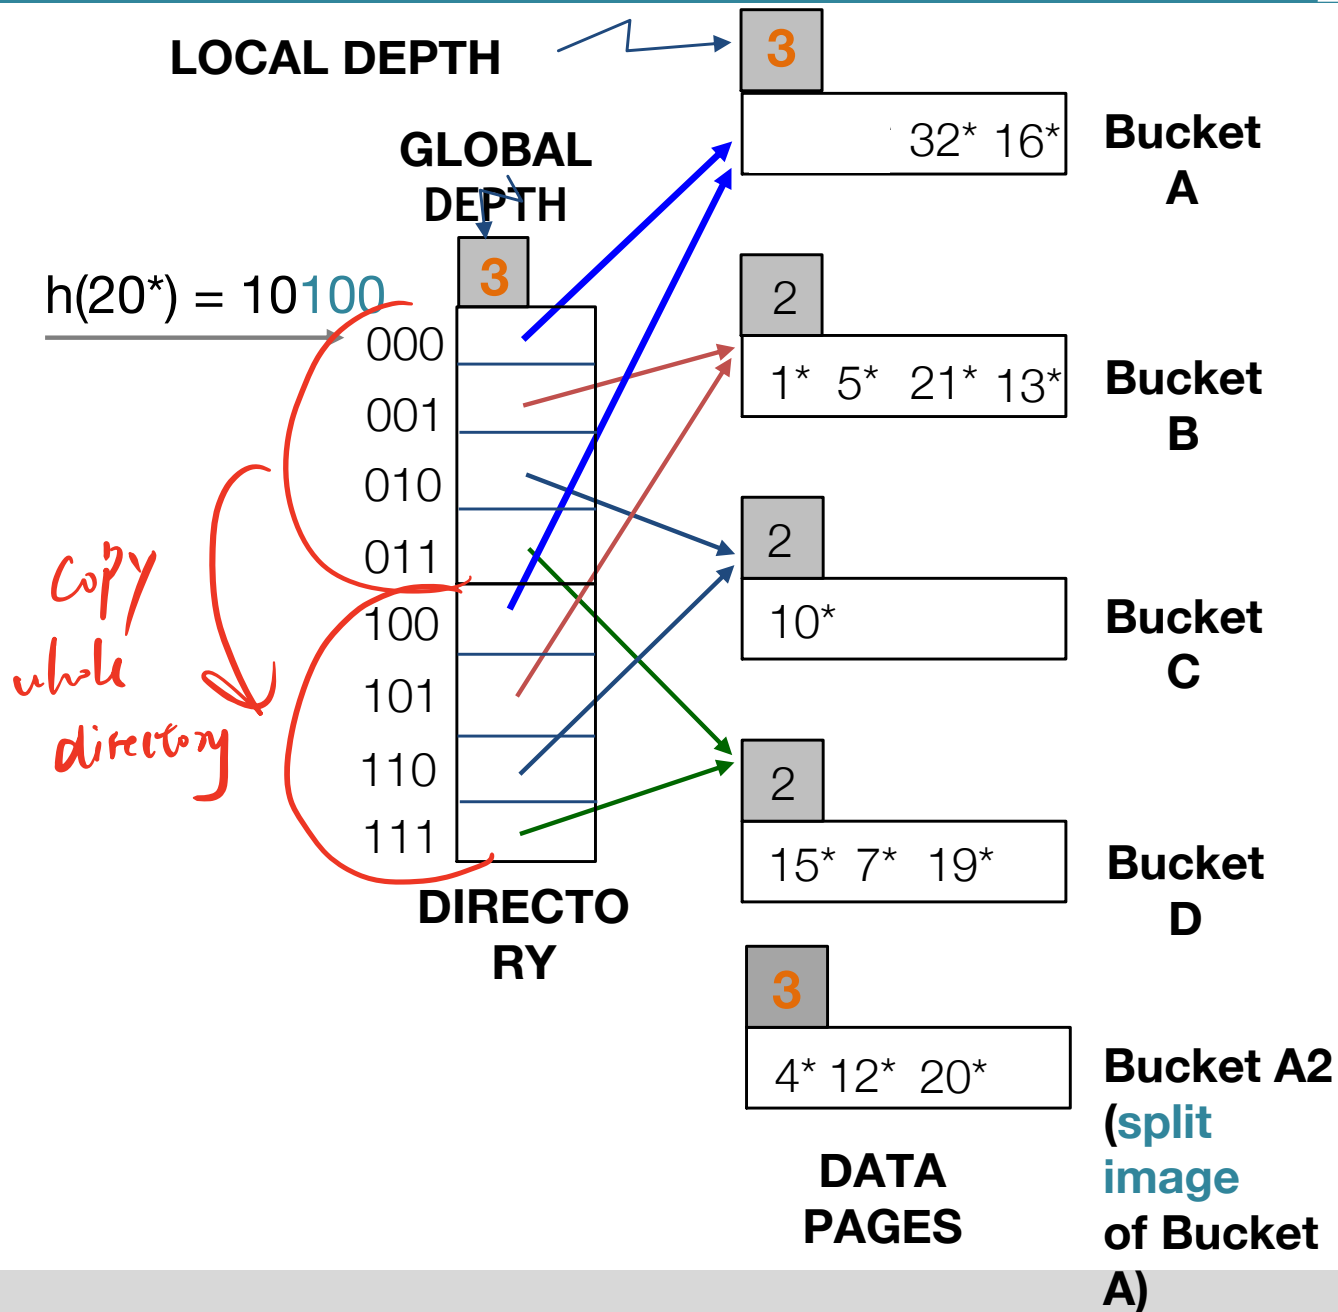
$$4 = (100)_2$$

$$12 = (1100)_2$$

$$20 = (10100)_2$$

Example: Insertion into a full bucket

Insert $20 = (10100)_2$



$$32 = (100000)_2$$

$$16 = (10000)_2$$

Notice that splitting a bucket only requires doubling the directory when new LD > old GD

Directory doubled by *copying it over*

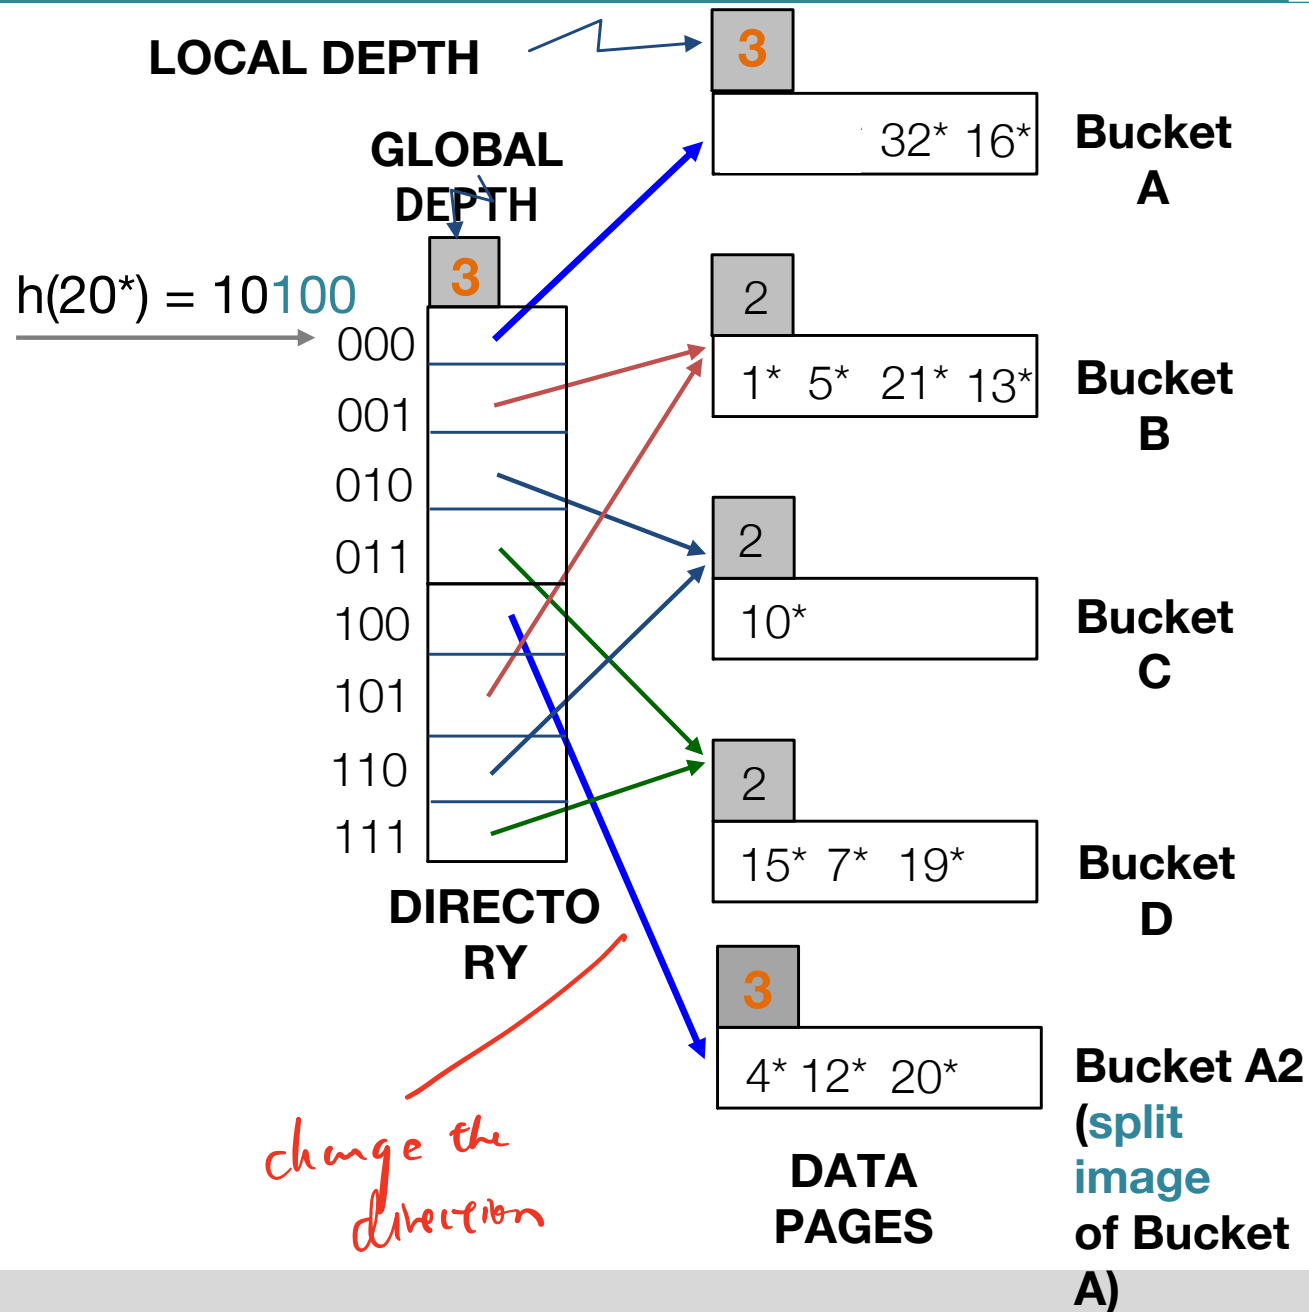
$$4 = (100)_2$$

$$12 = (1100)_2$$

$$20 = (10100)_2$$

Example: Insertion into a full bucket

Insert $20 = (10100)_2$



$$32 = (100000)_2$$

$$16 = (10000)_2$$


Fix the pointer for the split image bucket

$$4 = (100)_2$$

$$12 = (1100)_2$$

$$20 = (10100)_2$$

Insert $18 = (10010)_2$



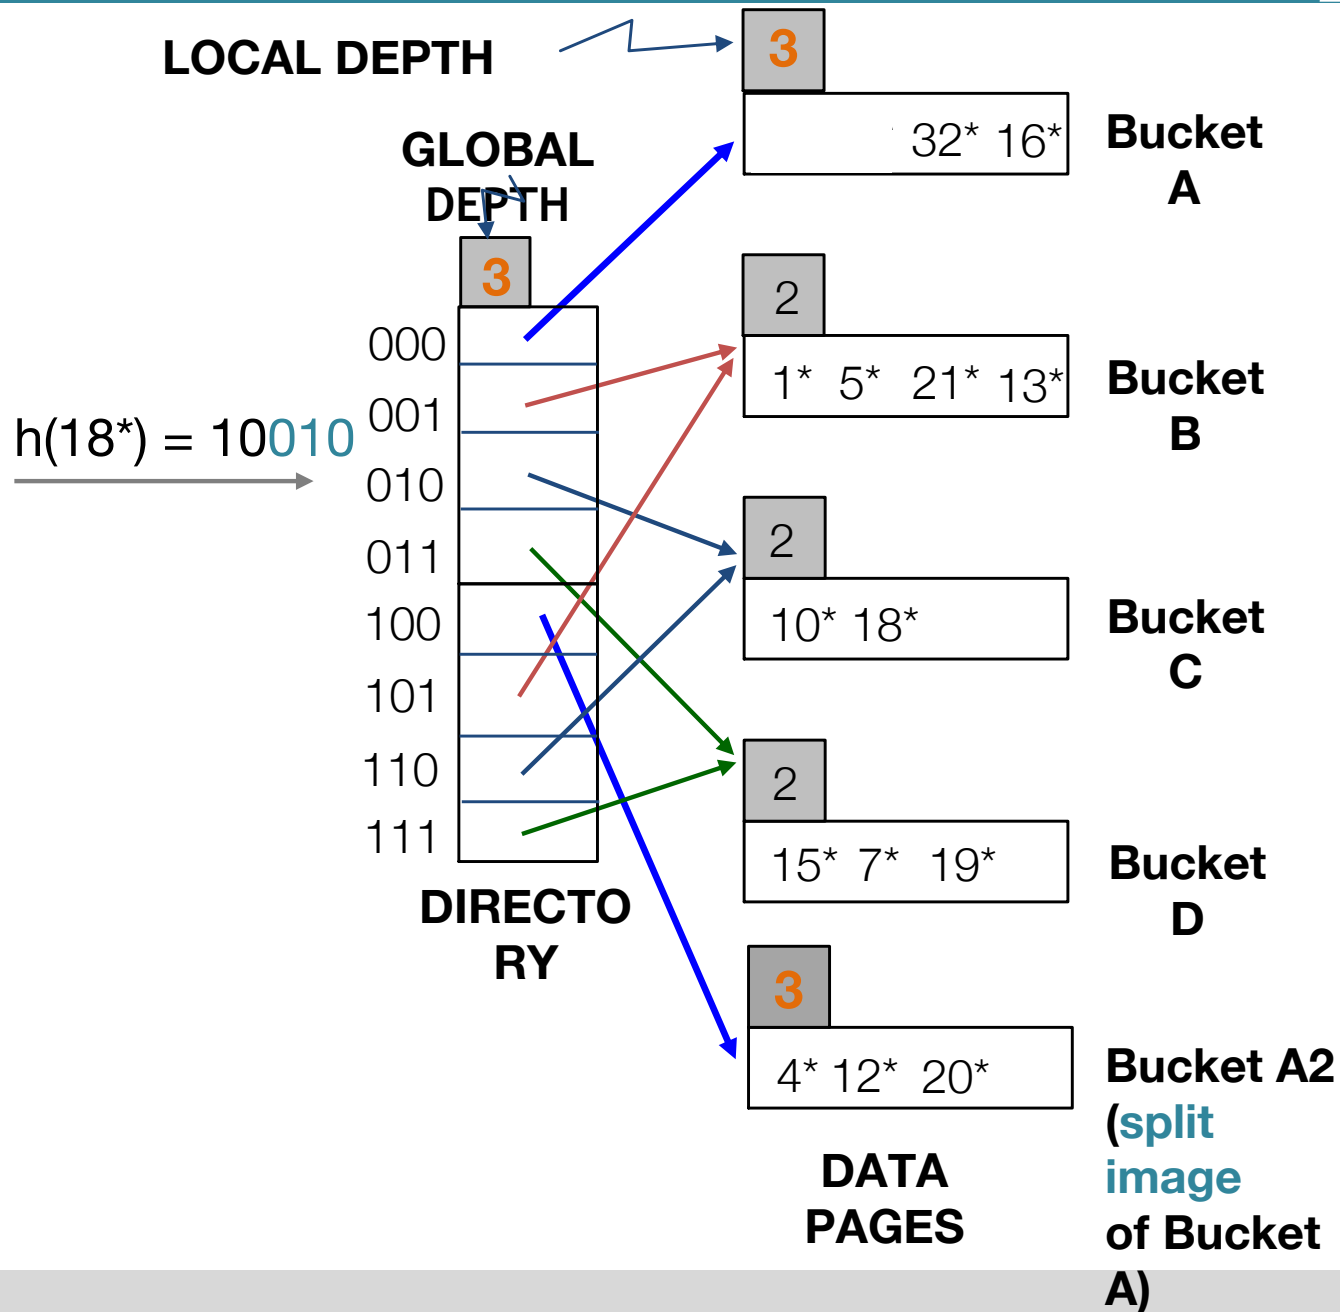
Use last 3 bits to hash

gloбал дупећ

$$\begin{aligned} 4 &= (100)_2 \\ 12 &= (1100)_2 \\ 20 &= (10100)_2 \end{aligned}$$

Example: Insertion into a free bucket

Insert $18 = (10010)_2$



$$32 = (100000)_2$$
$$16 = (10000)_2$$

Use last 3 bits to hash

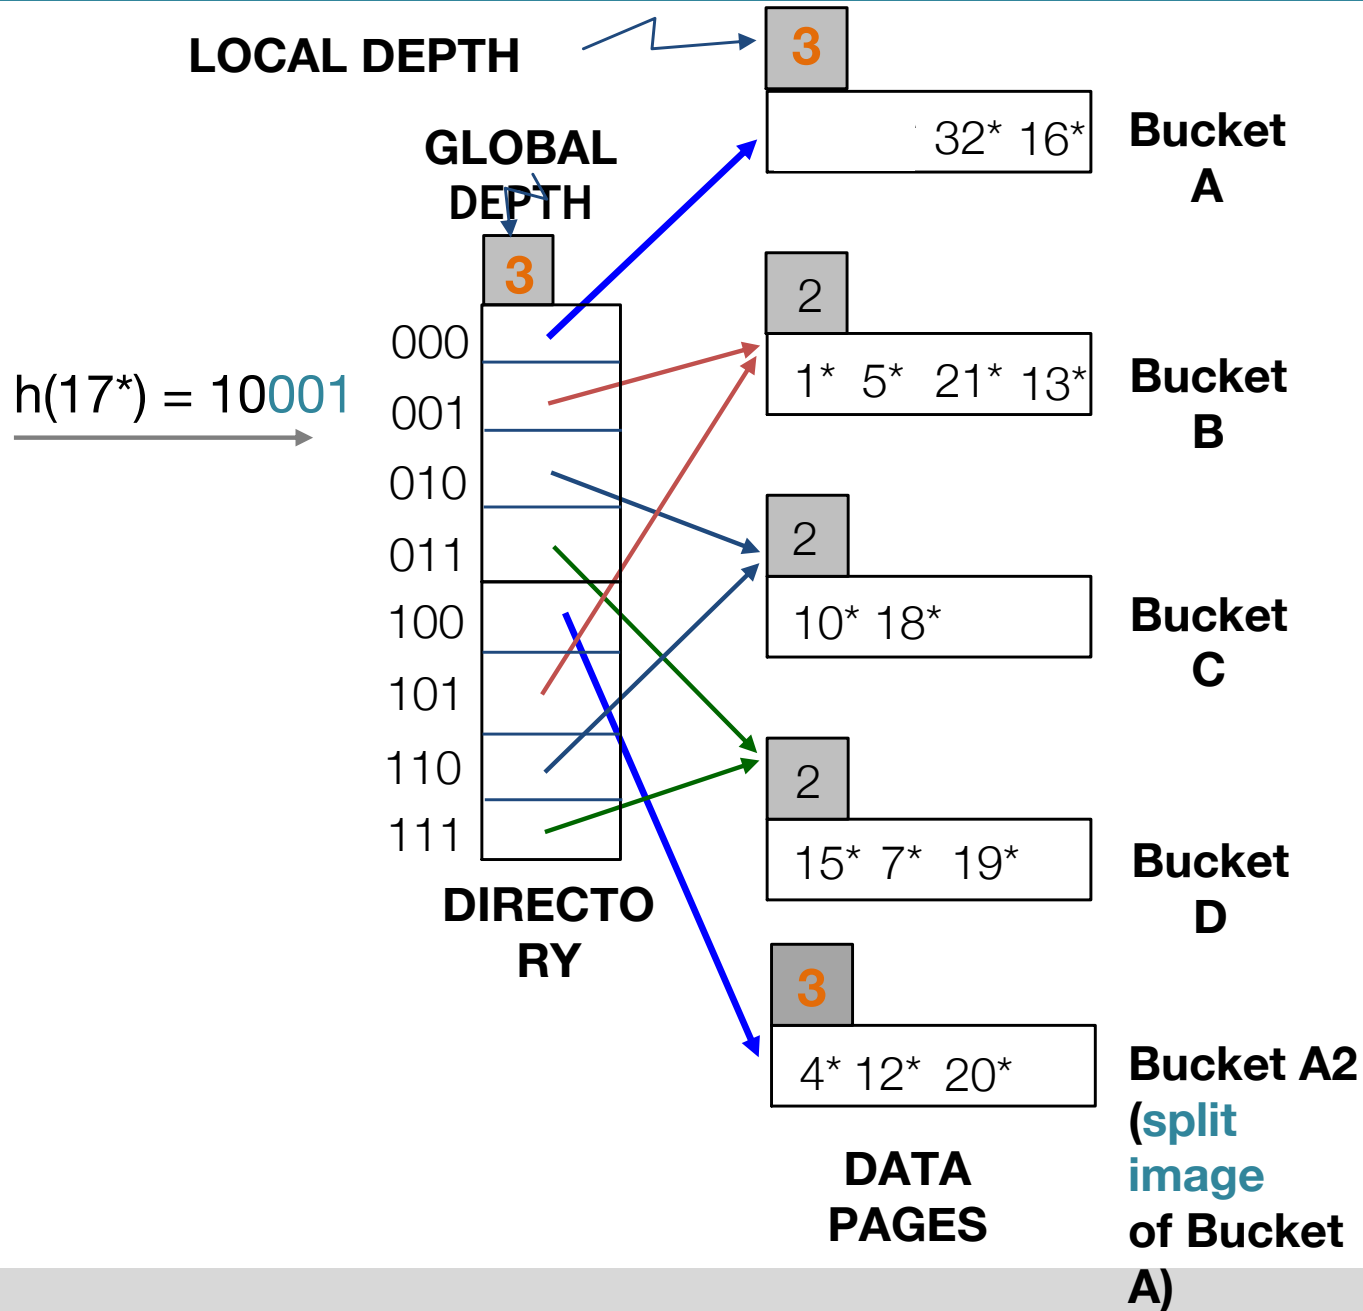
$$4 = (100)_2$$
$$12 = (1100)_2$$
$$20 = (10100)_2$$

Question??

Local depth is

- A. Always less than or equal to global depth ✓
- B. Always greater than or equal to global depth
- C. Can be less than, greater than, or equal to global depth
- D. None of the above

Question: What will happen if 17 is inserted?



$$1 = (000001)_2$$

$$5 = (000101)_2$$

$$21 = (010101)_2$$

$$13 = (001101)_2$$





Comments on Extensible Hashing

- one page look up*
 - How many disk accesses for equality search?
 - One if directory fits in memory, else two
 - Directory grows in spurts, and, if the distribution of hash values is skewed, directory can grow large

*lots of with low local depth
but a few with high local depth*

Question??

Extensible hashing still needs overflow pages

A. True ✓

B. False

↓
split and redistribute
but what if all entries
are the same?

Answers

- Multiple entries with same hash value cause problems => We could still need overflow pages

2 insert page count: one read ^① one write ^① (no split)


3 (with split): read, insert, create new copy, write two pages

- Delete: Reverse of inserts – see textbook

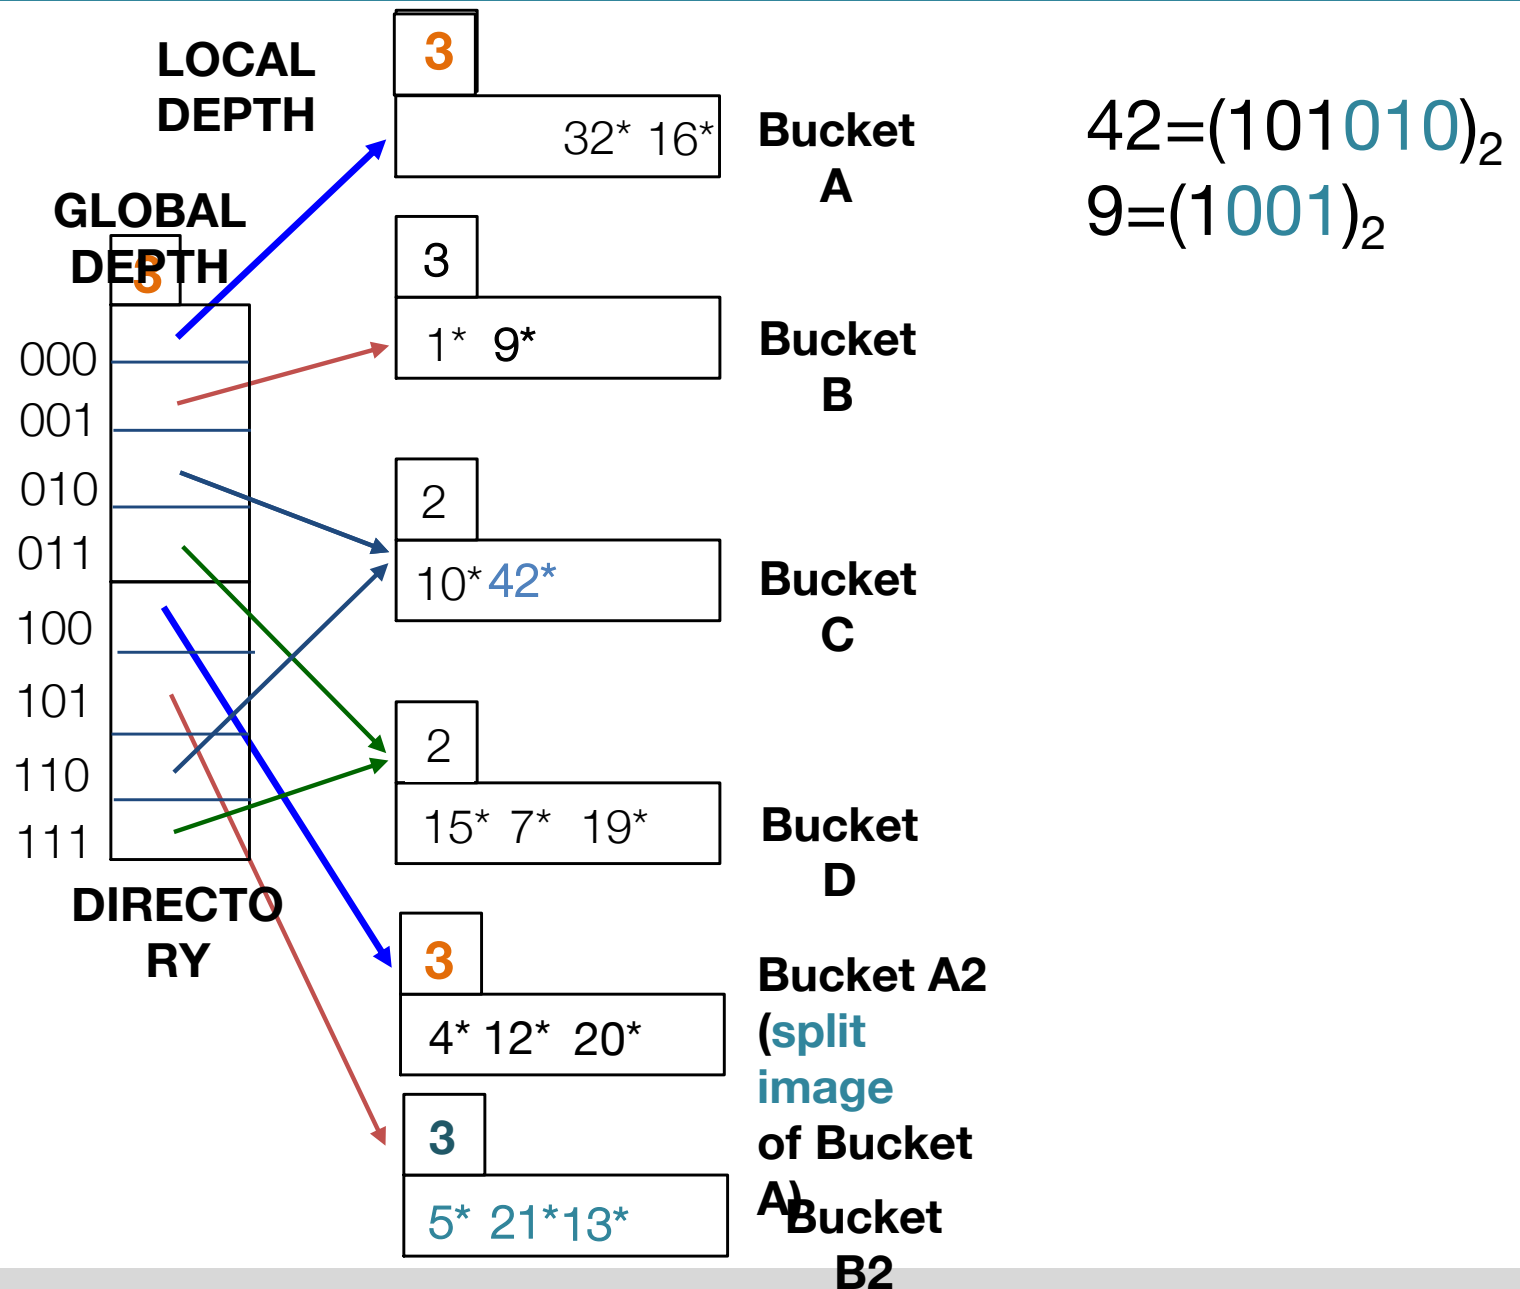
②

Extensible Hashing Exercise:

Insert 42, 9



Answer: Insert 42, 9




2nd technique: Linear Hashing

↓ assume a good hash
fun distribute things
evenly

- Overflow pages are allowed
- Buckets are split in a deterministic order (a Next pointer is maintained), rather than the overflowing bucket
- Upon an **overflow anywhere**, next bucket in the **deterministic order** is split.
- This means that overflow chains can occur
- Eventually, all buckets will split, **making the overflow chains rare**

2nd Technique: Linear Hashing


- Works in rounds
- Grows more gradually



Yellow and Orange are split image buckets.


Next identifies the bucket to be split next.

Buckets During Round i




Split image buckets
(created by splitting other
buckets in this round). h_{i+1}
uses an extra bit over h_i

Upon an overflow anywhere




Add an image bucket,
advance Next.


Overflow: Split next bucket, advance Next



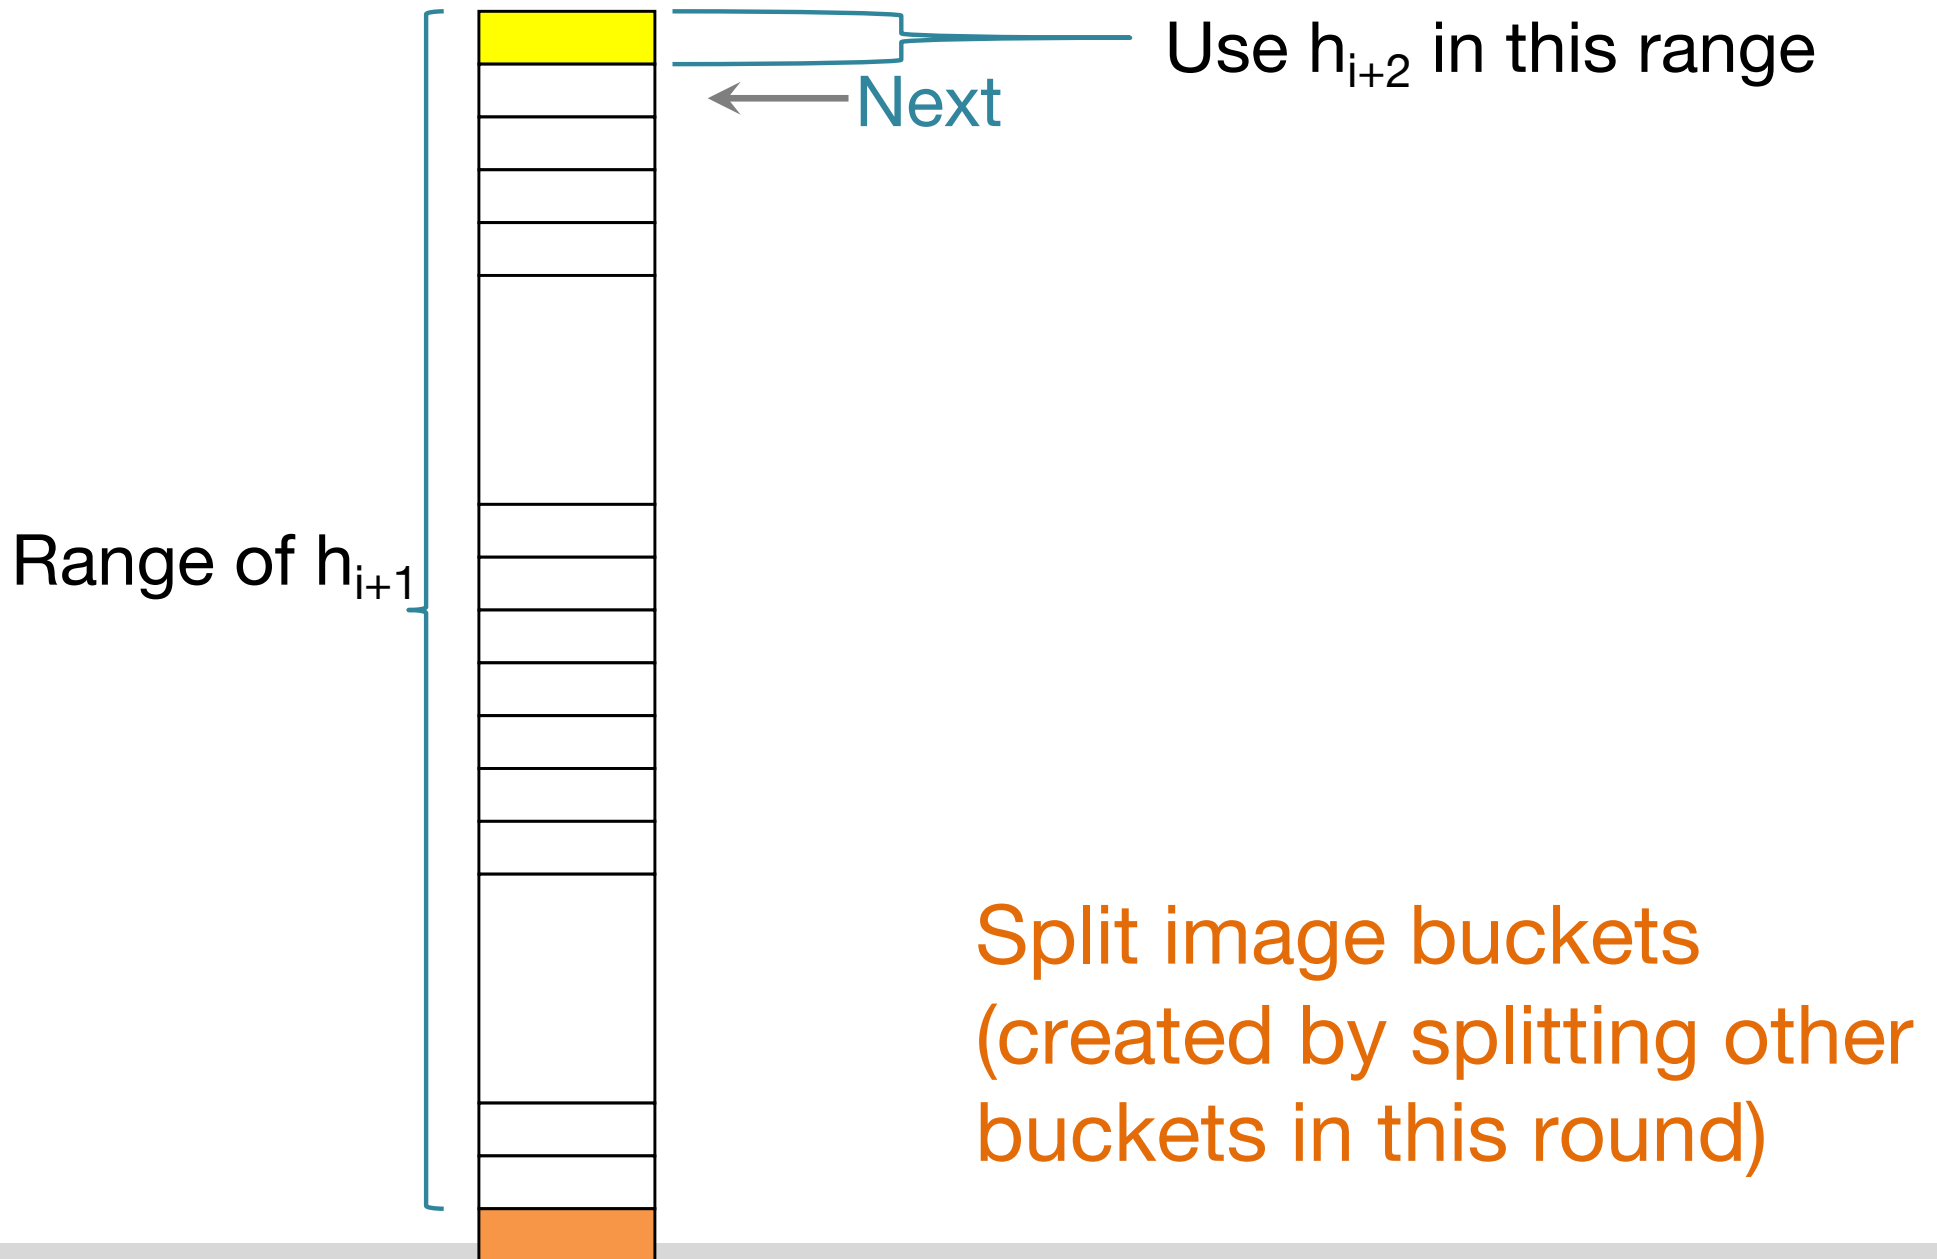
End of Round i = Start of Round $i+1$



End of Round i = Start of Round $i+1$



Round $i+1$ after the First Split



Linear Hashing: Details

- Splitting proceeds in rounds, starting from Level 0
 - During round $Level$, **only** h_{Level} and $h_{Level+1}$ are in use
 - Variables
 - $Level$: *(like global depth)* Initialized to 0 *level is not # of bits in extensible hashing*
 - $Next$: Pointer to the bucket next to be split
 - At the beginning of round # $Level$,
 - # buckets (some virtual) in the file = $N_{Level} = N * 2^{Level}$ *N : base directory size to start with, never start with 1*
- where N is initial number of buckets

Question??

In Linear Hashing, you always split the page pointed to by *Next*, even if the split is caused by an overflow to a different page.

A. True ✓

B. False

Linear Hashing: Hash Functions

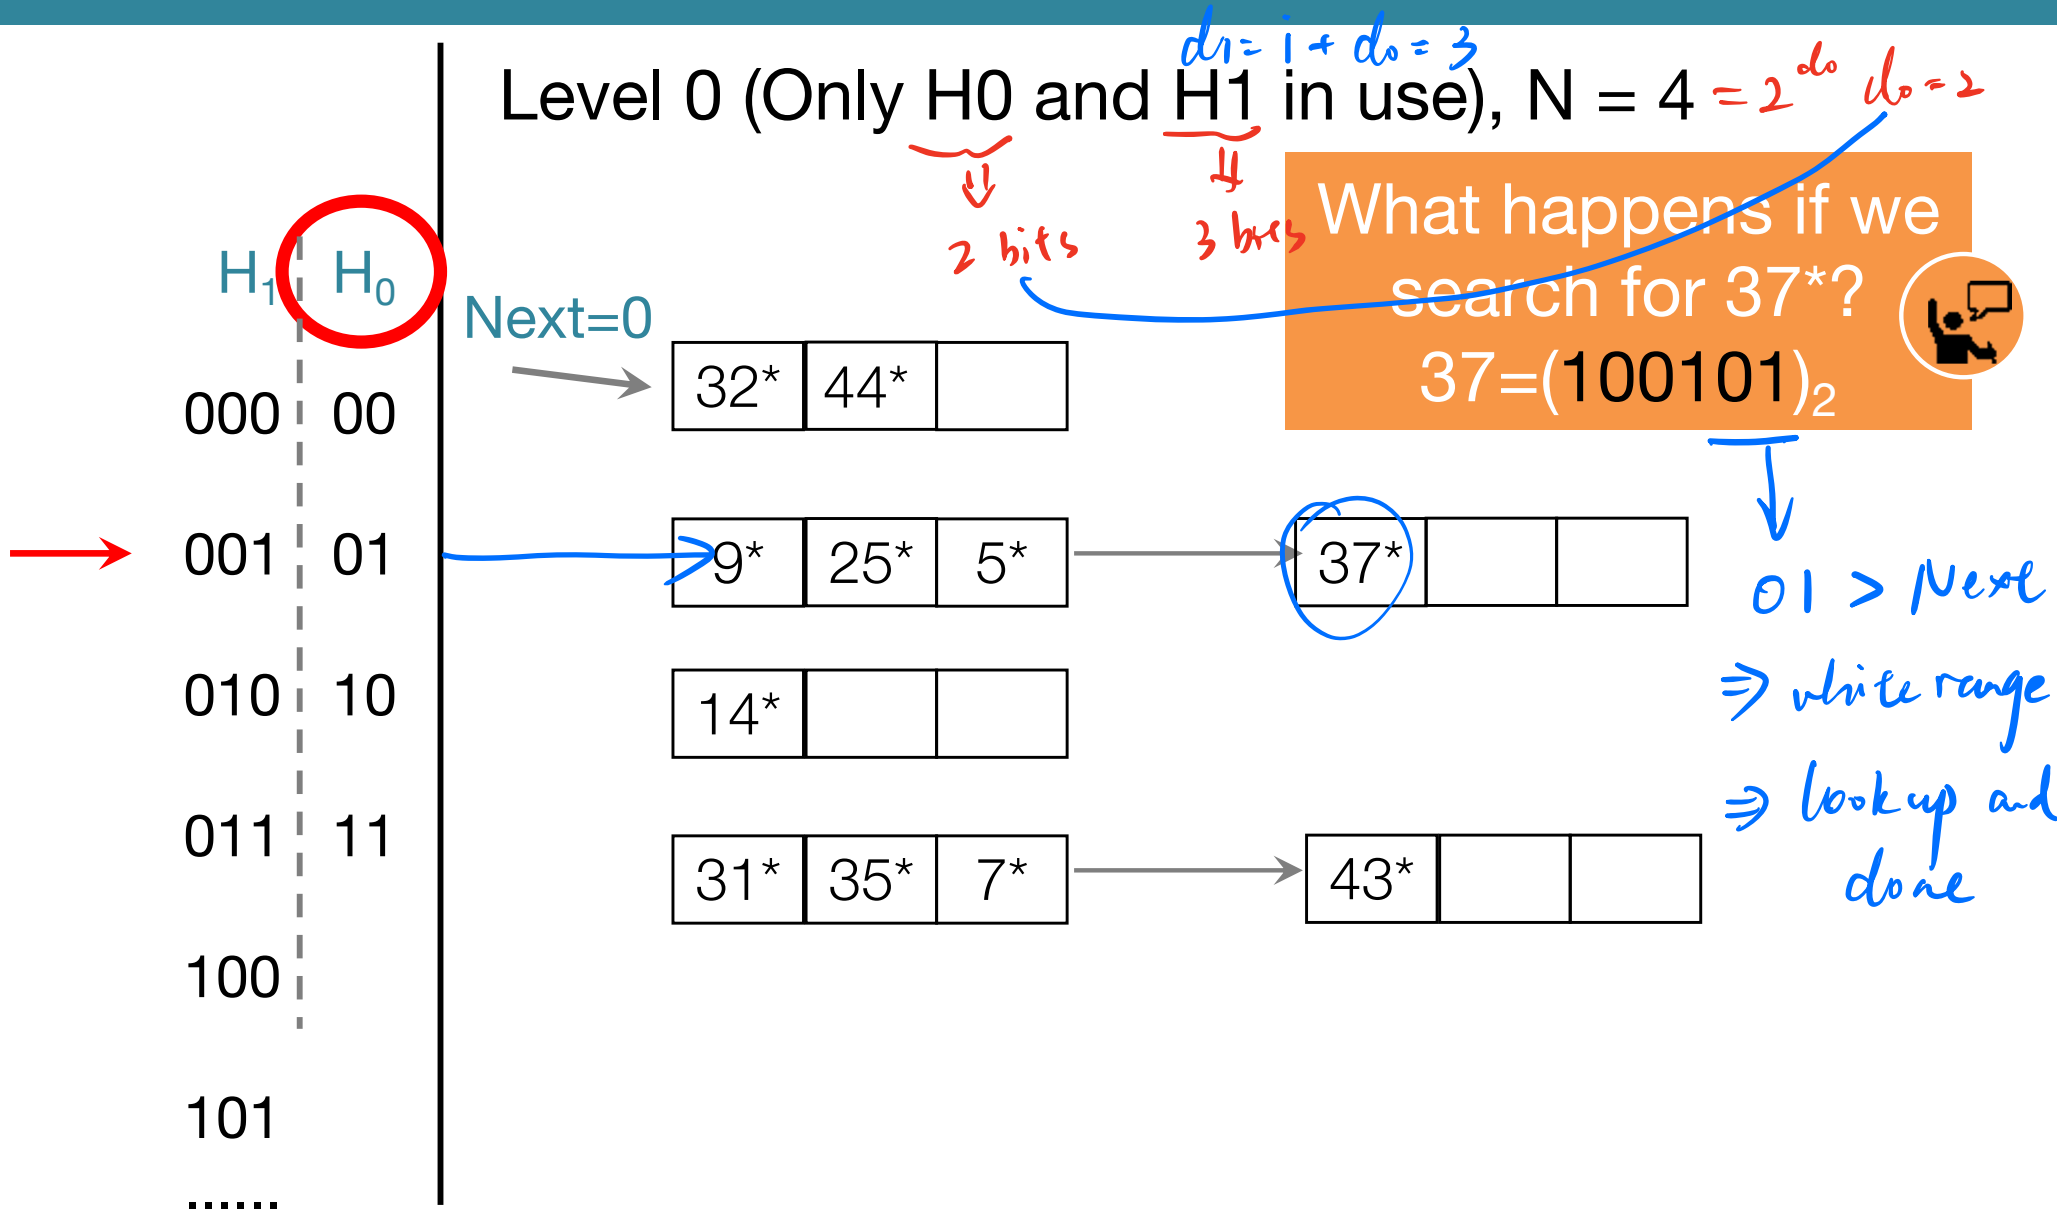
- Use a family of hash functions: H_0, H_1, H_2, \dots
 - H_{i+1} doubles the range of H_i (similar to directory doubling)
 - Hash family typically obtained by choosing a hash function $h()$ and initial number of buckets $N=2^{d_0}$
 - To compute $H_i(\text{value})$, apply hash function $h()$, and look at last d_i bits of the result where $d_i = d_0 + i$
 - In other words: $H_i(\text{value}) = h(\text{value}) \bmod (2^i N)$

Linear Hashing Algorithm

Search:

- Find bucket by applying H_{Level} *first*.
- Apply $H_{Level+1}$ if the value returned by $H_{Level} < \text{Next}$
- Search for the item in the bucket (possibly a chain of pages)

Linear Hashing Example #1



Linear Hashing Search Example #2

Level 0, $N = 4$

H_1	H_0
000	00
001	01
010	10
011	11
100	
101	
.....	

32*		
-----	--	--

9*	25*	17*
----	-----	-----

Next=2

14*		
-----	--	--

31*	35*	7*
-----	-----	----

43*		
-----	--	--

44*		
-----	--	--

5*	37*	
----	-----	--

What happens if we search for 37*?
 $37 = (100101)_2$

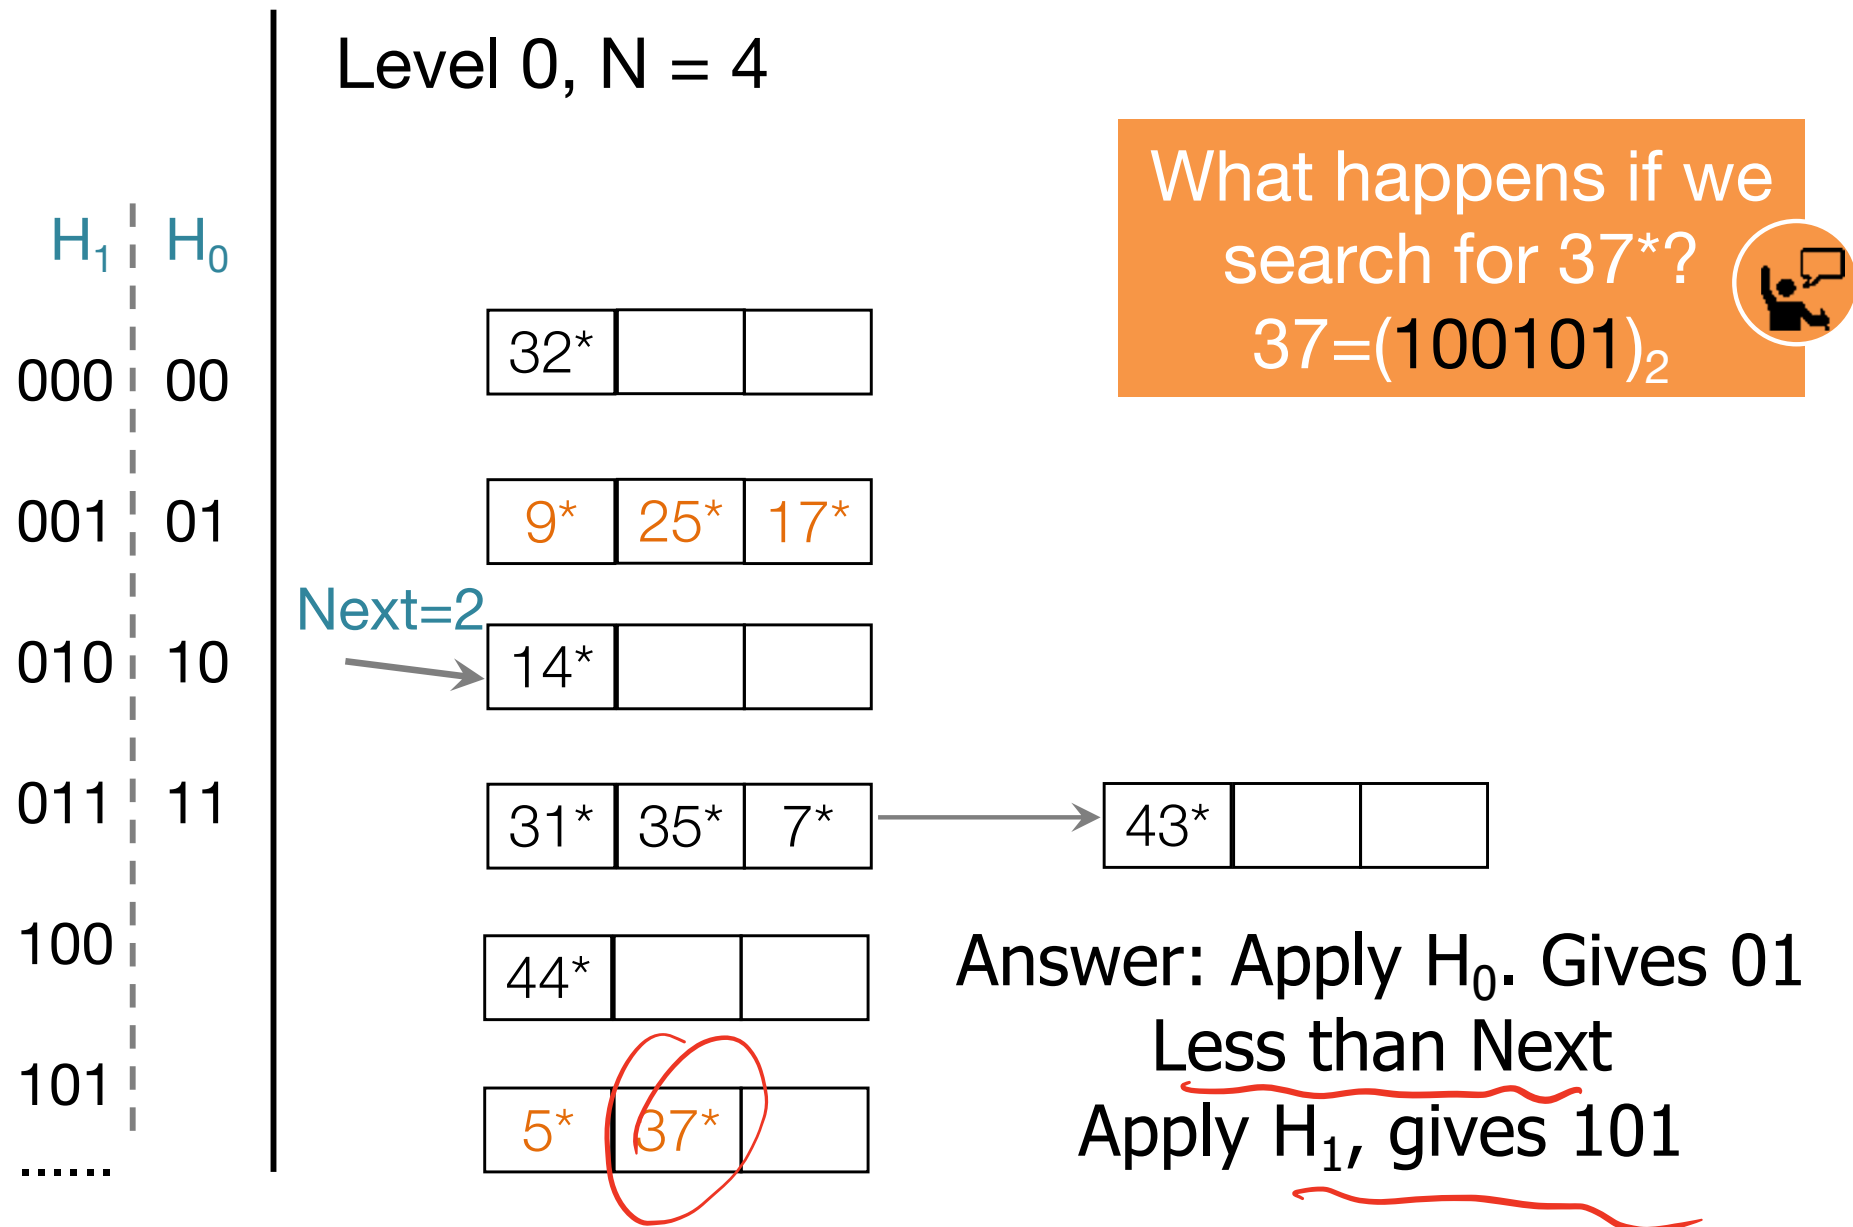


01 < Next : Yellow / Orange

go to 101 : look up



Linear Hashing Search Example #2



Linear Hashing -- Insert

- Find bucket by applying h_{Level}
 - Apply $h_{Level+1}$ if the value returned by $h_{Level} < \text{Next}$
- If bucket being inserted into is not full, insert normally. DONE
- Else, full page.
 - Add overflow page to the bucket and insert. (see special case #1)
 - Split *Next* bucket using $h_{Level+1}$ and increment *Next* (see special case #2)

Special case #1: If full bucket is *Next*, split first. (You may not need to also add an overflow page in the original bucket)

Special case #2: If $\text{Next} = N_{Level} - 1$ and a split is triggered






1. Split the bucket *Next* (i.e. the last bucket before the split image buckets)
2. Reset $\text{Next} = 0$
3. Start next round. $\text{Level} = \text{Level} + 1$

Linear Hashing Example

Level 0 (Only H0 and H1 in use), N = 4

H ₁	H ₀
000	00
001	01
010	10
011	11
100	
101	
.....	

Next=0



What happens if we add 43*?

$$43 = (101011)_2$$

11 > Next

Overflow! Bucket at Next needs to be split

split what next points to
++Next

Linear Hashing Example

Level 0, $N = 4$

H_1	H_0
000	00
001	01
010	10
011	11
100	
101	
.....	

Next=1

32*		
-----	--	--

9*	25*	5*
----	-----	----

14*		
-----	--	--

31*	35*	7*
-----	-----	----

 →

43*		
-----	--	--

44*		
-----	--	--

What happens if we
insert 17^* ?

$$17 = (10001)_2$$



This is not actually
stored

Wrong Answer

Level 0, N = 4

H_1 H_0

000 00

→ 001 01

010 10

011 11

100

101

.....

Next=1

32*

→ 9* 25* 5*

14*

31* 35* 7*

44*

→ 37* 17*

→ 43*


What happens if we
now insert 17*?

$$17 = (10001)_2$$



This is not actually
stored

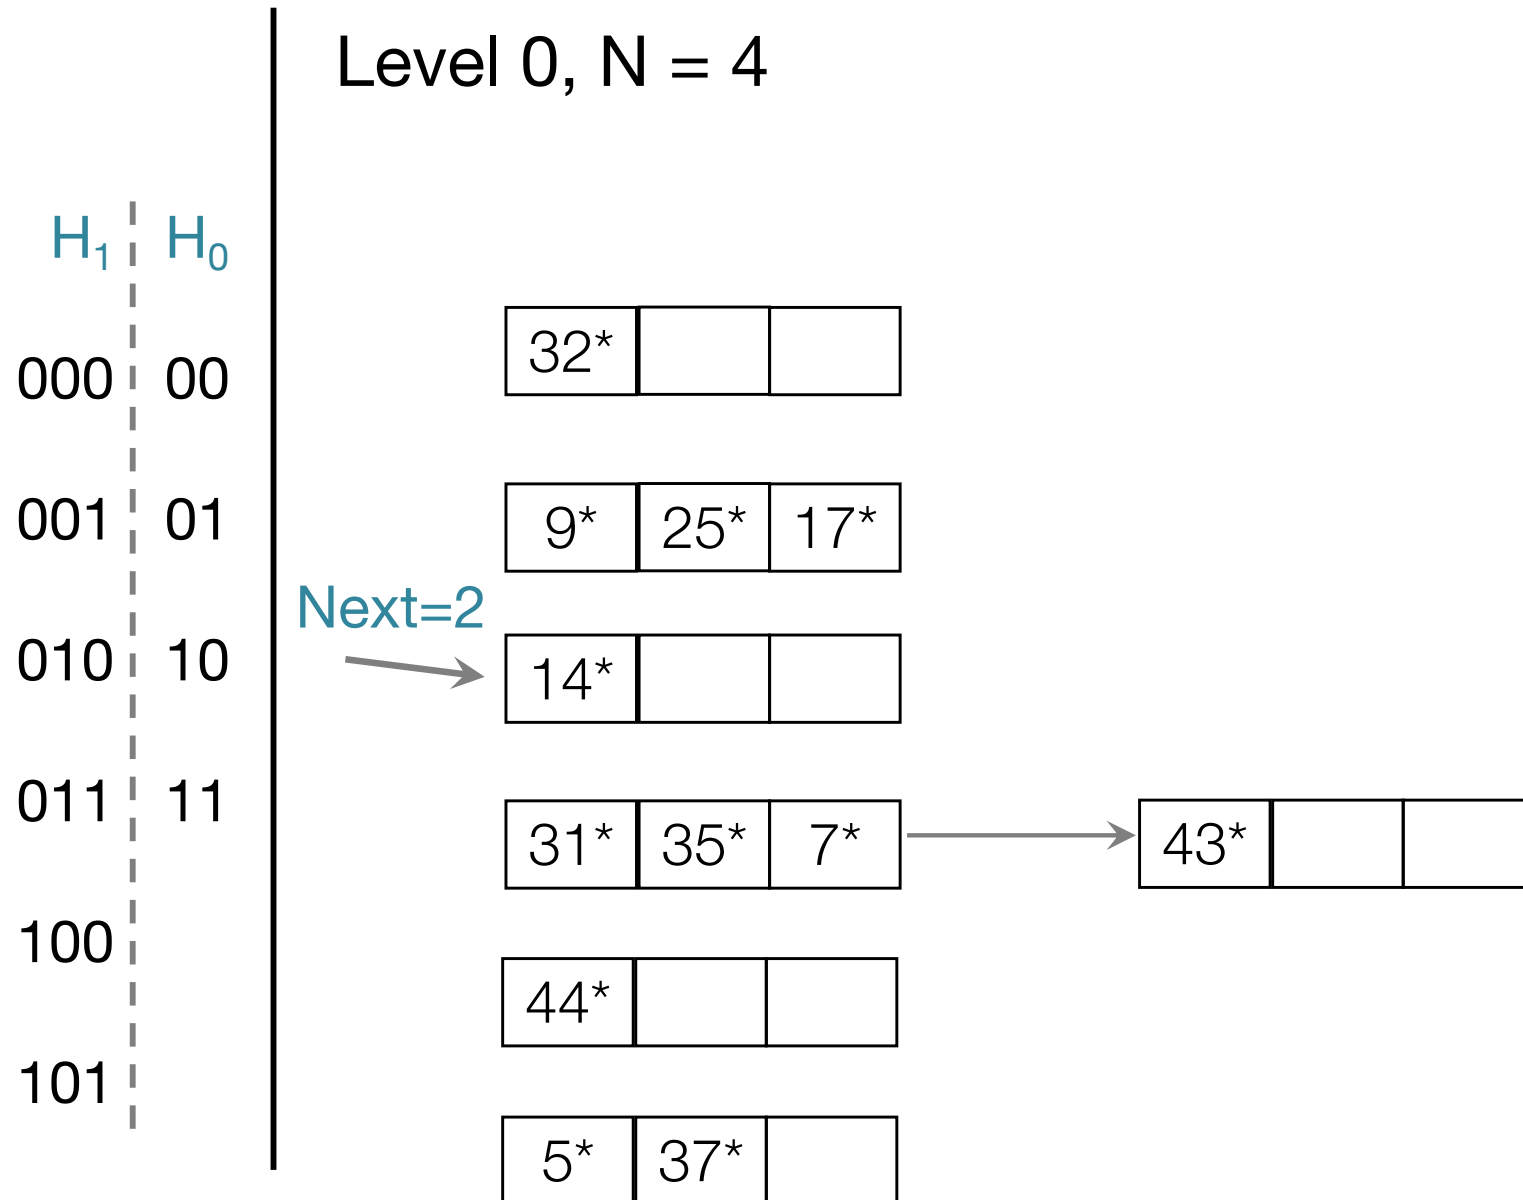
Correct Answer



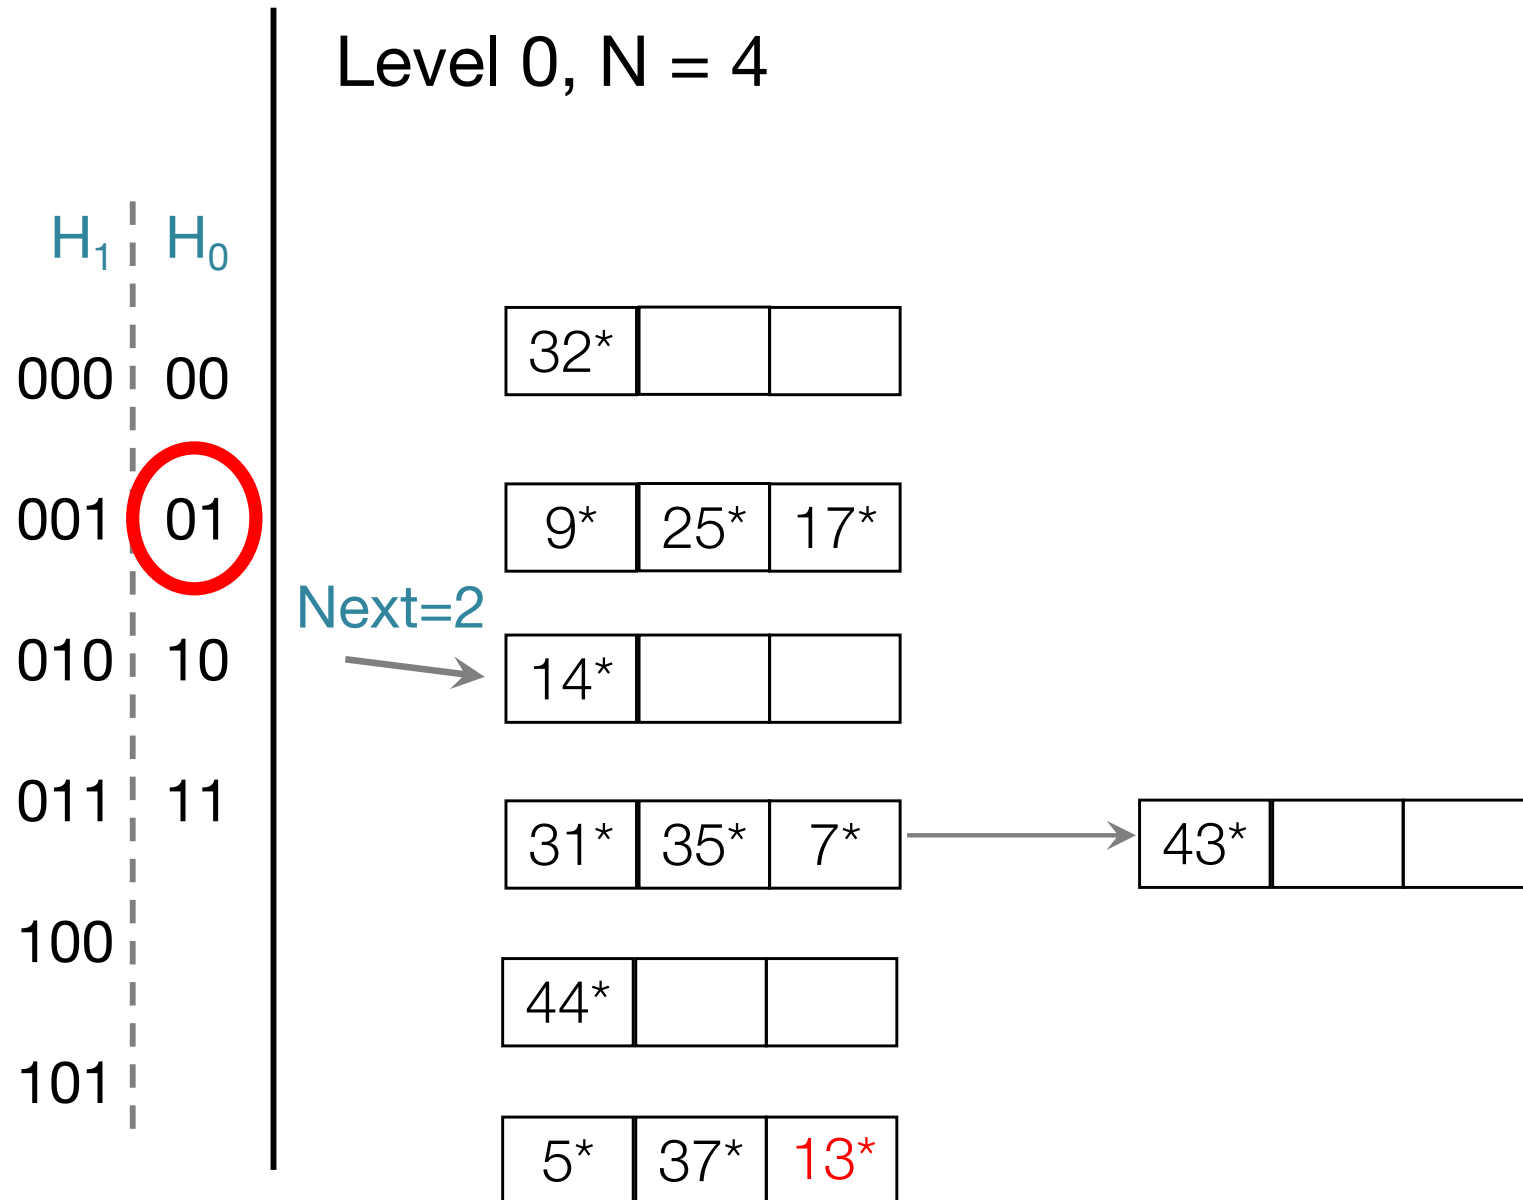
Linear Hashing (Contd.)

- Can choose any criterion to trigger split
 - ① e.g., split on an overflow (as in example) *← pay penalty of whole empty page*
 - ② e.g., space utilization on the page $> 90\%$ *⇒ pay penalty of little extra space you leave*
 - Since buckets are split round-robin, long overflow chains don't develop!
 - Deletes: see textbook
- need more buckets that 100% full
⇒ pay space penalty to avoid fetch penalty*


Exercise: Insert $13=(1101)_2$ & $29=(11101)_2$



Exercise: Insert $13=(1101)_2$ & $29=(11101)_2$



Exercise: Insert $13=(1101)_2$ & $29=(11101)_2$



Summary

- Discussed 3 kinds of hash-based indexes
- Static Hashing can lead to long overflow chains
- Extensible Hashing
 - Directory to keep track of buckets, doubles periodically
 - Always splits the “right” bucket
- Linear Hashing
 - Split buckets round-robin, and use overflow pages
 - Space utilization could be lower than Extensible Hashing

Optional Exercises

11.1, 11.3, 11.7, 11.9