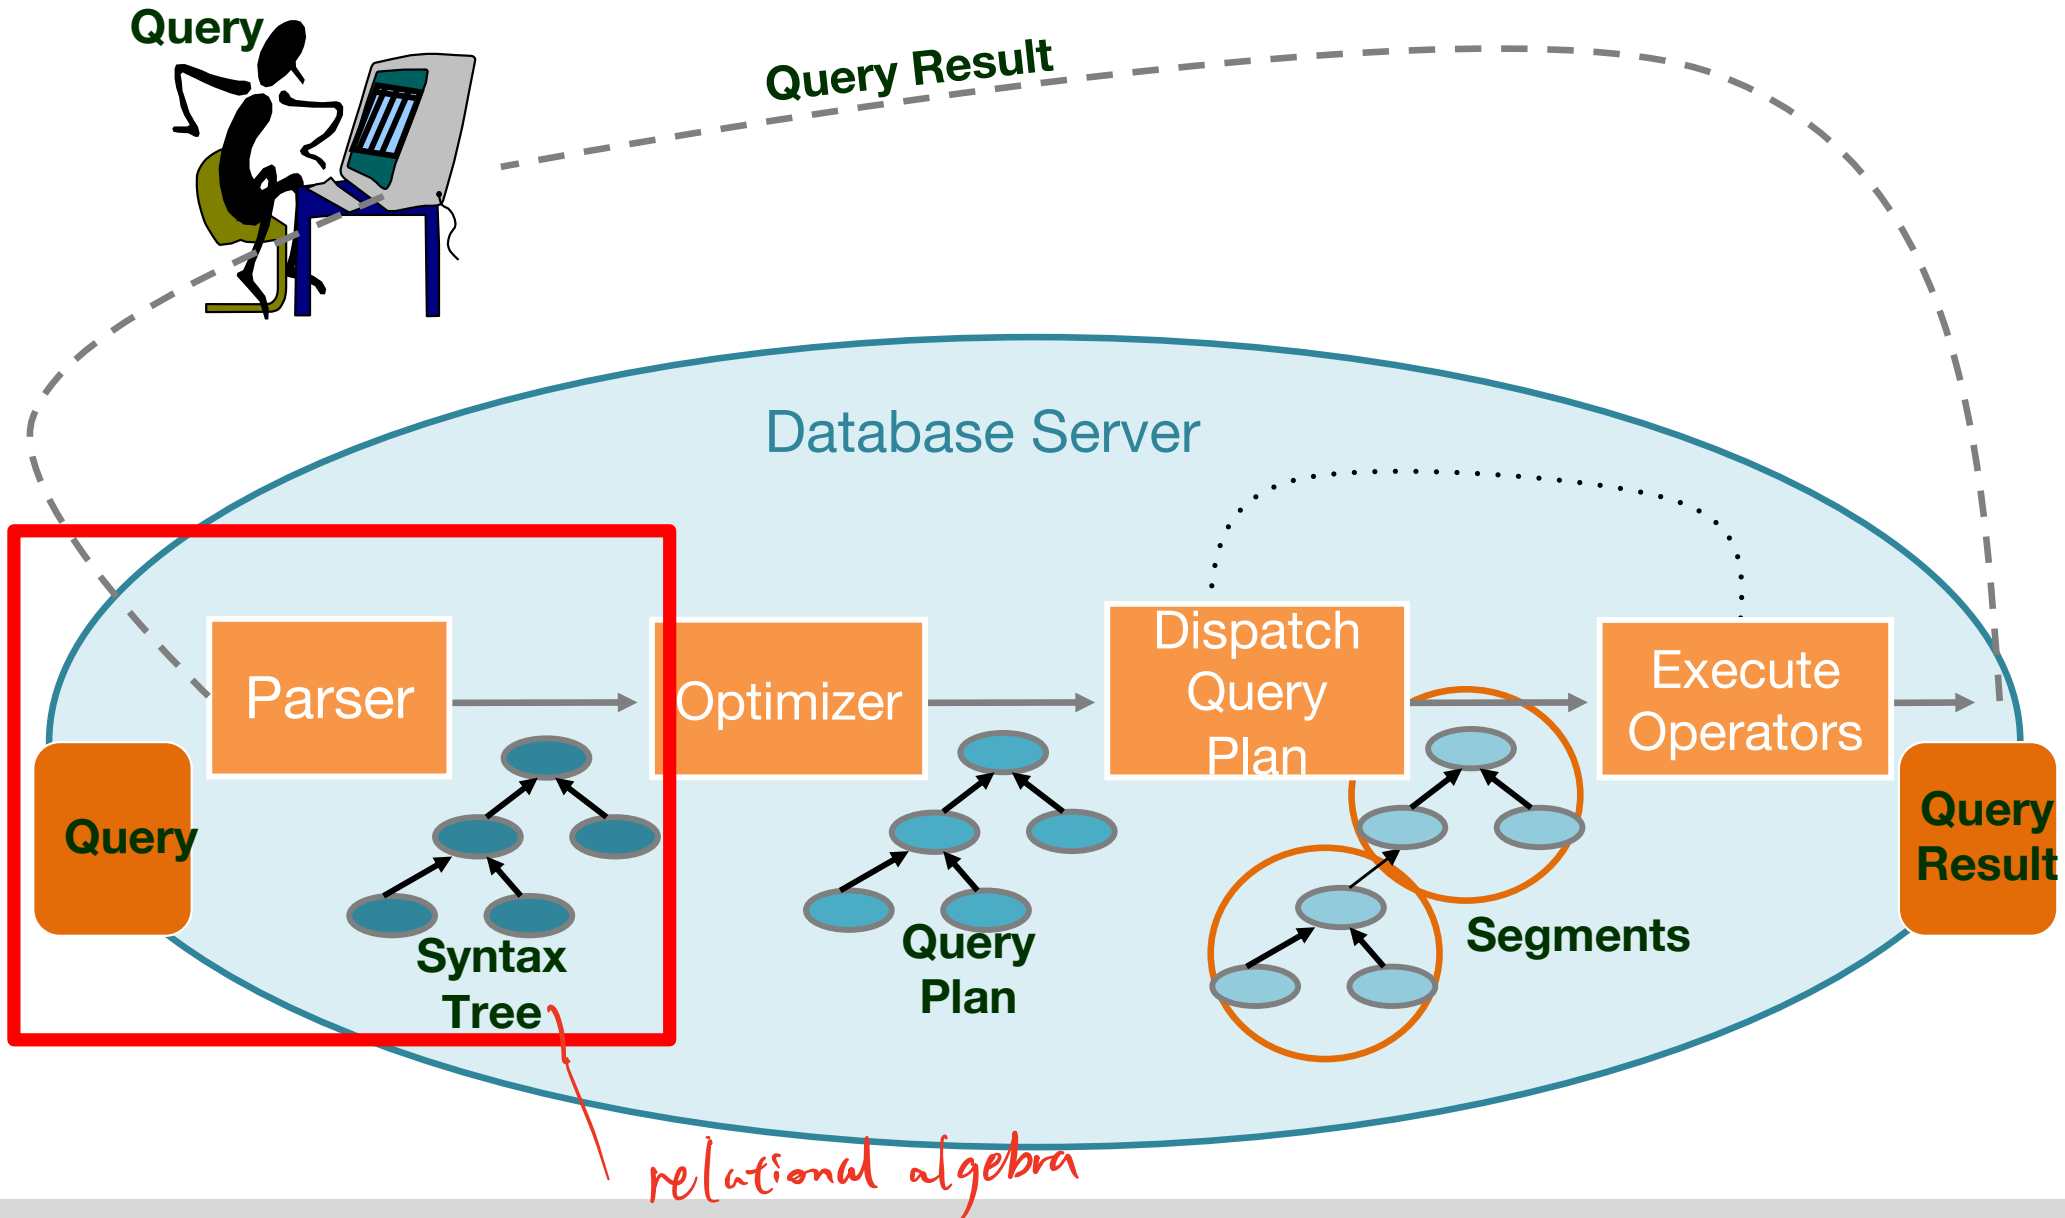


Evaluation of Relational Operations: Selection

Chapter 12


Query Execution Life-Cycle



Query Parsing



- Consult the meta information stored in **system catalogs**
- Create an initial **syntax tree** based on relational algebra tree

```
SELECT E.ename  
FROM Emp E, Dept D  
WHERE D.dname = 'Toy' AND  
      D.did = E.did
```



Question??

Which plan is cheaper?



System Catalogs

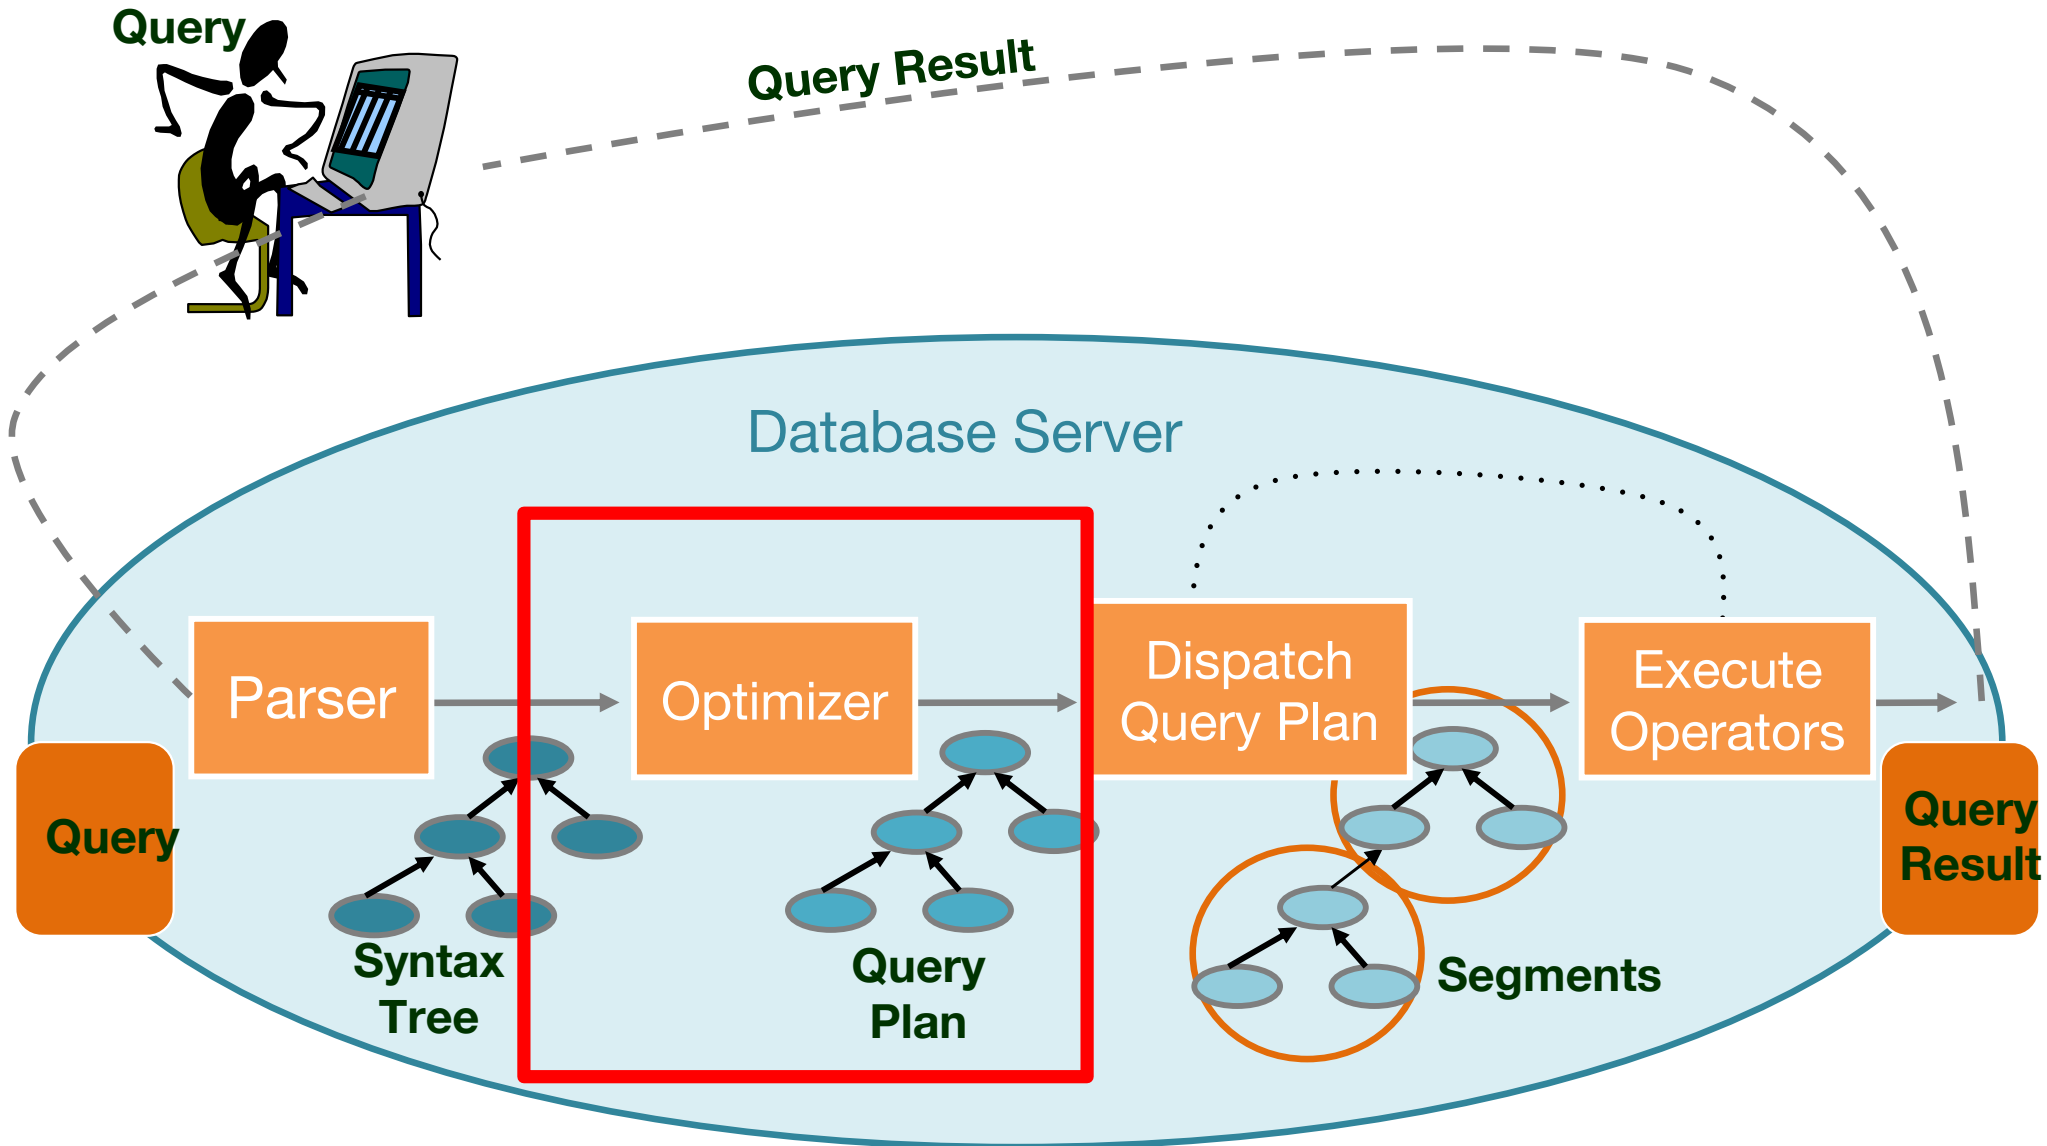
- To help optimize queries, the system keeps information on each **relation** *data records*
 - *table name* name, file name, file structure (e.g. heap file)
 - attribute name and type, for each attribute
 - index name, for each index
 - integrity constraints
- For each **index**:
 - structure (e.g. B+ tree) and search key fields
- For each **view**:
 - view name and definition
- Plus **statistics**, **authorization**, **buffer pool size**, etc.

Example Catalog: Attribute Catalog

attrName	relName	type	position
attrName	Attribute_cat	string	1
relName	Attribute_cat	string	2
type	Attribute_cat	string	3
position	Attribute_cat	integer	4
sid	Students	string	1
name	Students	string	2
login	Students	string	3
age	Students	integer	4
gpa	Students	real	5
fid	Faculty	string	1
fname	Faculty	string	2
sal	Faculty	real	3

Catalogs are themselves stored as relations

Query Execution Life-Cycle



Query Evaluation Plan



EMP (ssn, ename, addr, sal, did)

DEPT (did, dname, floor, mgr)

```
SELECT E.ename
FROM Emp E, Dept D
WHERE D.dname = 'Toy' AND
      D.did = E.did
```

EMP

ssn	ename	addr	sal	did
13b	Mary	First	9,000	1
29g	Jackie	Main	3,500	2

DEPT

did	dname	floor	mgr
1	Adv	5	Depp
2	Toy	10	Brown


Query Evaluation Plan



EMP (ssn, ename, addr, sal, did)

DEPT (did, dname, floor, mgr)

```
SELECT E.ename
FROM Emp E, Dept D
WHERE D.dname = 'Toy' AND
      D.did = E.did
```



EMP

ssn	ename	addr	sal	did
13b	Ma			
29g	Jackie Smith	Main	3,500	2



DEPT

did	dname	floor	mgr
			Depp
2	Sales	10	Brown


Query Optimizer selects the evaluation plan

Query Optimization Example

Modify to:



Query Optimization



2. Will compare different evaluation algorithms for SELECT, JOIN, and PROJECT operator based on whether an index is available, table sizes, etc.

Query Optimization

Query optimizer selects an evaluation plan with the least cost in two steps:

1. Plan Enumeration

- Different query plans
- Different implementations (i.e., evaluation algorithms) for each operator ~~☆~~

In 2 weeks

Today →

2. Cost Estimation

- Cost of each operator
- Overall cost of the plan

Operator Evaluation

- How to implement common operators?



- Selection

- Matching Indexes (*previously created*),
- Join
- Projection (optional DISTINCT)
- Set Difference
- Union
- Aggregate operators (SUM, MIN, MAX, AVG)
- GROUP BY

Selection (on one table)

- Access path defines a strategy to do a selection on a table, possibly utilizing an index

- Example of a selection condition

- a predicate: $\text{gpa} > 3.0$ and $\text{age} = 21$

- Some possible access paths

- File scan

- Index that *matches* a selection in the query. Examples:

- B+tree index on the $\langle \text{gpa}, \text{age} \rangle$ attributes \Rightarrow more leaf

- B+tree index on gpa

- B+tree index on age

- Hash index on age



hash index
(equality index \Rightarrow not usable for $>$)

first
sort
key
secondary
sort
key

\Rightarrow more leaf
pages in
the tree
but more efficient
data access

Selection Cost

- `WHERE R.a op value`
- Options:
 - Heap file Cost: $O(N)$
 - Sorted File Cost: $O(\log_2 N) + \dots$
 - Index
 - Hash Cost: $O(1) + \dots$
 - B+Tree: Clustered/Unclustered Cost: $O(\log_F N) + \dots$

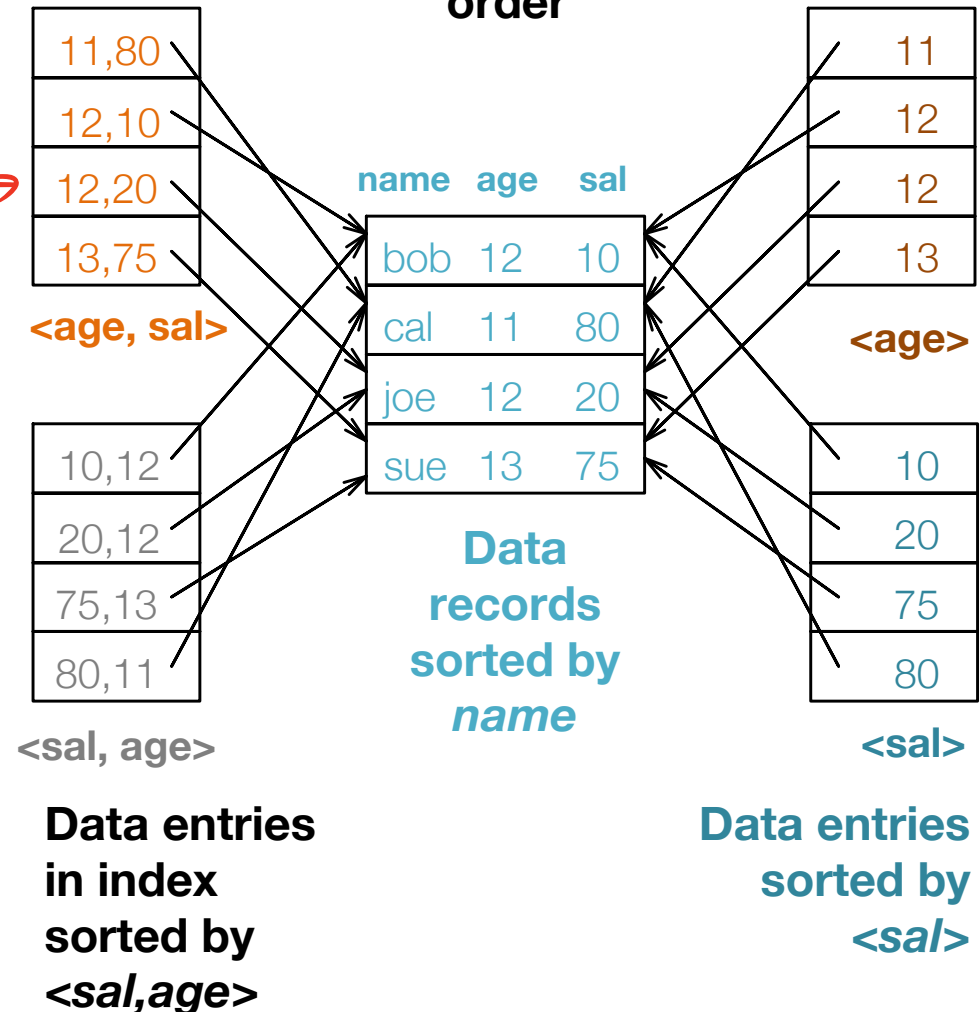
Composite Search Keys

- Index on <age>: search key is a single-attribute age
- Index on <name, age>: search key is **composite** (name, age) pair
 - For B+-tree, name is the primary comparison attribute and age matters only when names are equal
 - For hash-index, $h((\text{name}, \text{age}))$ used – both name and age are needed to hash

Indexes with Composite Search Keys

- **Composite Search Keys:** Search on a combination of fields
 - **Equality query:** Every field value is equal to a constant value. e.g. wrt $\langle \text{sal}, \text{age} \rangle$ index:
 - age = 12 and sal = 75
 - **Range query:** Some field value is not a constant
 - e.g. age = 12 and sal > 10
- Data entries in index sorted by search key to support range queries

Examples of composite key indexes using lexicographic order



Index Matching

- When can we use an index to evaluate a selection predicate? *when match*
- An index *matches* a predicate if the index can be used to evaluate (at least part of) the predicate

Index Matching

- Index on $\langle a, b, c \rangle$

- $a=5$ and $b=3$

- $a > 5$ and $b < 3$

- $b=3$

Index matches (part of) a predicate if:

- Conjunction of terms involving only attributes (no disjunctions)
- **Hash:** only equality operation, predicate has all index attributes
- **Tree:** Attributes are a prefix of the search key, any ops

Index Matching

- Index on $\langle a, b, c \rangle$

sorted by a, b, c
→

Tree Idx

Hash Idx

- $a=5$ and $b=3$

• yes

• no! *c is required*

- $a > 5$ and $b < 3$ *$a=6, b<3, a=7, b<3 \Rightarrow$*

• yes

• no!

- $b=3$

No

No

- $a=7$ and $b=5$ and $c=4$ and $d>4$

• yes

• yes

- $a=7$ and $c=5$

• yes

• no!

Index matches (part of) a predicate if:

- Conjunction of terms involving only attributes (no disjunctions)
- Hash: only equality operation, predicate has all index attributes
- Tree: Attributes are a prefix of the search key, any ops

like $a=7$ for index $\langle a, b, c \rangle$

Index Selectivity



- To retrieve Emp records with $\text{age}=30$ AND $\text{sal}=4000$, an index on $\langle \text{age}, \text{sal} \rangle$ would be better than an index on age or an index on sal
 - $\langle \text{age}, \text{sal} \rangle$ is more **selective** than just $\langle \text{age} \rangle$ or just $\langle \text{sal} \rangle$
 - It identifies fewer spurious records that will later be rejected
- If condition is: $\text{age}=80$ AND $3000 < \text{sal} < 20,000$:
 - $\langle \text{age} \rangle$ index much better than $\langle \text{sal} \rangle$ index!
- Ideally, the more **selective index preferred**
apply more selective index first

Example: Index Matching

- Predicate could match more than 1 index
- Hash index on $\langle a, b \rangle$ ^{specify 2 thing} and B+tree index on $\langle a, c \rangle$
- Predicate: $a=7$ and $b=5$ and $c=4$. ^{specify 3 things} Which index?

① Option 0: Neither. Simply use file scan

② Option 1: More selective one. Then, scan among the selected records

③ Option2: Use both!
 - hash index $\langle 7, 5 \rangle$
 - intersect and tree index $a=7$ $c=4$

- Algorithm: Intersect ^② rid sets
- Sort ^① rids, retrieve rids in both sets

Example Choices

- Hash index on $\langle a \rangle$ and Hash index on $\langle b \rangle$

- $a=7$ or $b>5$

force to do full scan for b pure
Which index?



*⇒ can check $a=7$ during scan
⇒ no need hash look up for a*

- ✓ Neither! File scan required for $b>5$

no hash index on b

- Hash index on $\langle a \rangle$ and B+tree on $\langle b \rangle$

- $a=7$ or $b>5$

Which index?



- Option 1: Neither

hash on $a=7$ ∪ tree on $b>5$.

- Option 2: Use both! Fetch rids and union

- Note: Option 1 could be better sometimes. (When?)

B+tree maybe unclustered

⇒ worse than full scan.

Question??

- Hash index on $\langle a \rangle$ and B+tree on $\langle b \rangle$

- $(a=7 \text{ or } c>5) \text{ and } b > 5$

no index \Rightarrow file scan \Rightarrow so \bigcirc has to full scan

Which index choices should we consider?

- A. None (file scan) and Tree
- B. None and Hash
- C. None, Hash, Tree, and Both
- D. None, Tree, and Both
- E. None, Hash, and Both

Question??

- Hash index on <a> and B+tree on
 - (a=7 **or** c>5) **and** b > 5

Which index choices should we consider?

- A. None (file scan) and Tree**
- B. None and Hash
- C. None, Hash, Tree, and Both
- D. None, Tree, and Both
- E. None, Hash, and Both

Why is Hash not appropriate?

Among None and Tree, what factors will determine the choice? In other words, when is the optimizer likely to prefer the tree-based Index and when it is likely to simply use a file Scan?

When to Use a B+tree Index?

- Consider a relation with 1M tuples
- 100 tuples on a page
- 500 (key, rid) pairs on a page
- Fill factor in index: 0.67
- Plan a selection query for which the index matches

1,000,000

(data records)
data pages
= $1M / 100 = 10K$ pages
leaf idx pgs (leaf page)
= $1M / (500 * 0.67)$
~ 3K pages

in leaf of b+ tree
⇒

3K x 1%

	1% SELECTIVITY	10% SELECTIVITY
CLUSTERED	~ 30 + 100	~ 300 + 1000
NON-CLUSTERED	~ 30 + 10,000	~ 300 + 100,000
NC + SORT RIDS by page ID.	~ 30 + (~10,000)	~ 300 + (~10,000)

non-clustered
(table still
not clustered)

⇒ Choice of Index access plan, consider:
1. Index Selectivity 2. Clustering
⇒ Similar consideration for hash-based indices

Summary

- DBMS use catalogs, which are relations themselves.
- Query Optimizer selects the evaluation plan after evaluating the cost of *many* plans. It evaluates the potential algorithms and chooses the best one per operator.
- Single vs. composite search keys
- Each relation may have multiple indexes
 - Selectivity
 - Clustering

Optional Exercises

12.1 (1-4), 12.3, 12.5